

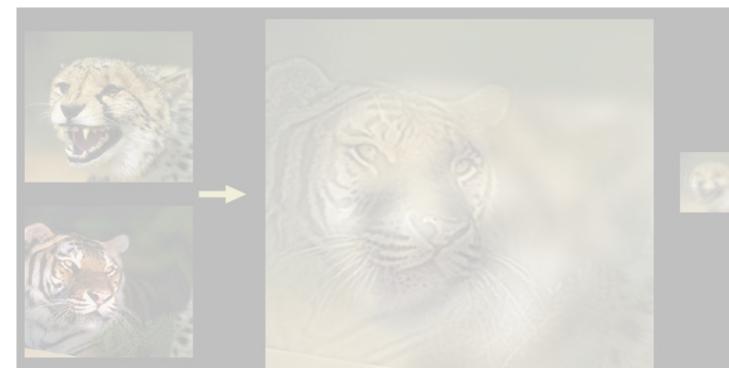


CS 4476 Computer Vision:

3. Image Filtering

Humphrey Shi

8/27/2024



Outline

- Logistics & Course Recap
- Image Filtering

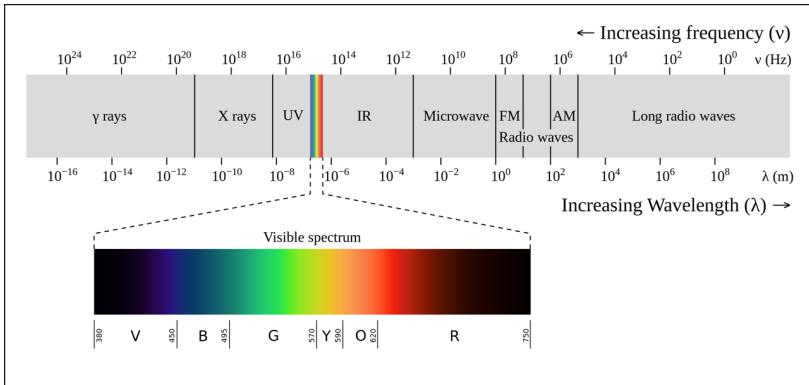
Logistics & Course Recap

Logistics

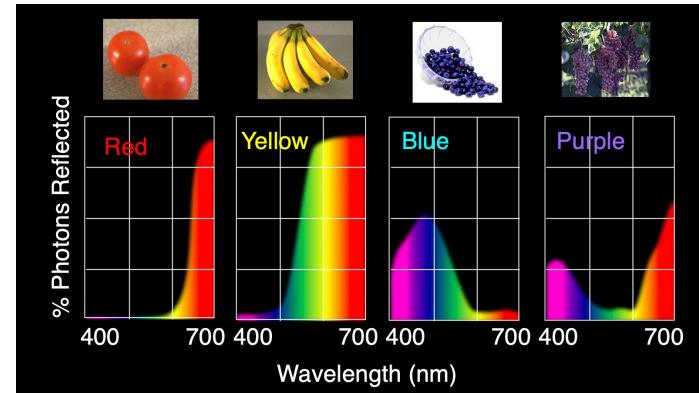
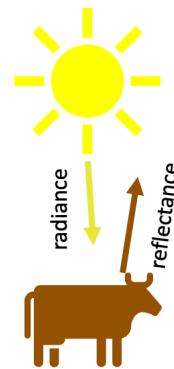
- Assignment 0 & 1 out on Canvas/GradeScope.
- TA office hours will be out this week.
- Claim your extra participation credit on Ed Discussion.

Recap: Light & Color

- **Color** is the result of interaction between physical **Light** in the environment and **our visual system**.



Visible light (400~700nm)



Light interacts with environments

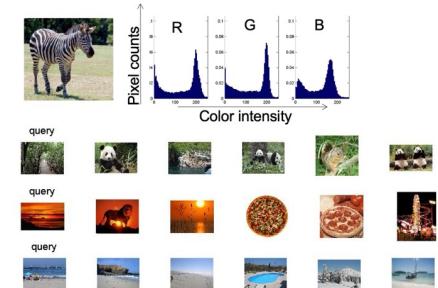
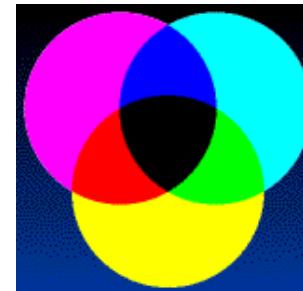
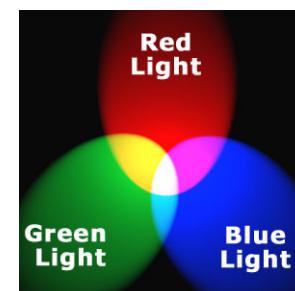
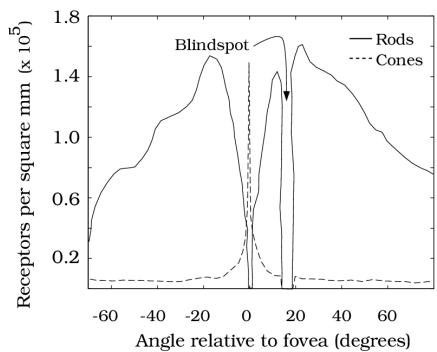
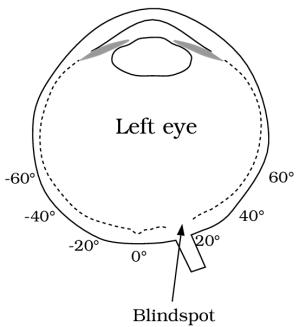


Image Filtering

Digital Image Formation

(Review) What is in an Image?



```
from PIL import Image
import numpy as np

image = Image.open('./tech.jpg')
image_array = np.array(image)
print(image_array.shape)
print(image_array)
```

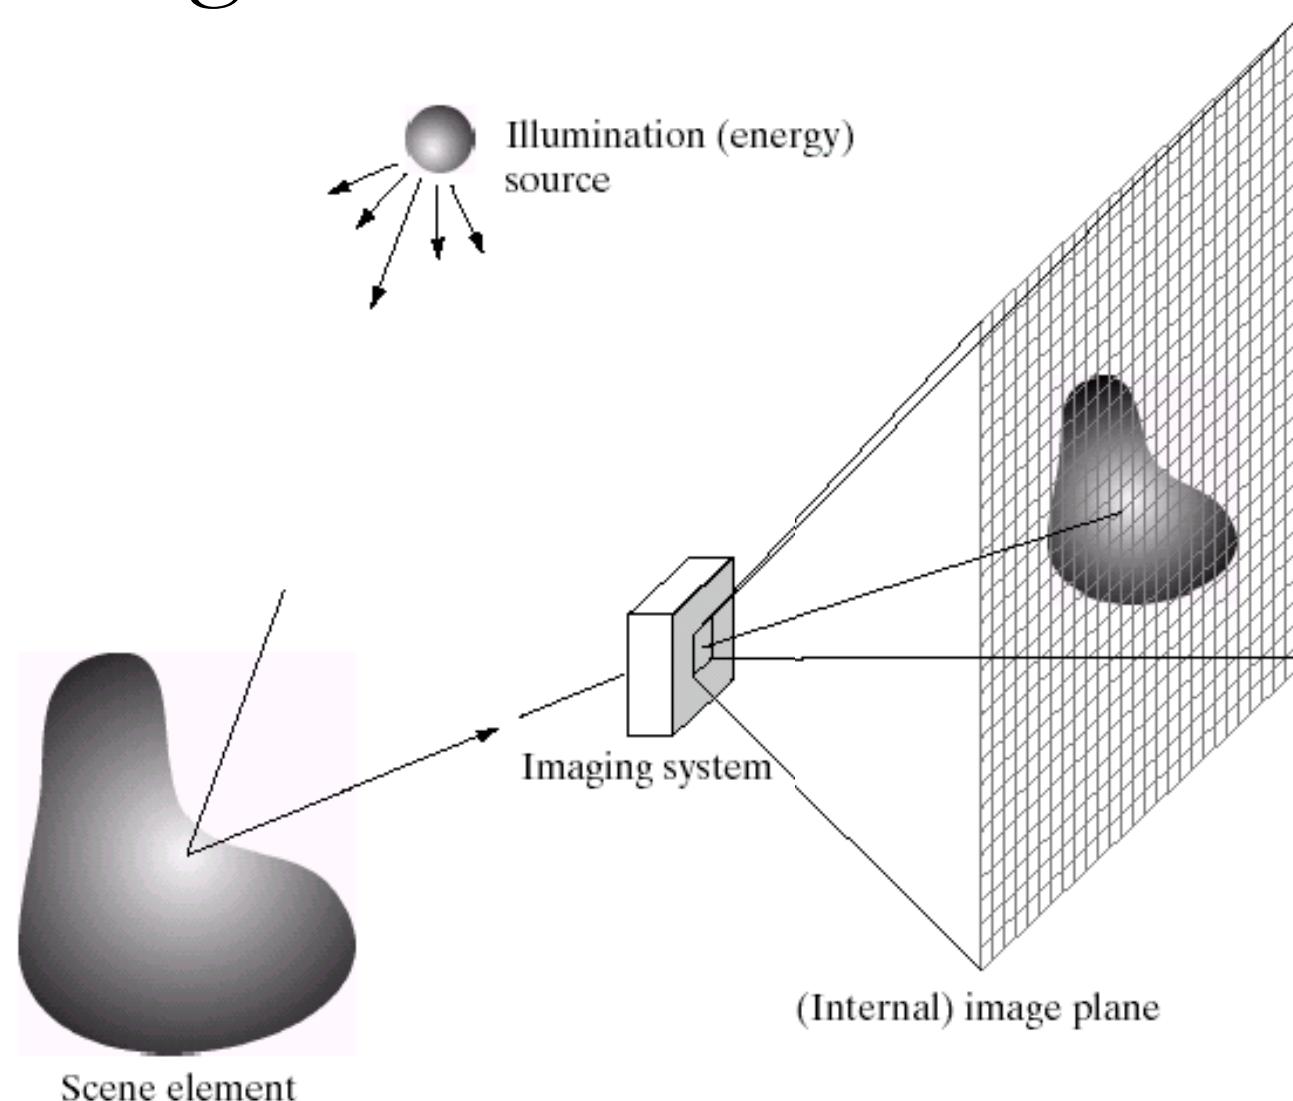
```
(1187, 949)
[[253 224 229 ... 216 216 216]
 [250 151 143 ... 112 112 112]
 [251 142 150 ... 112 112 112]
 ...
 [254 66 68 ... 80 80 80]
 [255 253 255 ... 254 254 254]
 [255 255 248 ... 255 255 255]]
```



```
(1187, 949, 3)
[[[249 254 255]
 [217 222 251]
 [224 226 255]
 ...
 [188 222 255]
 [188 222 255]
 [188 222 255]]
 [[247 251 255]
 [144 148 185]
 [135 138 193]
 ...
 [ 76 116 188]
 [ 74 116 188]
 [ 76 116 188]]
 [[248 251 255]
 [135 136 199]
 [142 142 214]
 ...
 [ 63 119 204]
 [ 62 120 204]
 [ 63 119 204]]
 ...
 [[254 254 254]
 [ 66 66 66]
 [ 68 68 68]
 ...
 [ 76 80 89]
 [ 74 81 87]
 [ 76 80 89]]
 [[255 255 255]
 [253 253 253]
 [255 255 255]
 ...
 [252 255 255]
 [252 255 255]
 [252 255 255]]
 [[255 255 255]
 [255 255 255]
 [248 248 248]
 ...
 [254 255 255]
 [254 255 255]
 [254 255 255]]]
```

- Images are made by arrays of numbers of brightness/colors.

How are Images Formed?



Digital Camera

A digital camera replaces film with a sensor array

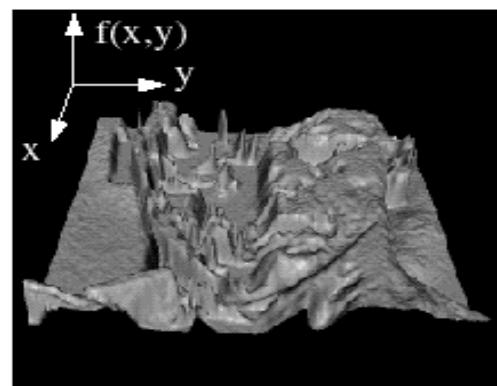
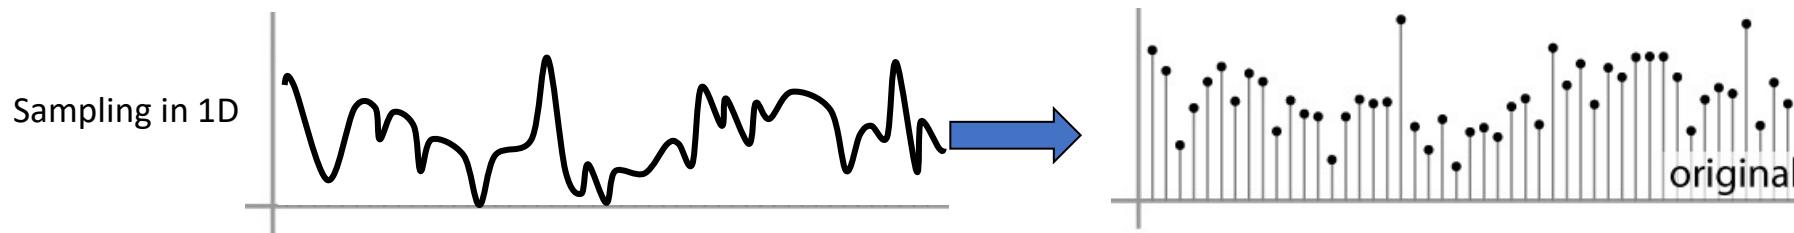
- Each cell in the array is light-sensitive diode that converts photons to electrons
- <http://electronics.howstuffworks.com/digital-camera.htm>



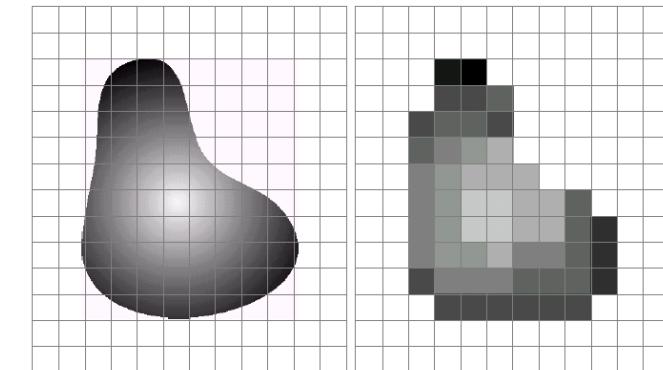
Image source: <https://www.pinterest.se/pin/849984129648395414/>

Digital Images

- **Sample** the 2D space on a regular grid
- **Quantize** each sample (round to nearest integer) **Why?**
- Image thus represented as a matrix of integer values.



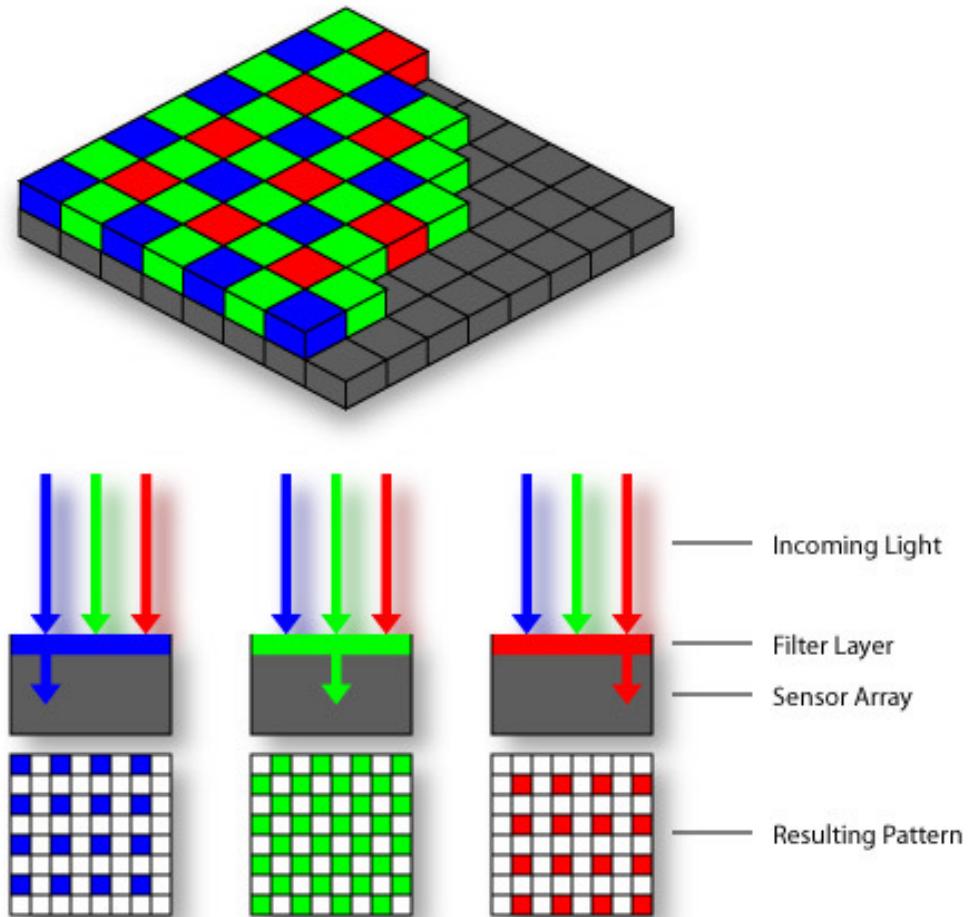
j							
	i						
62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30



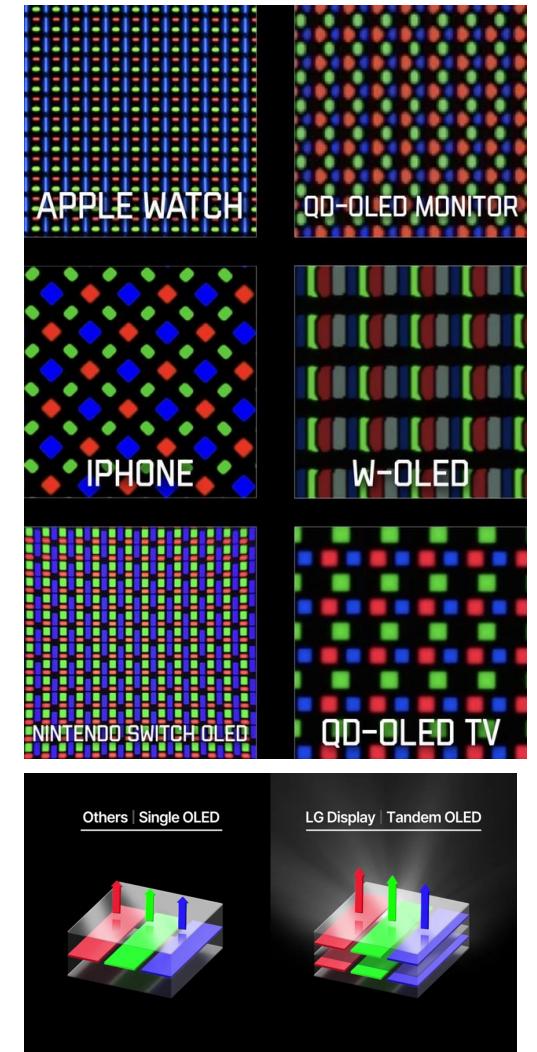
Continuous Image

Quantized/Sampled Image

Digital Color Images



Estimate RGB
at 'G' cells from
neighboring values



Digital color images

Color images,
RGB color space



R



G



B

Image Filtering

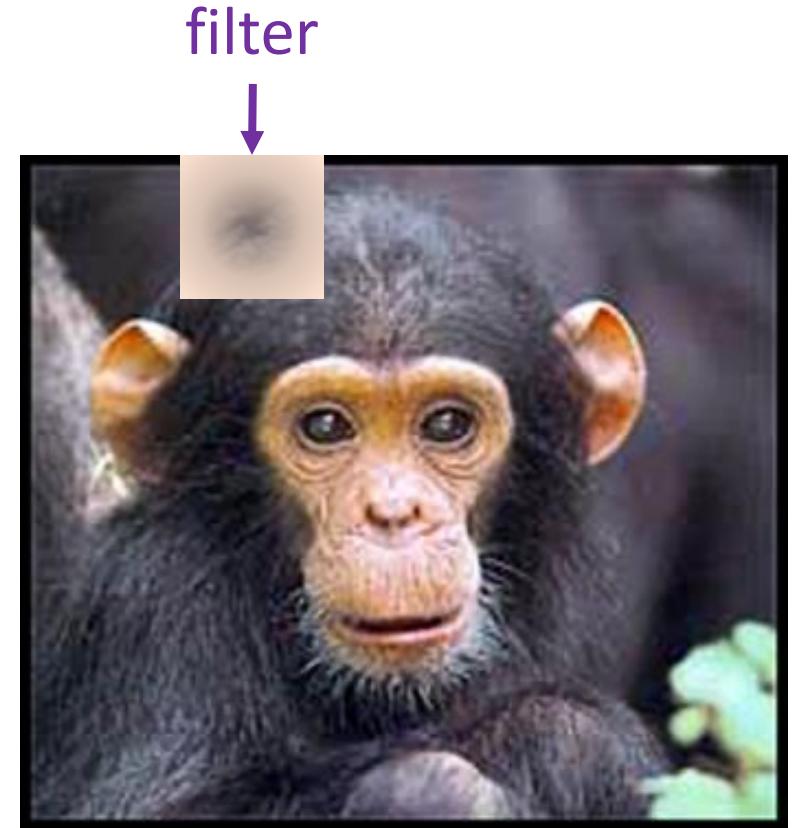
Image Noise, Filtering and Convolution

Image Filtering

- **Idea:** Compute a function of the local neighborhood at each pixel in the image
 - Function specified by a “filter” or mask saying how to combine values from neighbors.

- **Uses of filtering:**

- Enhance an image (denoise, resize, etc)
- Extract information (texture, edges, etc)
- Detect patterns (template matching)



Motivation: Image Noise reduction

- A long-studied area in signal/image processing and computer vision.

1972

- [j6] Grant B. Anderson, Thomas S. Huang:
Correction to "Piecewise Fourier Transformation for Picture Bandwidth Compression". IEEE Trans. Commun. 20(3): 488-492 (1972)

1971

- [j5] Thomas S. Huang:
How the Fast Fourier Transform Got its Name. Computer 4(3): 15 (1971)
- [j4] Grant B. Anderson, Thomas S. Huang:
Frequency-domain image errors. Pattern Recognit. 3(2): 185-192 (1971)

1970

- [j3] Thomas S. Huang:
Introduction to the special issue on image enhancement. Pattern Recognit. 2(2): 67-68 (1970)

[–] 1960 – 1969 ⓘ

1969

- [c1] Grant B. Anderson, Thomas S. Huang:
Errors in frequency-domain processing of images. AFIPS Spring Joint Computing Conference 1969: 173-185

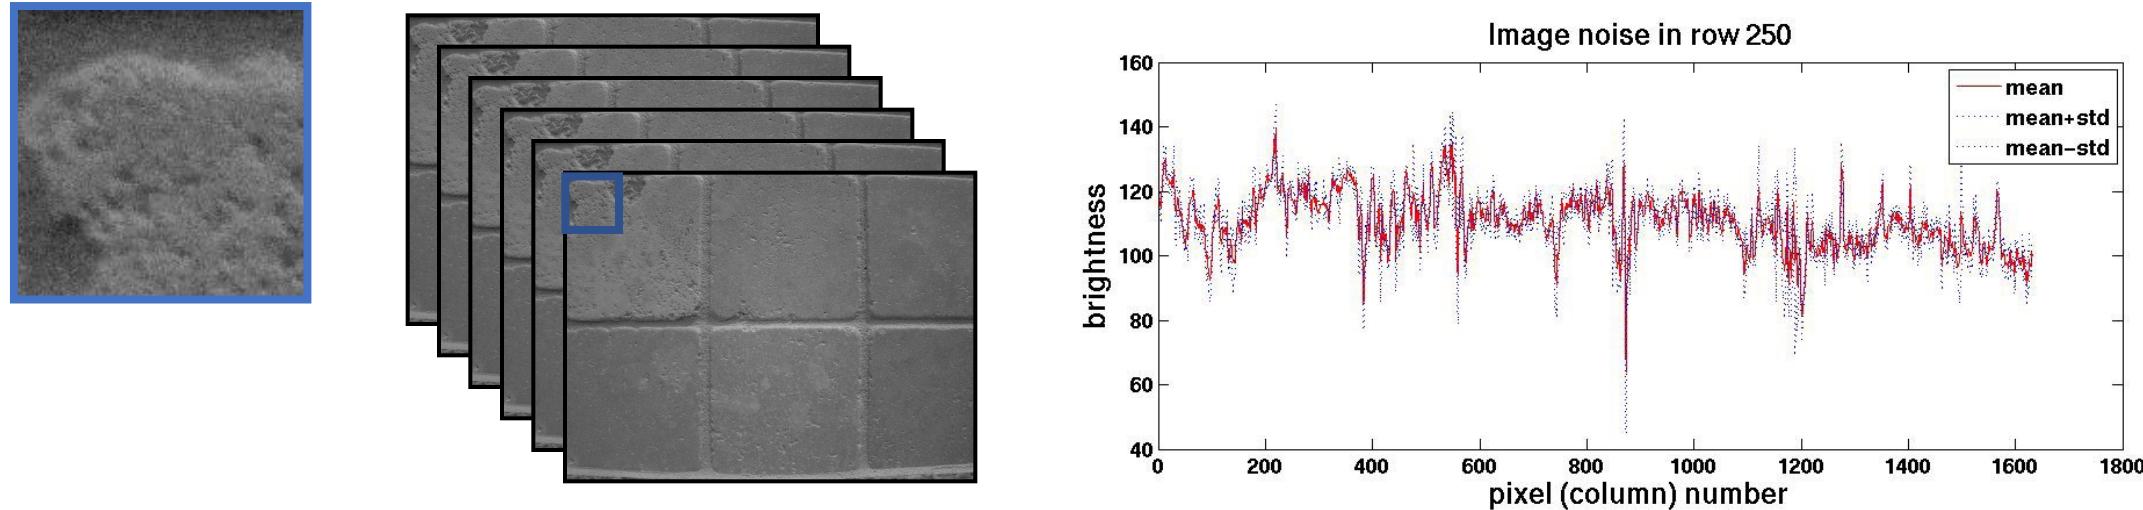
1967

- [j2] Thomas S. Huang, Mohamed Chikhaoui:
The effect of BSC on PCM picture quality. IEEE Trans. Inf. Theory 13(2): 270-273 (1967)

1965

- [j1] Thomas S. Huang:
The subjective effect of two-dimensional pictorial noise. IEEE Trans. Inf. Theory 11(1): 43-53 (1965)

Motivation: Image Noise Reduction



- Even multiple images of the **same static scene** will not be identical.

Common Types of Noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise

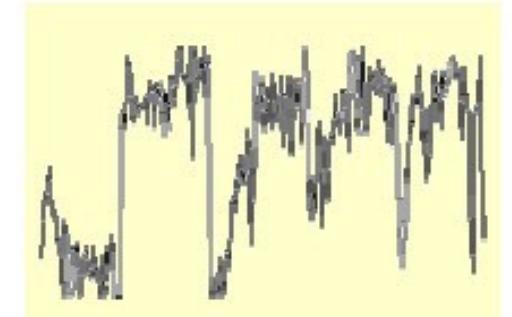
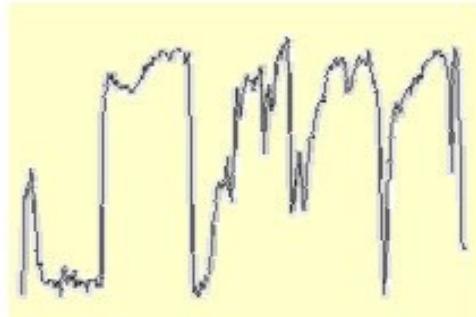
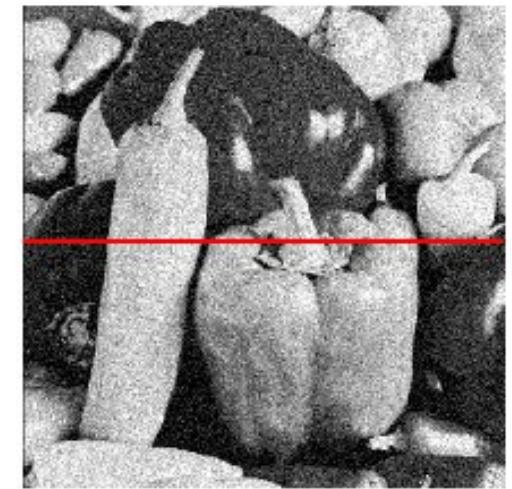


Gaussian noise

Gaussian noise

```
>> noise = randn(size(im)).*sigma;  
>> output = im + noise;
```

Image
Noise

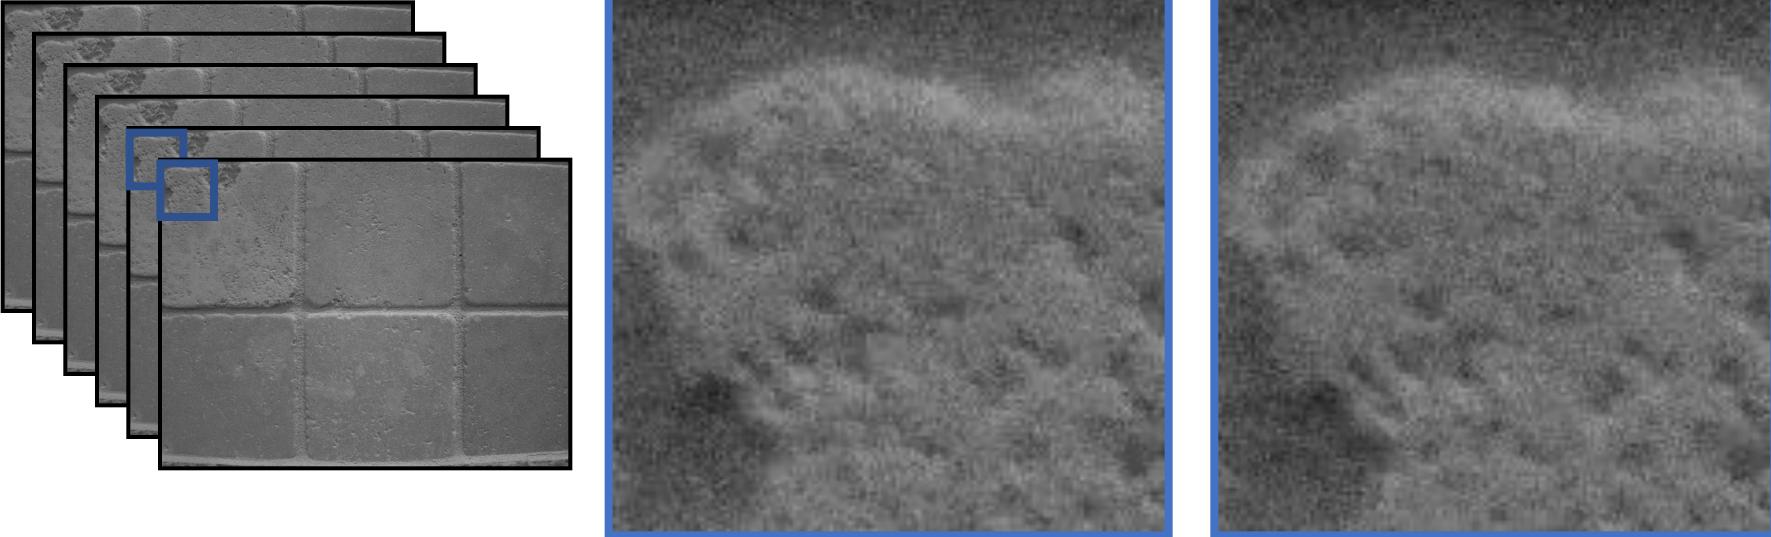


$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

What is impact of the sigma?

Motivation: Image Noise Reduction



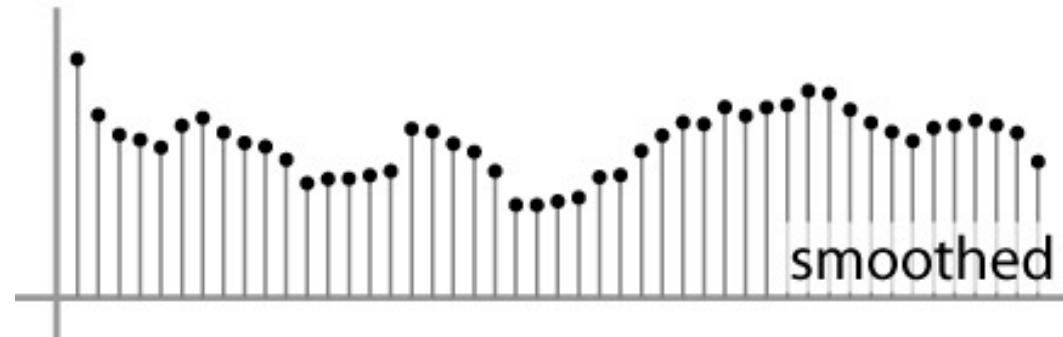
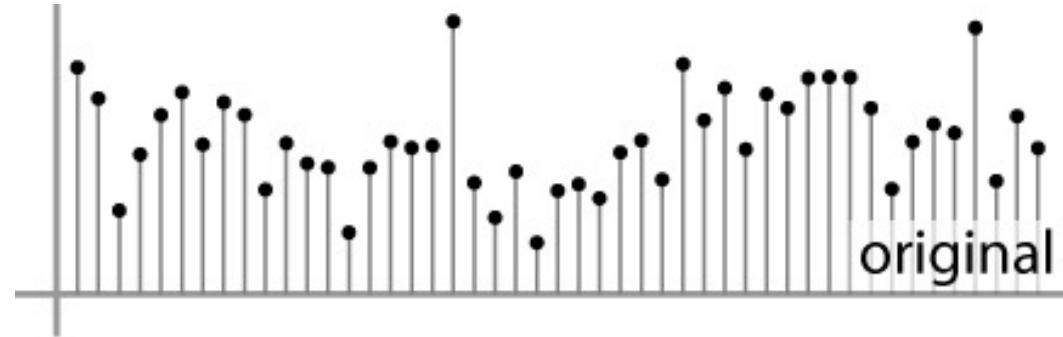
- Even multiple images of the same static scene will not be identical.
- How could we reduce the noise?
 - i.e., give an estimate of the true intensities
- **What if there's only one image?**

First Attempt at a Solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise processes to be independent from pixel to pixel

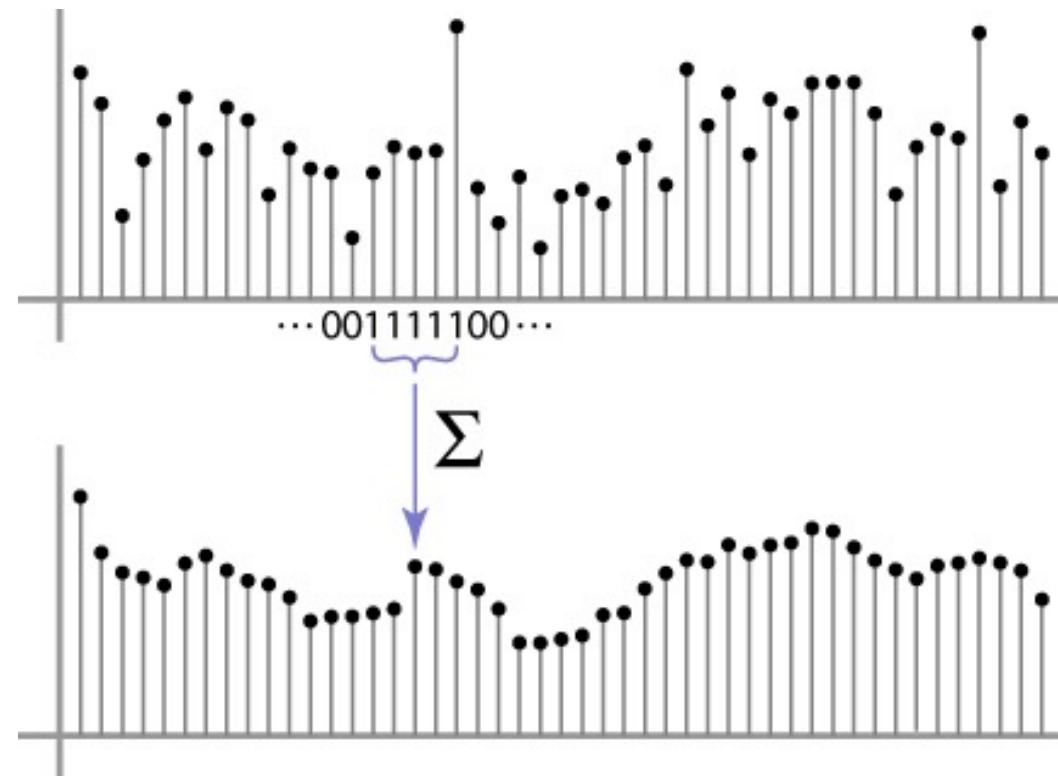
First Attempt at a Solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



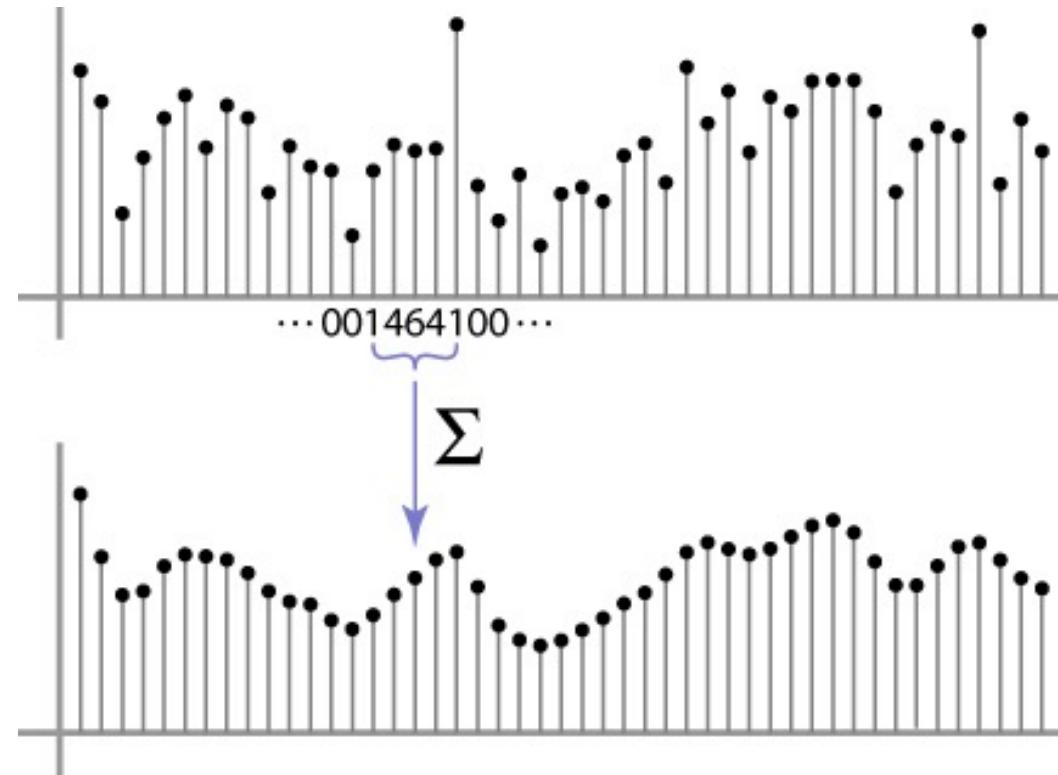
Weighted Moving Average

- Can add weights to our moving average
- *Weights* $[1, 1, 1, 1, 1] / 5$



Weighted Moving Average

- Non-uniform weights $[1, 4, 6, 4, 1] / 16$



Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0										

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10									

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

			0	10						

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

			0	10	20					

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving Average In 2D

$$F[x, y]$$

$$G[x, y]$$

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

			0	10	20	30	30			

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Correlation Filtering

Say the averaging window size is $2k+1 \times 2k+1$:

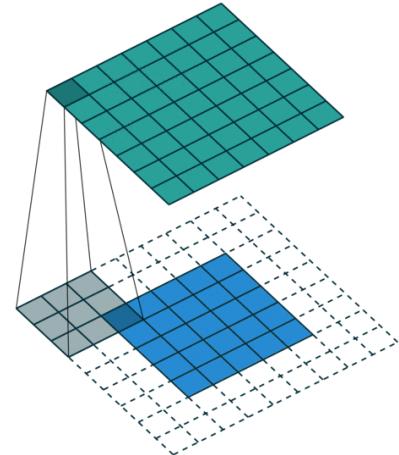
$$G[i, j] = \frac{1}{(2k+1)^2} \underbrace{\sum_{u=-k}^k}_{\text{Attribute uniform weight to each pixel}} \underbrace{\sum_{v=-k}^k}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]} F[i + u, j + v]$$

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] \underbrace{F[i + u, j + v]}_{\text{Non-uniform weights}}$$

Correlation Filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$



This is called cross-correlation, denoted $G = H \otimes F$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “kernel” or “mask” $H[u, v]$ is the prescription for the weights in the linear combination.

Averaging filter

- What values belong in the kernel H for the moving average example?

$$F[x, y] \otimes H[u, v] = G[x, y]$$

Diagram illustrating the convolution operation:

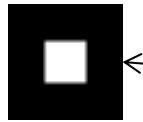
- Input Image ($F[x, y]$):** A 10x10 grid of grayscale values. A 3x3 neighborhood centered at position (4,4) is highlighted with a red border. The value 90 is present in the neighborhood.
- Kernel ($H[u, v]$):** A 3x3 grid of values. The bottom-right element is a question mark (?). The formula $\frac{1}{9}$ indicates that each input value is multiplied by $\frac{1}{9}$.
- Output Image ($G[x, y]$):** A 10x10 grid showing the result of the convolution. The value at position (4,4) is 30, indicated by a red box.

“box filter”

Talk to your neighbor

$$G = H \otimes F$$

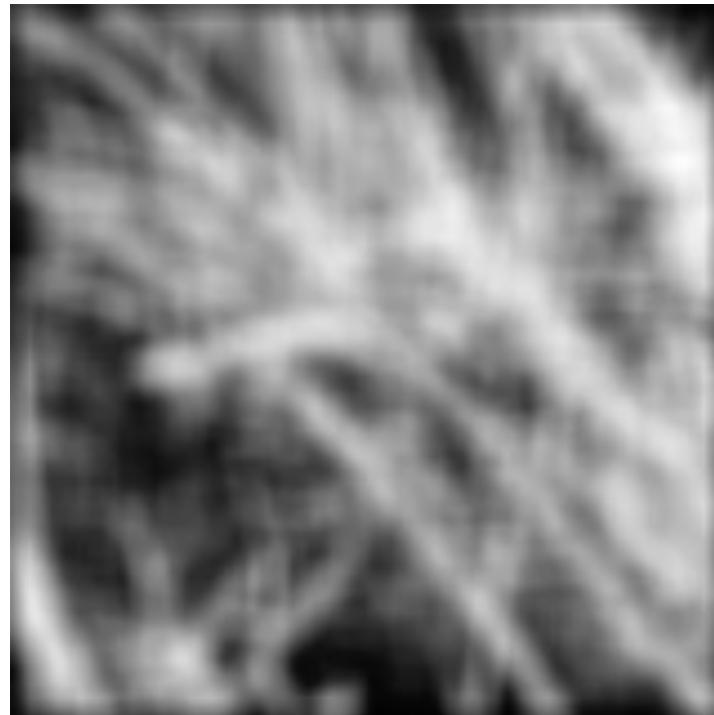
Smoothing by Averaging



depicts box filter:
white = high value, black = low value



original



filtered

What if the filter size was 5×5 instead of 3×3 ?

Gaussian Filter

- What if we want nearest neighboring pixels to have the most influence on the output?

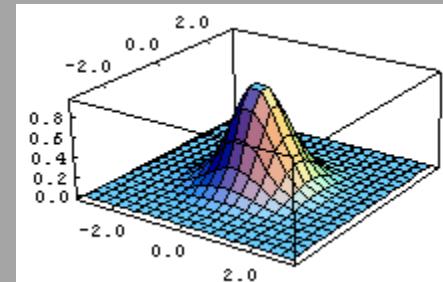
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$F[x, y]$$

$$\frac{1}{16} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$
$$H[u, v]$$

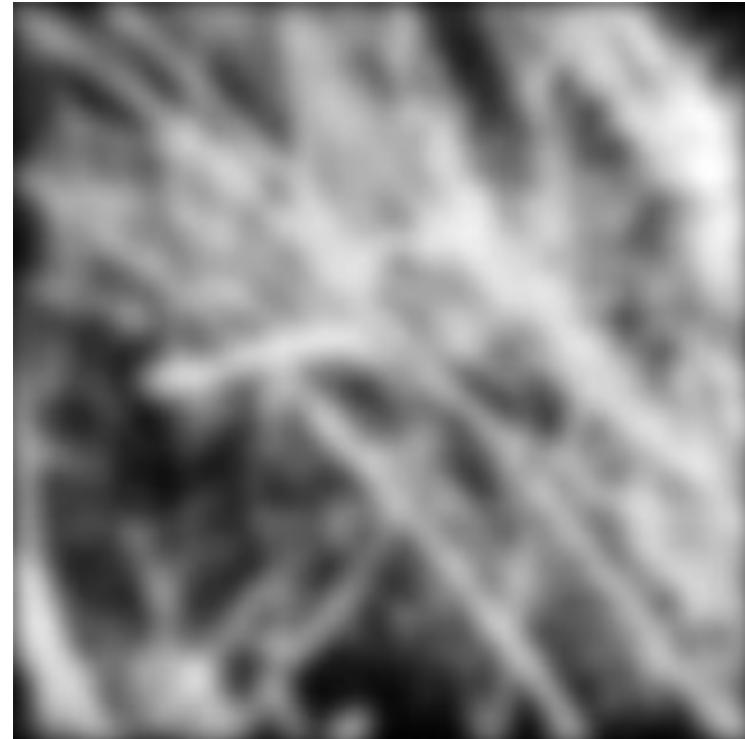
This kernel is an approximation of a 2d Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



- Removes high-frequency components from the image (“low-pass filter”).

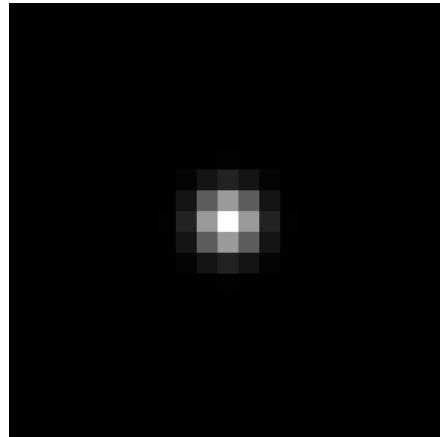
Smoothing with a Gaussian



Gaussian Filters

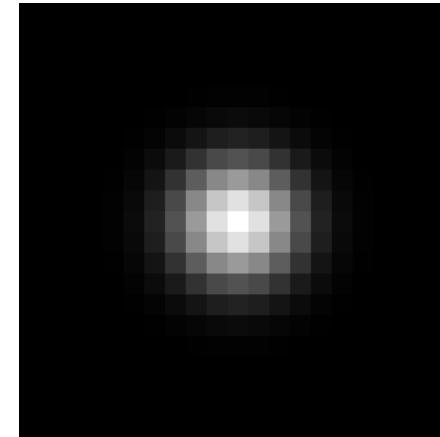
$\sigma = 1$

filter = 21x21



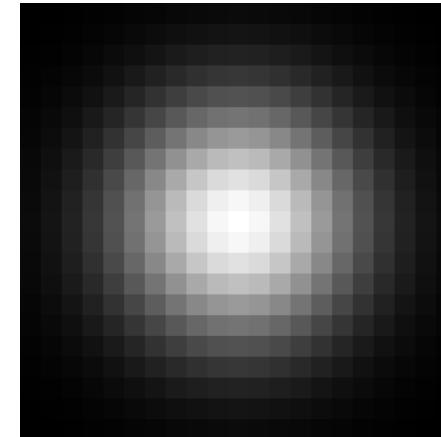
$\sigma = 2$

filter = 21x21



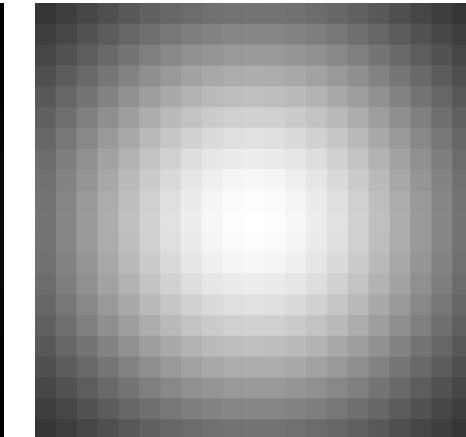
$\sigma = 4$

filter = 21x21



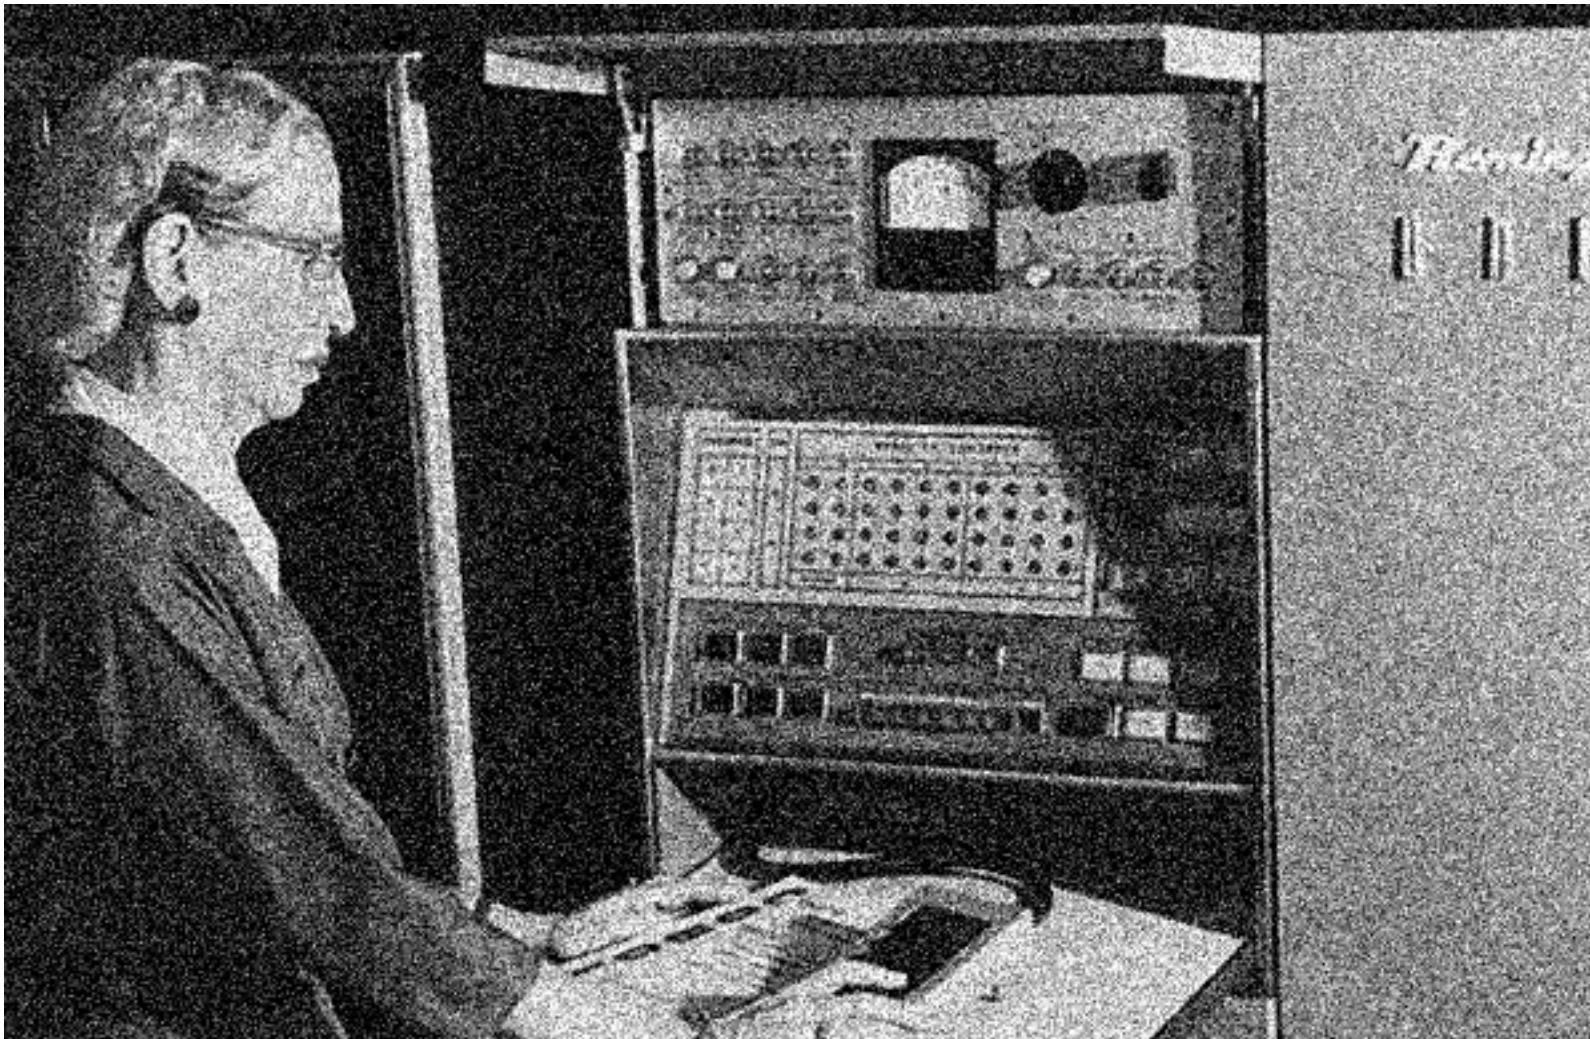
$\sigma = 8$

filter = 21x21



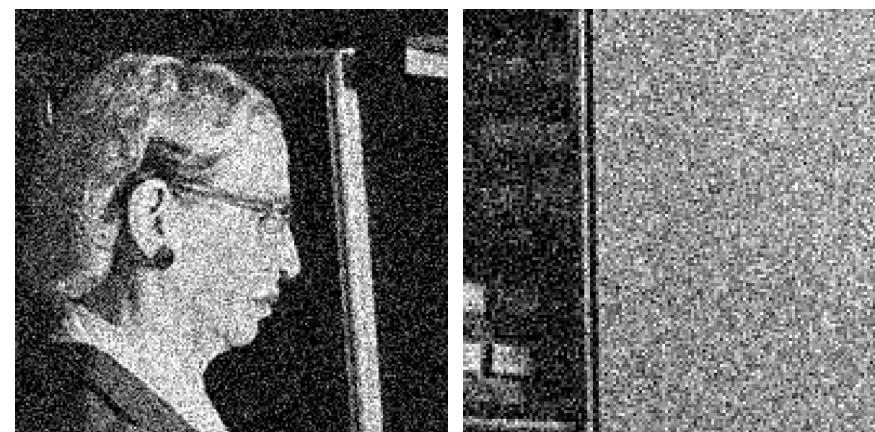
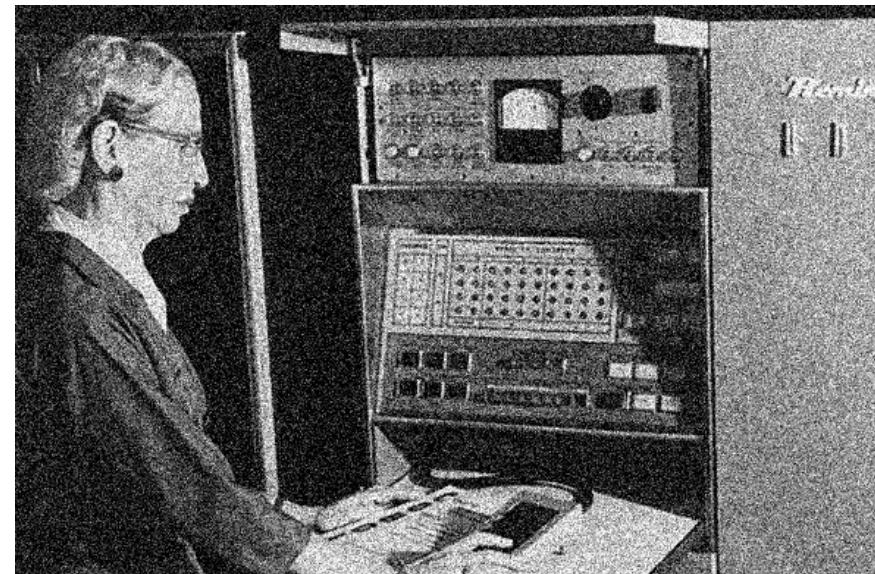
Note: filter visualizations are independently normalized throughout the slides so you can see them better

Applying Gaussian Filters



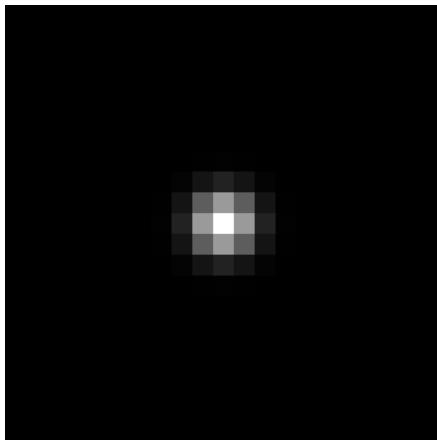
Applying Gaussian Filters

Input Image
(no filter)



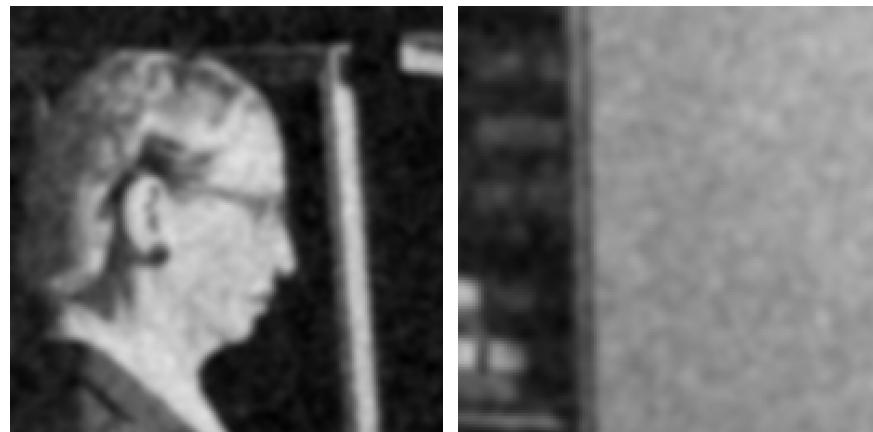
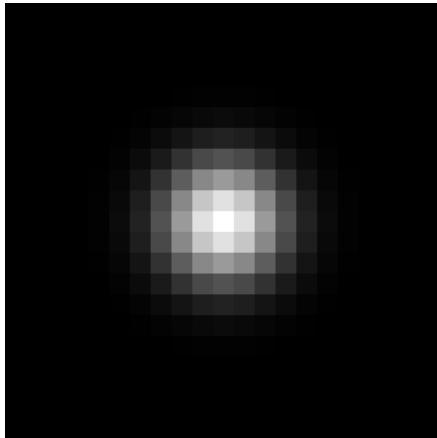
Applying Gaussian Filters

$$\sigma = 1$$



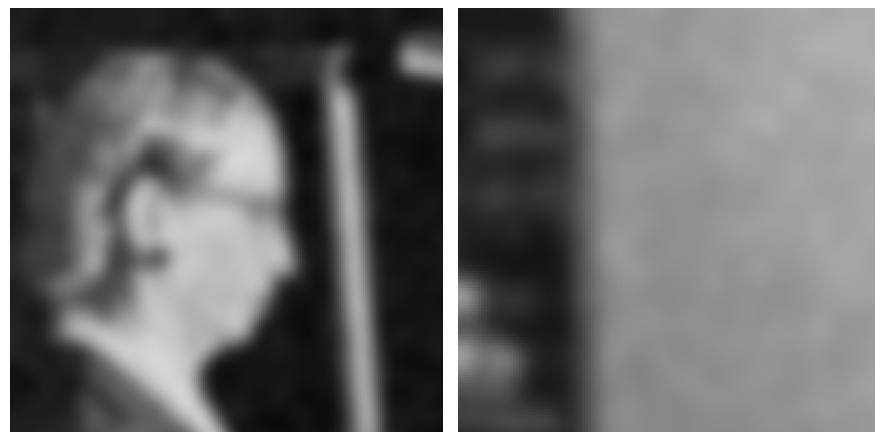
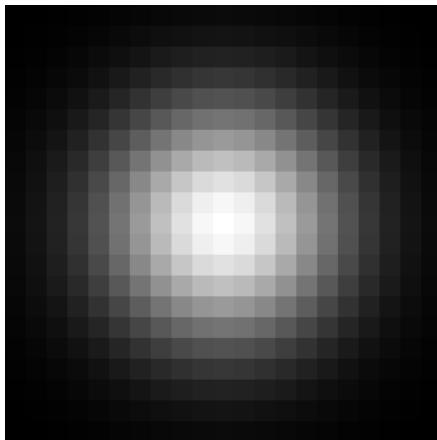
Applying Gaussian Filters

$$\sigma = 2$$



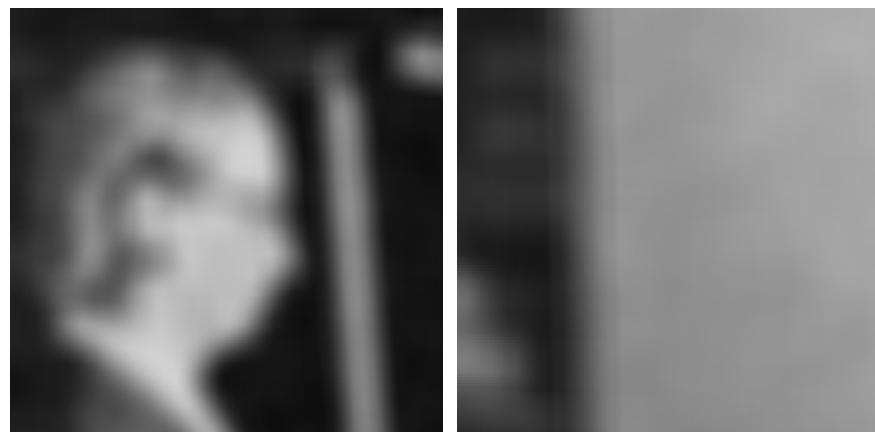
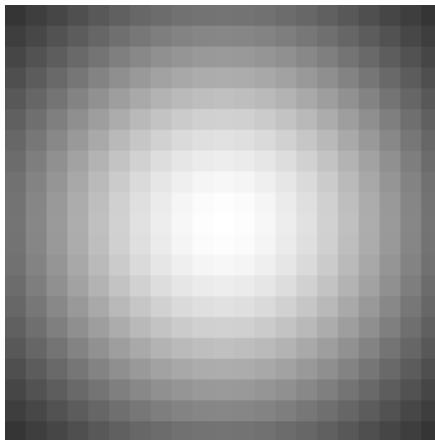
Applying Gaussian Filters

$$\sigma = 4$$



Applying Gaussian Filters

$$\sigma = 8$$

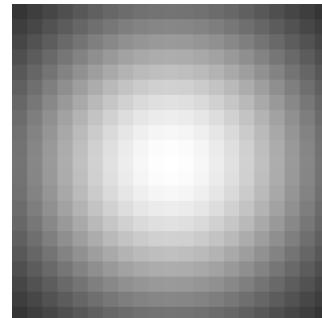


Picking a Filter Size

Too small filter → bad approximation

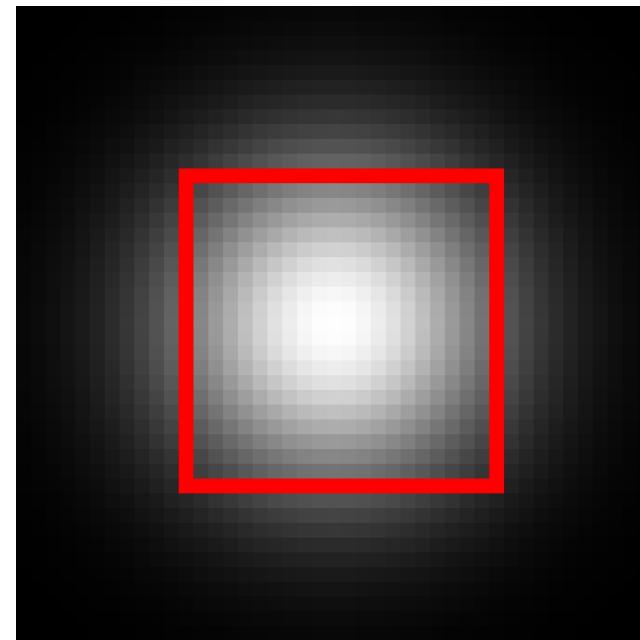
Approx Rule: Want filter size $\approx 6\sigma$ (99.7% of energy)

Far too small



$\sigma = 8$, size = 21

Slightly too small



$\sigma = 8$, size = 43

Runtime Complexity

Image size = $N \times N = 6 \times 6$

Filter size = $M \times M = 3 \times 3$

I11	I12	I13	I14	I15	I16
I21	F11	F12	F13	I25	I26
I31	F21	F22	F23	I35	I36
I41	F31	F32	F33	I45	I46
I51	I52	I53	I54	I55	I56
I61	I62	I63	I64	I65	I66

```
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterY in range(M):  
            for FilterX in range(M):  
                ...
```

Time: $O(N^2M^2)$

Separability

$\text{Conv}(\text{vector}, \text{transposed vector}) \rightarrow \text{outer product}$

$$\begin{array}{c|c} Fy1 \\ \hline Fy2 \\ \hline Fy3 \end{array} * \begin{array}{c|c|c} Fx1 & Fx2 & Fx3 \end{array} = \begin{array}{c|c|c} Fx1 * Fy1 & Fx2 * Fy1 & Fx3 * Fy1 \\ \hline Fx1 * Fy2 & Fx2 * Fy2 & Fx3 * Fy2 \\ \hline Fx1 * Fy3 & Fx2 * Fy3 & Fx3 * Fy3 \end{array}$$

Separability

$$Filter_{ij} \propto \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

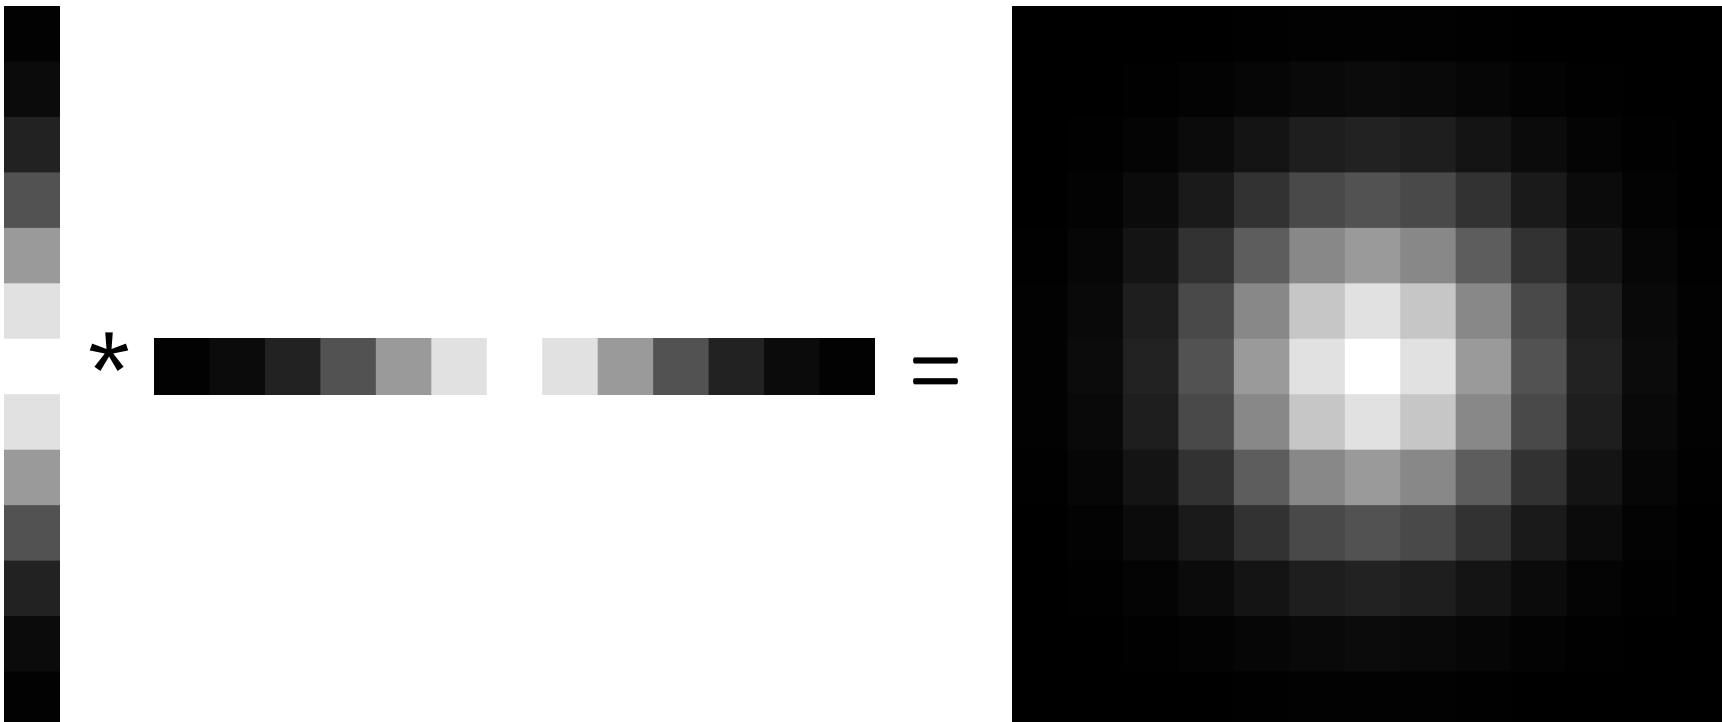


$$Filter_{ij} \propto \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)$$

Separability

$1D \text{ Gaussian} * 1D \text{ Gaussian} = 2D \text{ Gaussian}$

$\text{Image} * 2D \text{ Gauss} = \text{Image} * (1D \text{ Gauss} * 1D \text{ Gauss})$
 $= (\text{Image} * 1D \text{ Gauss}) * 1D \text{ Gauss}$



Runtime Complexity

Image size = $N \times N = 6 \times 6$

Filter size = $M \times 1 = 3 \times 1$

I ₁₁	I ₁₂	I ₁₃	I ₁₄	I ₁₅	I ₁₆
I ₂₁	F ₁	I ₂₃	I ₂₄	I ₂₅	I ₂₆
I ₃₁	F ₂	I ₃₃	I ₃₄	I ₃₅	I ₃₆
I ₄₁	F ₃	I ₄₃	I ₄₄	I ₄₅	I ₄₆
I ₅₁	I ₅₂	I ₅₃	I ₅₄	I ₅₅	I ₅₆
I ₆₁	I ₆₂	I ₆₃	I ₆₄	I ₆₅	I ₆₆

```
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterY in range(M):  
            ...  
    for ImageY in range(N):  
        for ImageX in range(N):  
            for FilterX in range(M):  
                ...
```

Time: $O(N^2M)$

Properties of smoothing filters

- Smoothing
 - Values positive
 - Sum to 1 → constant regions same as input
 - Amount of smoothing proportional to mask size
 - Remove “high-frequency” components; “low-pass” filter

Filtering an impulse signal

What is the result of filtering the impulse signal (image) F with the arbitrary kernel H ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G[x, y]$$

Talk with your neighbor

Filtering an impulse signal

What is the result of filtering the impulse signal (image) F with the arbitrary kernel H ?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

$F[x, y]$



a	b	c
d	e	f
g	h	i

$H[u, v]$

		i	h	g			
		f	e	d			
		c	b	a			

$G[x, y]$

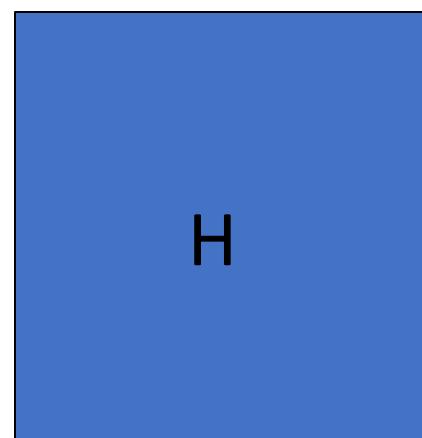
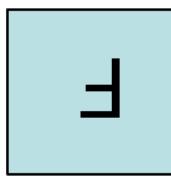
Convolution

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

↑
Notation for convolution operator



Convolution vs. Correlation

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

Convolution vs. correlation

	Padded f	Initial position for w	Correlation result	Full correlation result
Origin f	0 0 0 0 0 0 0 0 0 0 0 0 0 0	1 2 3 4 5 6 7 8 9 0 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 9 8 7 0 0 6 5 4 0 0 3 2 1 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 9 8 7 0 0 0 0 6 5 4 0 0 0 0 3 2 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
(a)	(b)	(c)	(d)	(e)

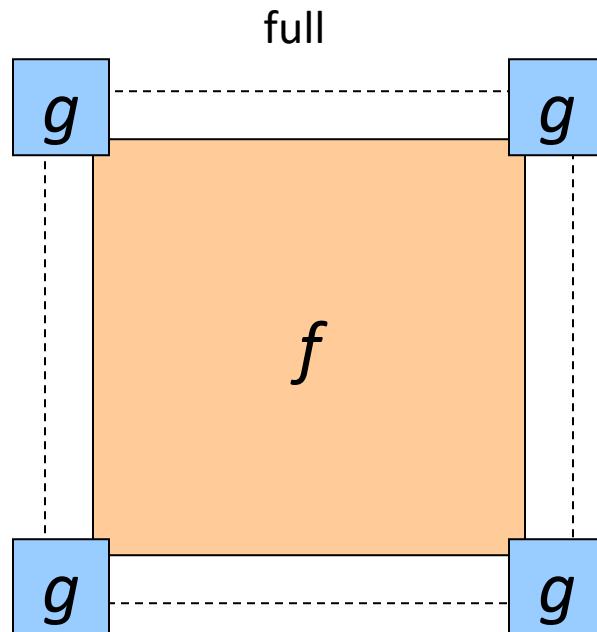
	Rotated w	Convolution result	Full convolution result
	9 8 7 6 5 4 3 2 1 0 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 1 2 3 0 0 4 5 6 0 0 7 8 9 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 2 3 0 0 0 0 4 5 6 0 0 0 0 7 8 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
	(f)	(g)	(h)

Boundary Issues

Convolution doesn't keep the whole image.

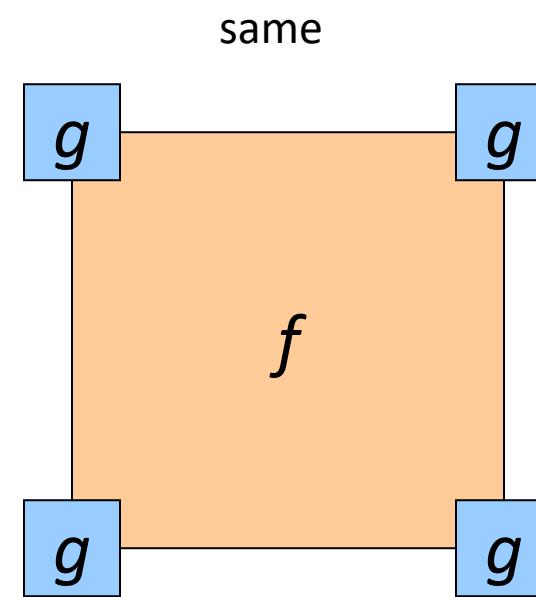
Suppose **f** is the image and **g** the filter.

Full – any part of g touches f.

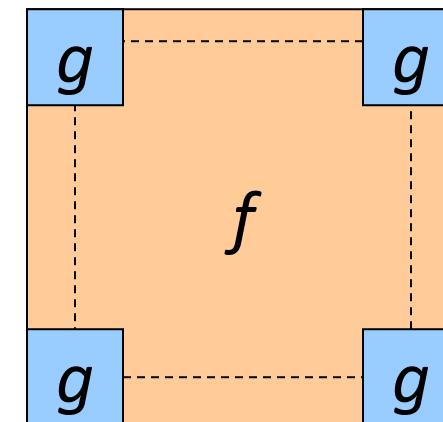


Same – same size as f;

Valid – only when filter doesn't fall off edge.

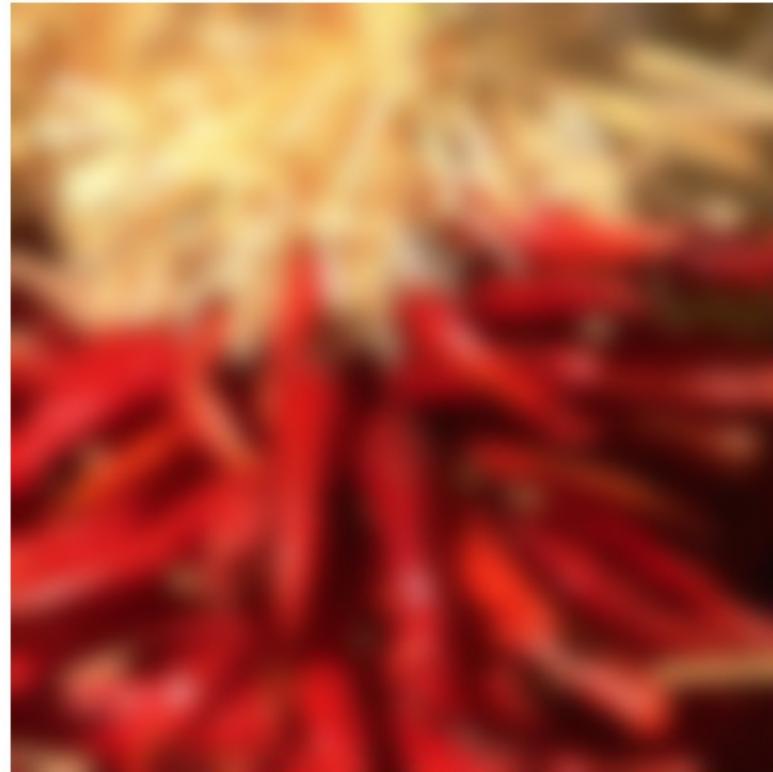


valid



Boundary issues

- What about near the edge?
 - the filter window falls off the edge of the image
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Properties of Convolution: Linear

Given: Image (I), Filters (f_1, f_2)

Linear: $\text{apply}(I, f_1 + f_2) = \text{apply}(I, f_1) + \text{apply}(I, f_2)$

I is a white box on black, and f_1, f_2 are rectangles

$$A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array} + \begin{array}{|c|}\hline \square \\ \hline\end{array}) = A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array}) = \begin{array}{|c|}\hline \bullet \\ \hline\end{array}$$

$$A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array}) + A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array}) = \begin{array}{|c|}\hline \bullet \\ \hline\end{array} + \begin{array}{|c|}\hline \bullet \\ \hline\end{array} = \begin{array}{|c|}\hline \bullet \\ \hline\end{array}$$

Note: I am showing filters un-normalized and blown up. They're a smaller box filter (i.e., each entry is $1/(\text{size}^2)$)

Properties of Convolution: Shift-Invariant

Given: Image (I), Filter (f)

Shift-invariant: $\text{shift}(\text{apply}(I, f)) = \text{apply}(\text{shift}(I, f))$

Intuitively: only depends on filter neighborhood

$$A(\begin{array}{|c|} \hline \blacksquare \\ \hline \end{array}, \boxed{\quad}) = \begin{array}{|c|} \hline \text{gray square} \\ \hline \end{array}$$

$$A(\begin{array}{|c|} \hline \quad \\ \hline \blacksquare \\ \hline \end{array}, \boxed{\quad}) = \begin{array}{|c|} \hline \quad \\ \hline \text{gray square} \\ \hline \end{array}$$

Properties of Convolution

- Any shift-invariant, linear operation is a convolution ($*$)
- Commutative: $f * g = g * f$
- Associative: $(f * g) * h = f * (g * h)$
- Distributes over $+$: $f * (g + h) = f * g + f * h$
- Scalars factor out: $k f * g = f * k g = k (f * g)$
- Identity (a single one with all zeros):

