

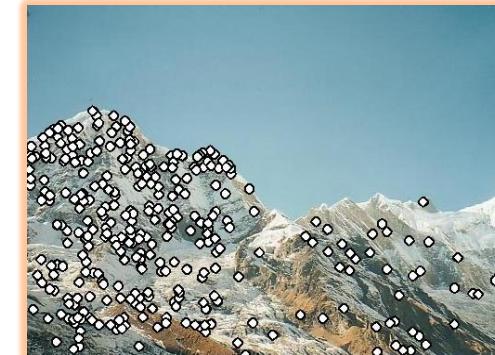
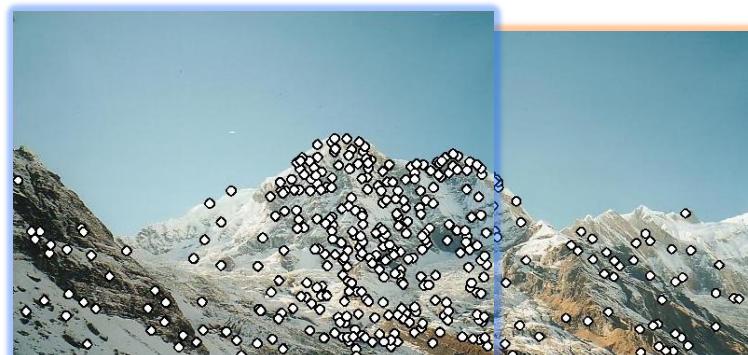


# CS 4476 Computer Vision:

## 5. Local Feature Detection

Humphrey Shi

9/3/2024



# Outline

- Logistics & Course Recap
- Local Feature Detection

# Logistics & Course Recap

# Logistics

- Assignment 1 – Due this Sunday.
- Estimated weekly effort ~ 10-15hrs.



# Logistics

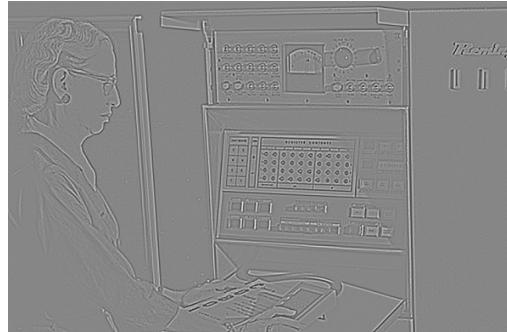
- Schedule for TA office hours
  - Monday 10 am to 12 pm: Anh, Fengzhe
  - Tuesday 9 am to 11 am: Dhruv,
  - Wednesday 3pm to 5pm: Chieh, Yue
  - Thursday 9am to 11am: Aditya, Shubham
  - Friday 9am to 11 am: Manushree, Bogi

# Recap: Image Filtering (continued)



Image

$+ \alpha$

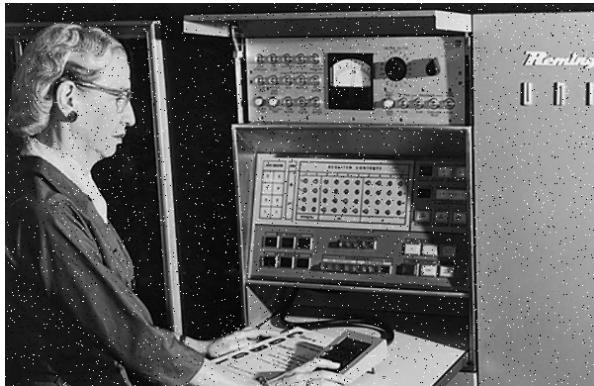


Details

=



"Sharpened"  $\alpha=2$



Noisy image



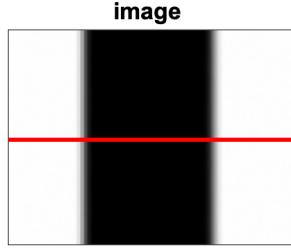
Gaussian filter,  $\sigma = 1$



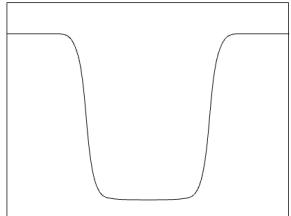
Median filter, size = 3



# Recap: Gradients & Edges

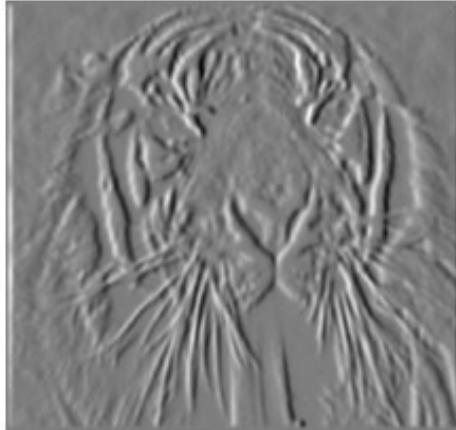


intensity function  
(along horizontal scanline)



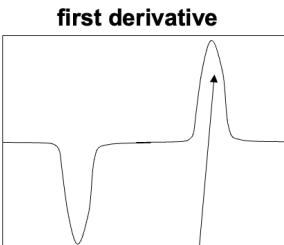
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

1	
-1	



edges correspond to  
extrema of derivative

$$M_x = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$



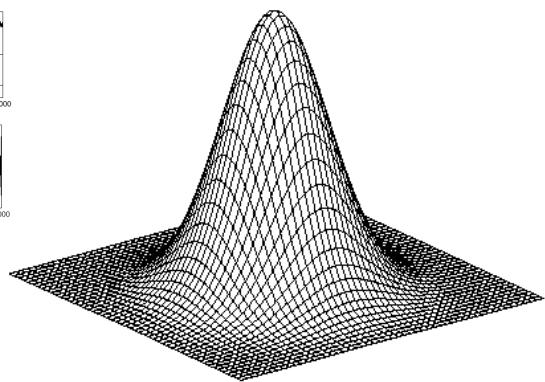
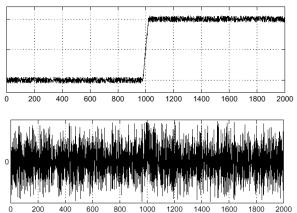
$$M_y =$$

-1	-2	-1
0	0	0
1	2	1

Sobel Filter

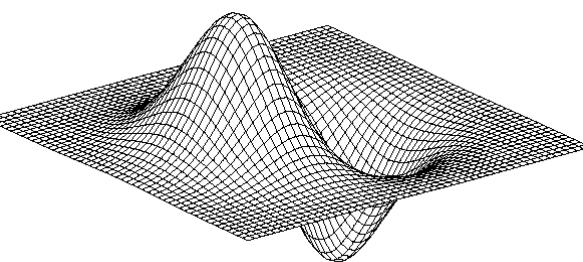
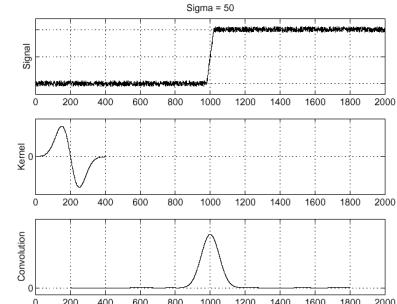
# Recap: Edge Detection

- 2D Edge Detection Filters



**Gaussian**

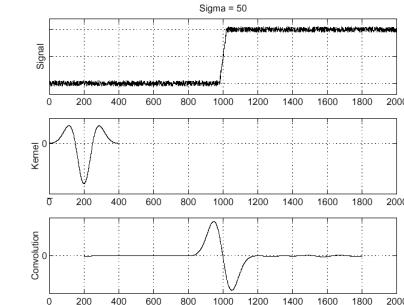
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



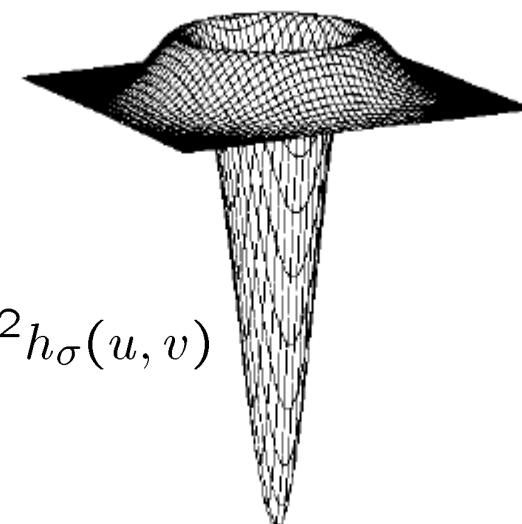
**derivative of Gaussian**

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$$\nabla^2 h_\sigma(u, v)$$



**Laplacian of Gaussian**



- Canny Edge Detector
  - How?

GT



# Noise Reduction by Smoothing



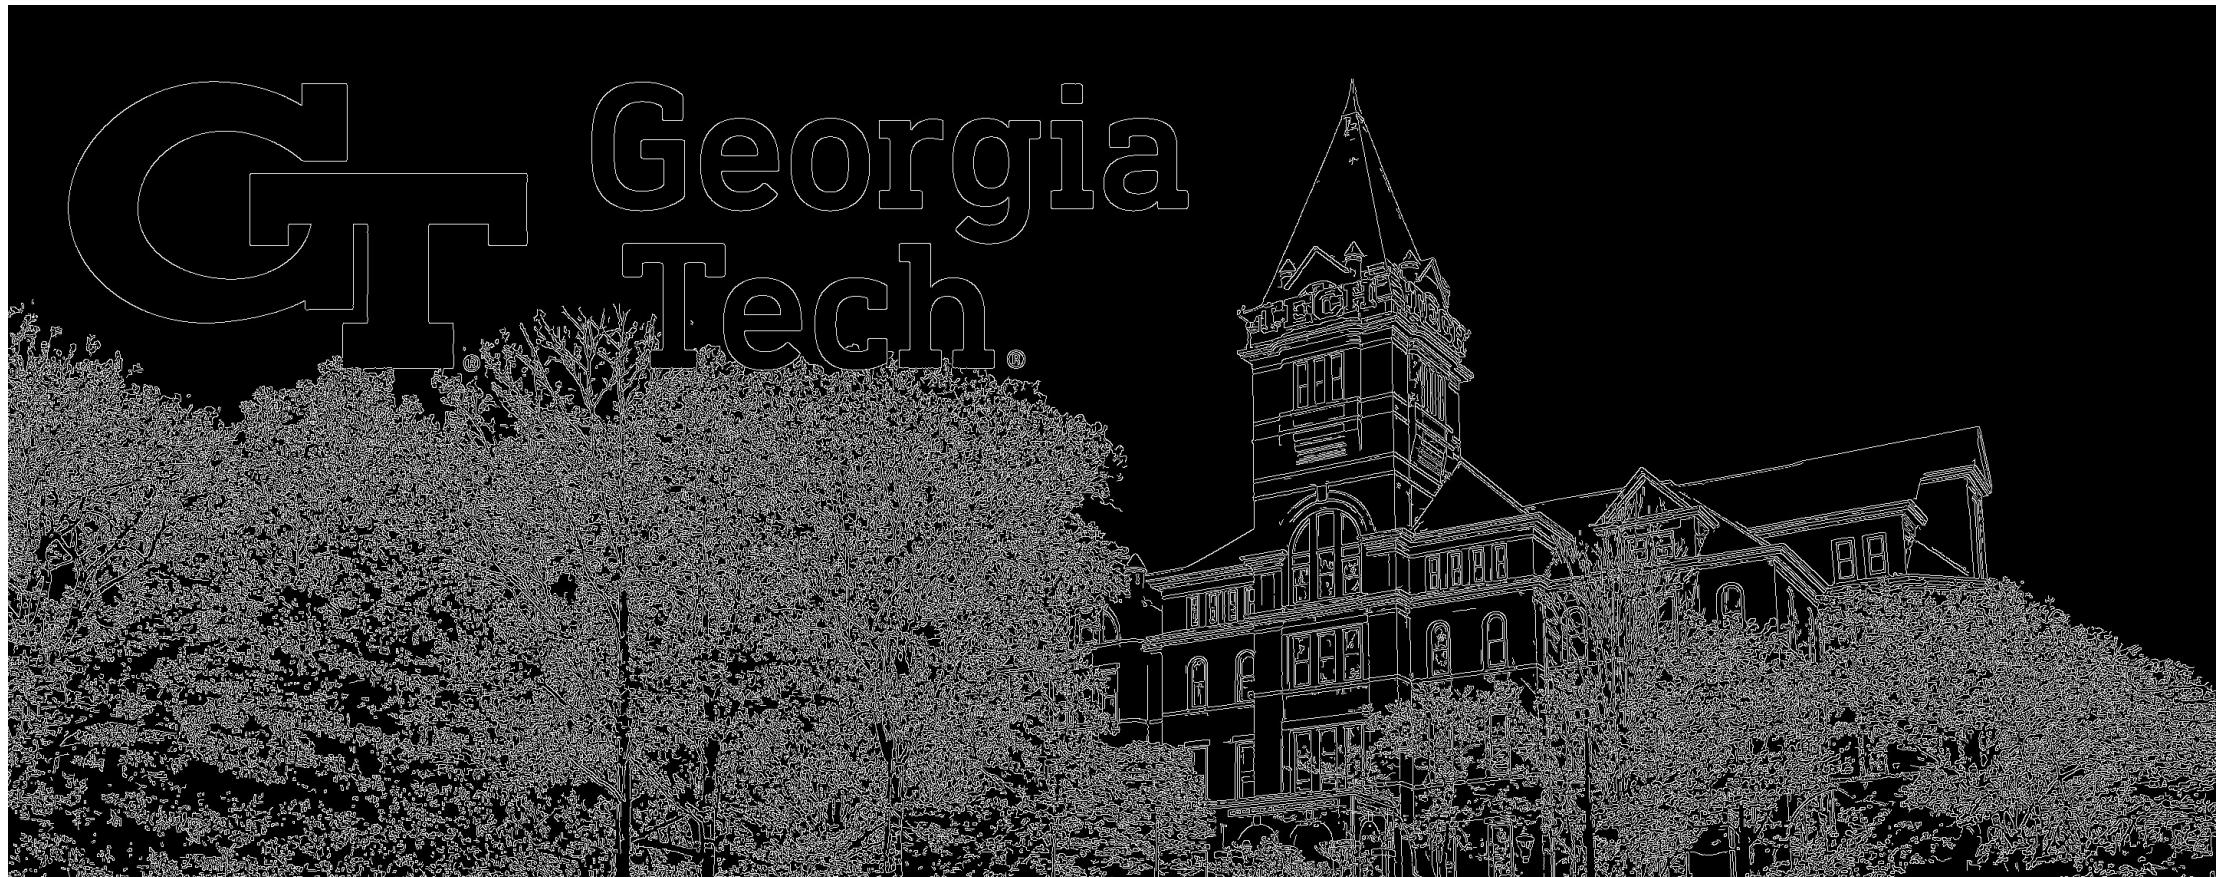
# Finding Image Gradients



# Non-Maximum Suppression to Thin Edges



# Edge Tracking by Hysteresis Thresholding



GT



# Local Feature Detection

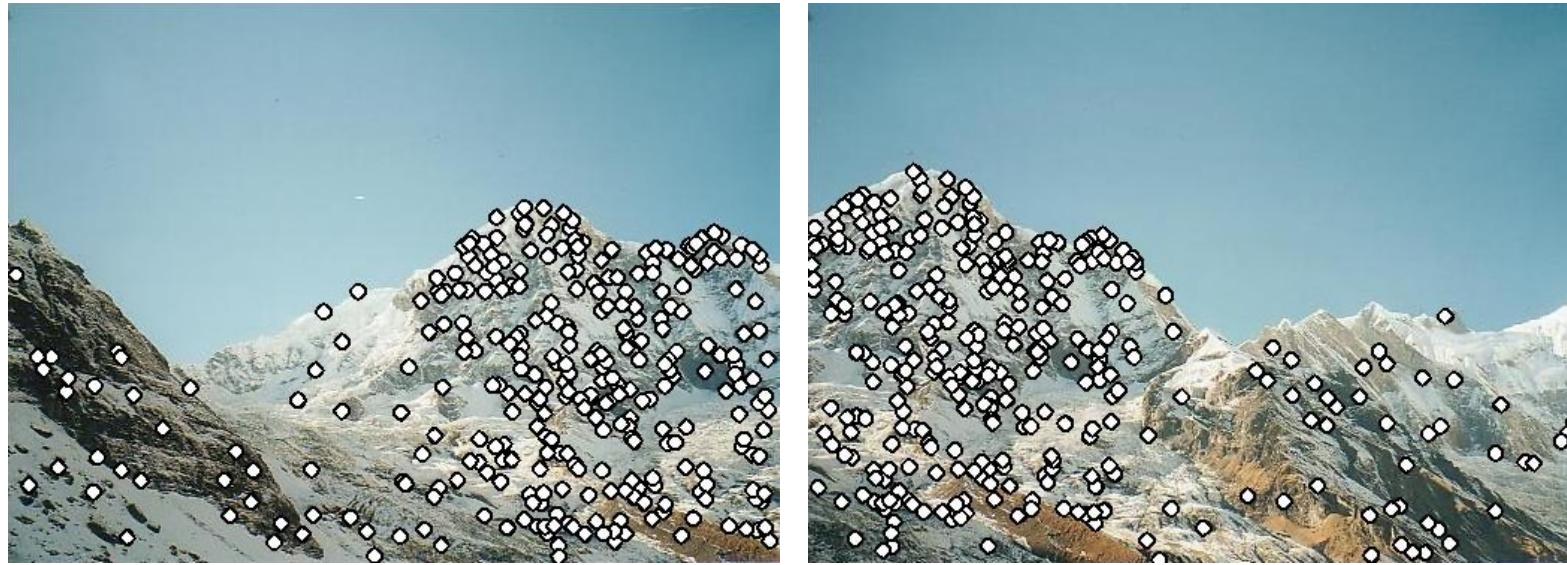
Image Alignment & Local Features

# Image Alignment



Question: How do we combine the two images?

# Robust feature-based alignment



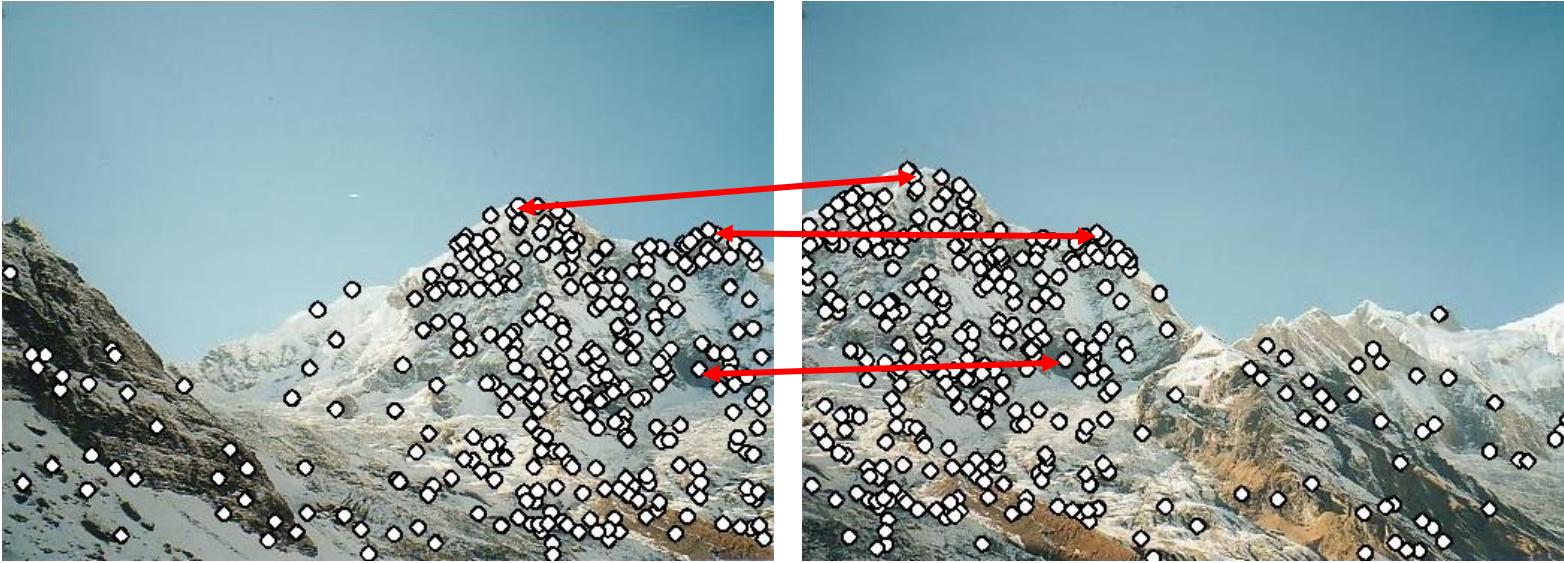
- Extract **keypoint** features

# Robust feature-based alignment



- Extract keypoint features
- Compute *potential matches*

# Robust feature-based alignment

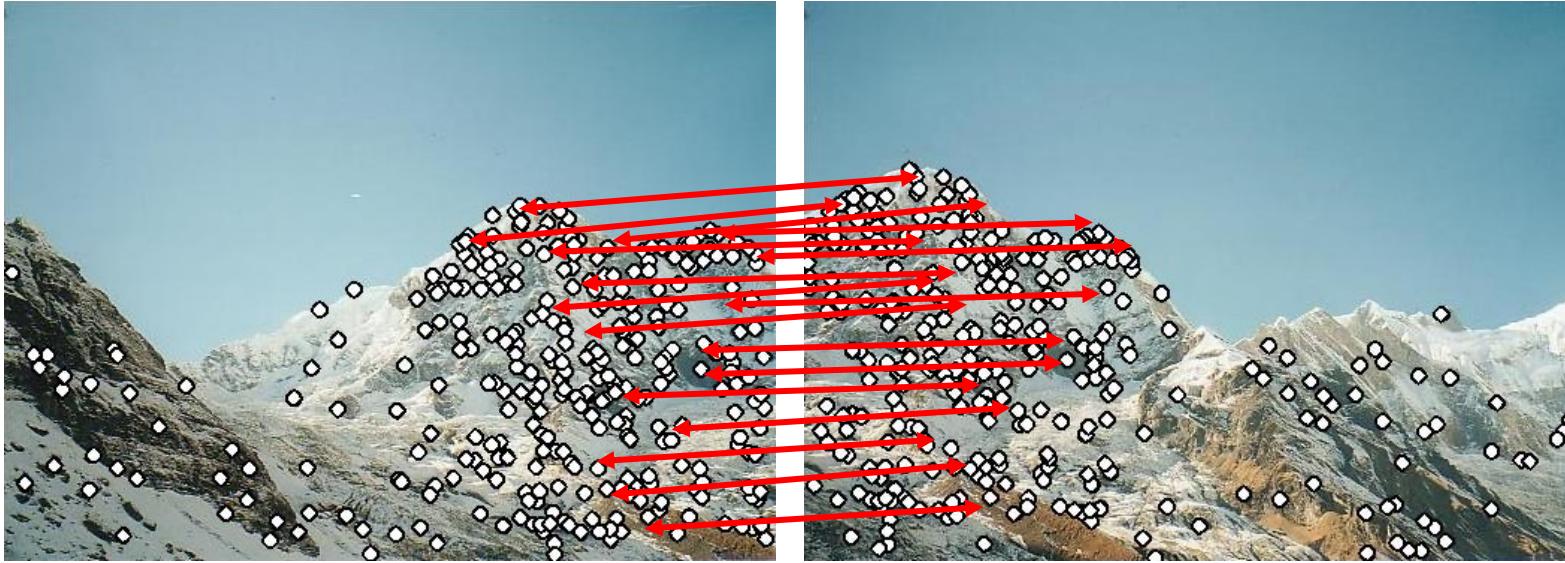


- Extract keypoint features
- Compute *potential matches*

**Loop:**

*Hypothesize* transformation  $T$  (small group of potential matches that are related by  $T$ )

# Robust feature-based alignment



- Extract keypoint features
- Compute *potential matches*

**Loop:**

*Hypothesize* transformation  $T$  (small group of potential matches that are related by  $T$ )

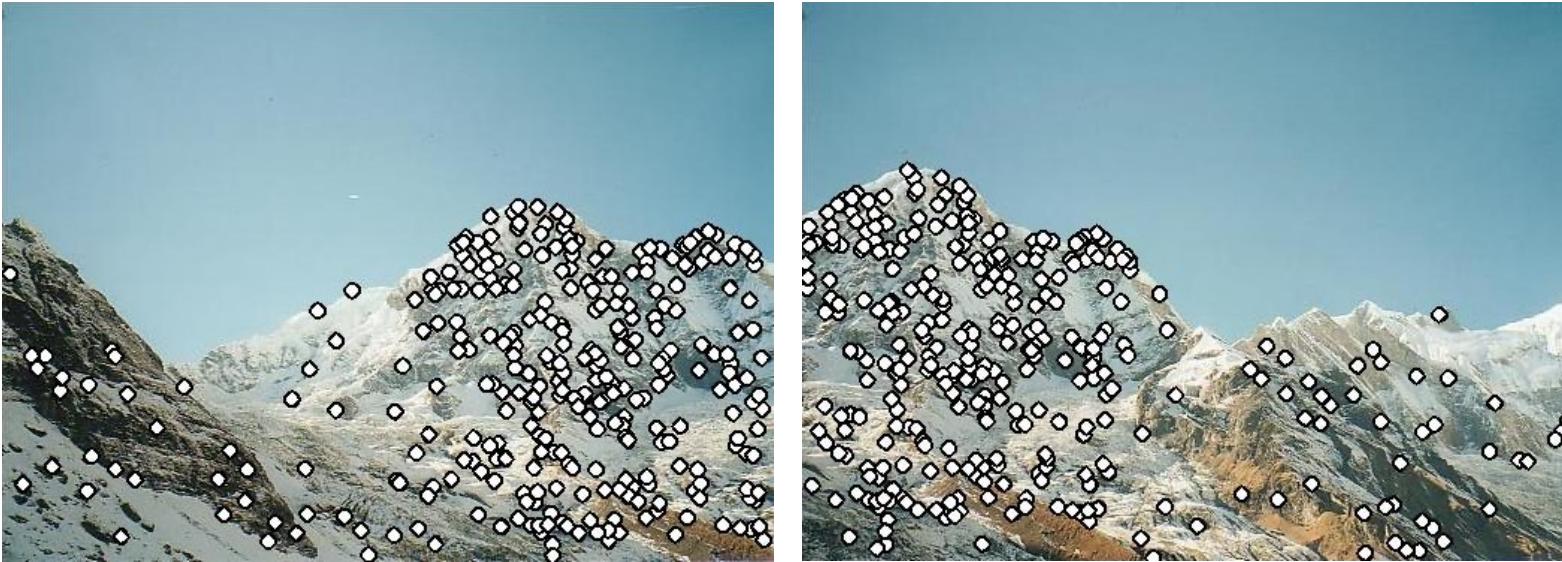
*Verify* transformation (search for other matches consistent with  $T$ )

# Robust feature-based alignment



- Extract keypoint features
- Compute *potential matches*
- Align the images

# Today: Detecting Features



How to detect *which features* to match?

# Local features: desired properties

- Repeatability
  - The same feature can be found in several images despite geometric and photometric transformations
- Saliency
  - Each feature has a distinctive description
- Compactness and efficiency
  - Many fewer features than image pixels
- Locality
  - A feature occupies a relatively small area of the image; robust to clutter and occlusion

# Goal: interest operator repeatability

- We want to detect (at least some of) the same points in both images.

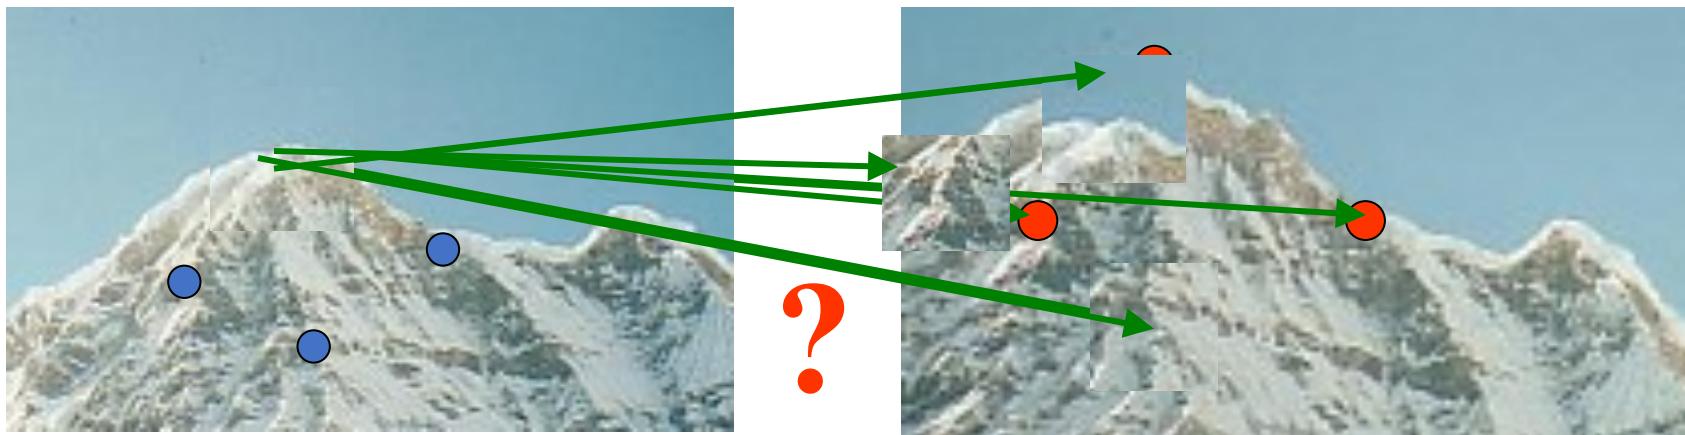


No chance to find true matches!

- Yet we have to be able to run the detection procedure *independently* per image.

# Goal: descriptor distinctiveness

- We want to be able to reliably determine which point goes with which.



- Must provide some invariance to geometric and photometric differences between the two views.



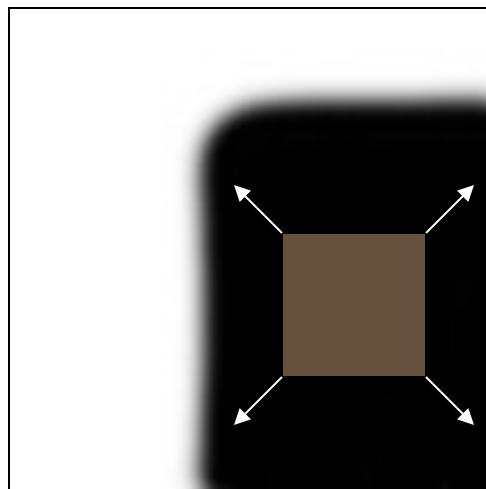
- Question: What keypoints should we choose?

# Local Feature Detection

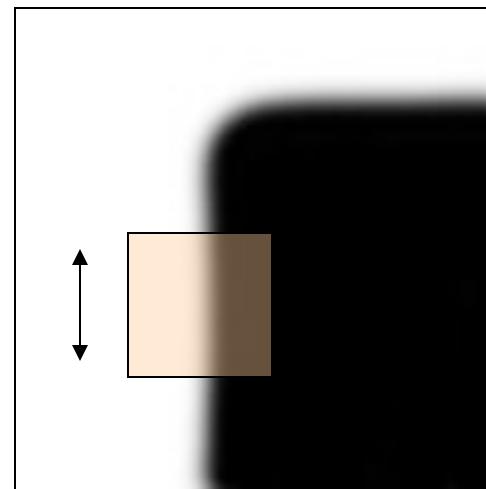
Corners as Interest Points

# Corners as distinctive interest points

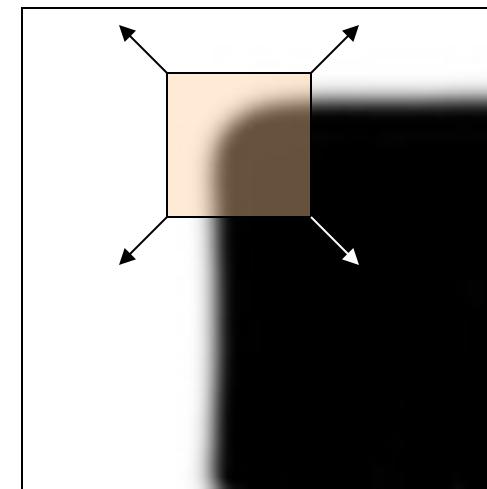
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity



“flat” region:  
no change in  
all directions

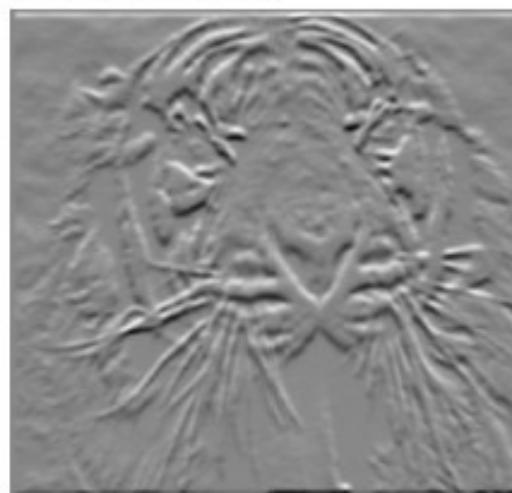
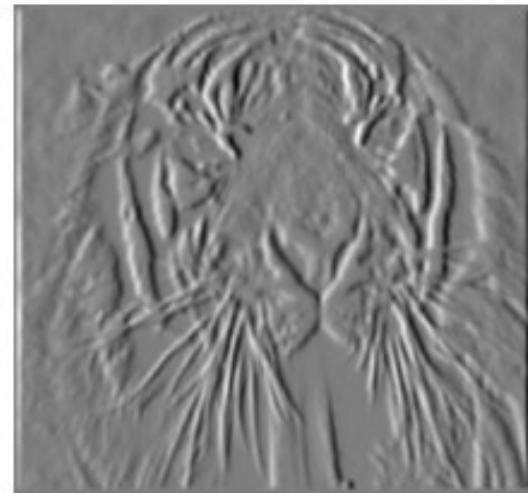


“edge”:  
no change  
along the edge  
direction



“corner”:  
significant  
change in all  
directions

Recall: Partial derivatives measure changes along direction



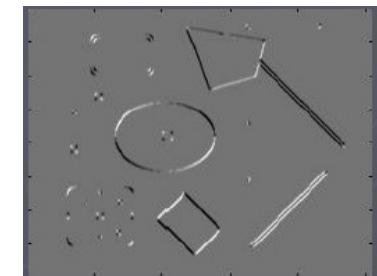
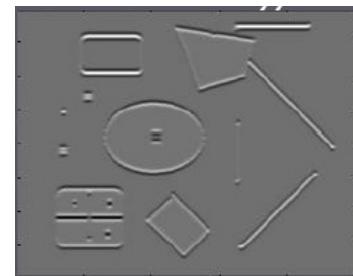
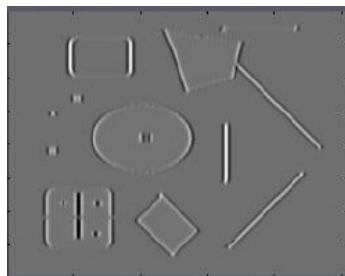
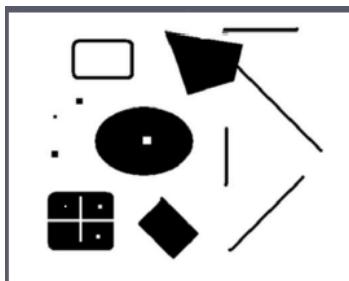
$$I_x = \frac{dI}{dx} \\ = \frac{\partial f(x, y)}{\partial x}$$

$$I_y = \frac{dI}{dy} \\ = \frac{\partial f(x, y)}{\partial y}$$

# Corners as distinctive interest points

$$M = \sum_{(i,j) \in W} \begin{bmatrix} I_x(i,j)I_x(i,j) & I_x(i,j)I_y(i,j) \\ I_x(i,j)I_y(i,j) & I_y(i,j)I_y(i,j) \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point  $(i,j)$  ).



Notation:

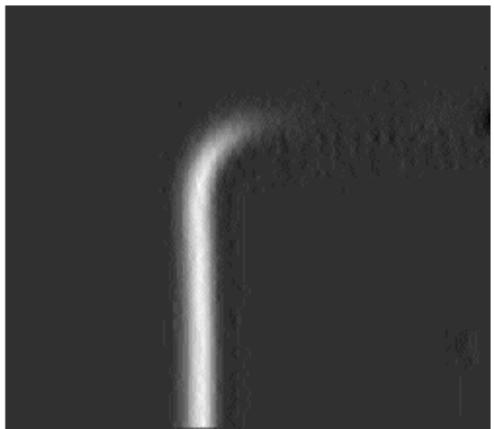
$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

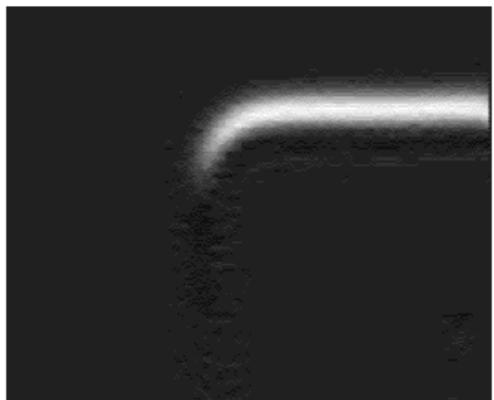
$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# What does this matrix reveal?

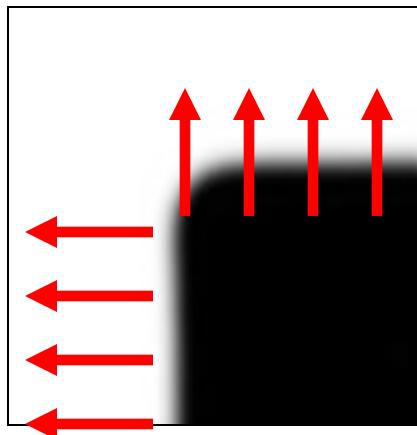
First, consider an axis-aligned corner:



$I_x$

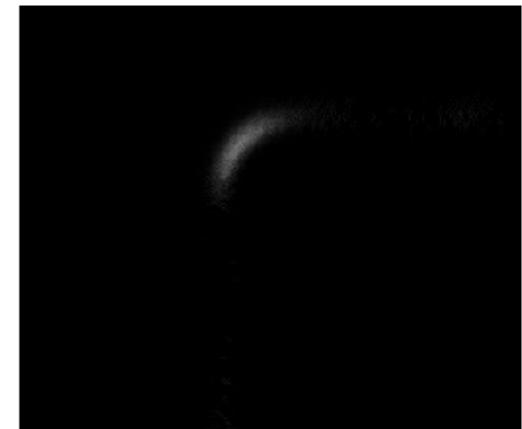


$I_y$



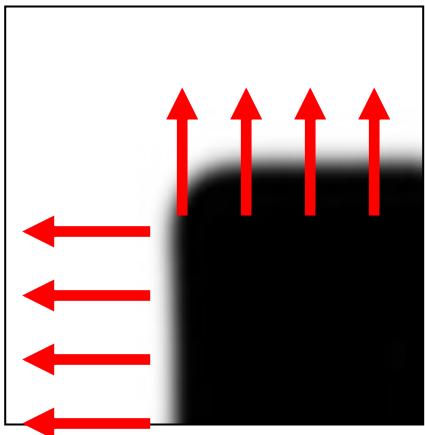
$$M = \sum_{(i,j) \in W} \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

What will  $I_x I_y$  look like?



# What does this matrix reveal?

First, consider an axis-aligned corner:



$$M = \sum_{(i,j) \in W} \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix} \approx \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

This means dominant gradient directions align with x or y axis

Look for locations where **both**  $\lambda$ 's are large.

If either  $\lambda$  is close to 0, then this is **not** corner-like.

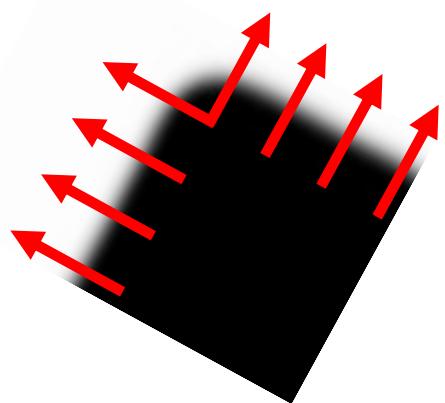
What if we have a corner that is not aligned with the image axes?

# What does this matrix reveal?

Since  $M$  is symmetric, we have

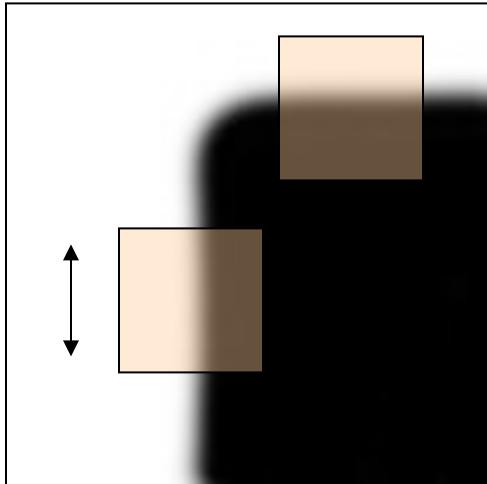
$$M = X \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} X^T$$

$$Mx_i = \lambda_i x_i$$



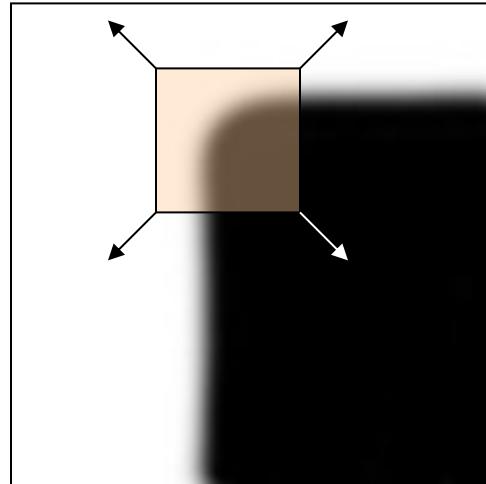
The *eigenvalues* of  $M$  reveal the amount of intensity change in the two principal orthogonal gradient directions in the window.

# Corner response function



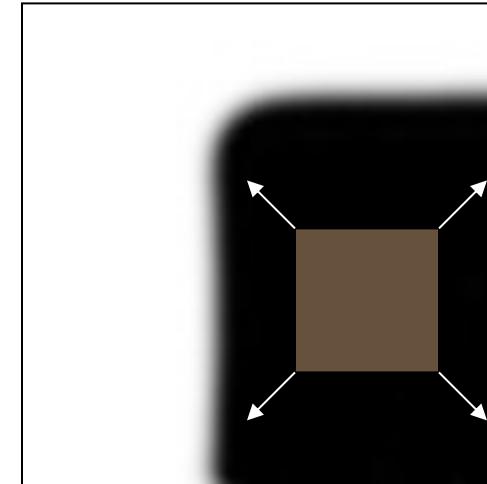
“edge”:

$$\begin{aligned}\lambda_1 &>> \lambda_2 \\ \lambda_2 &>> \lambda_1\end{aligned}$$



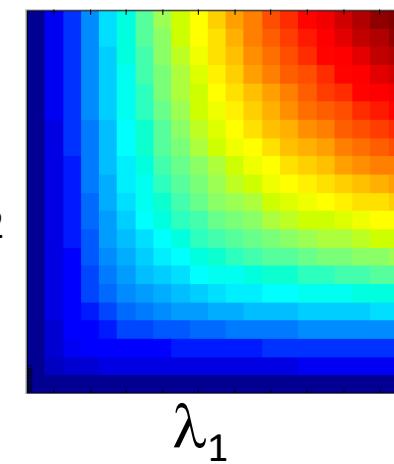
“corner”:

$\lambda_1$  and  $\lambda_2$  are large,  
 $\lambda_1 \sim \lambda_2$ ;



“flat” region

$\lambda_1$  and  $\lambda_2$  are small;



$$f = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2}$$

# Harris corner detector

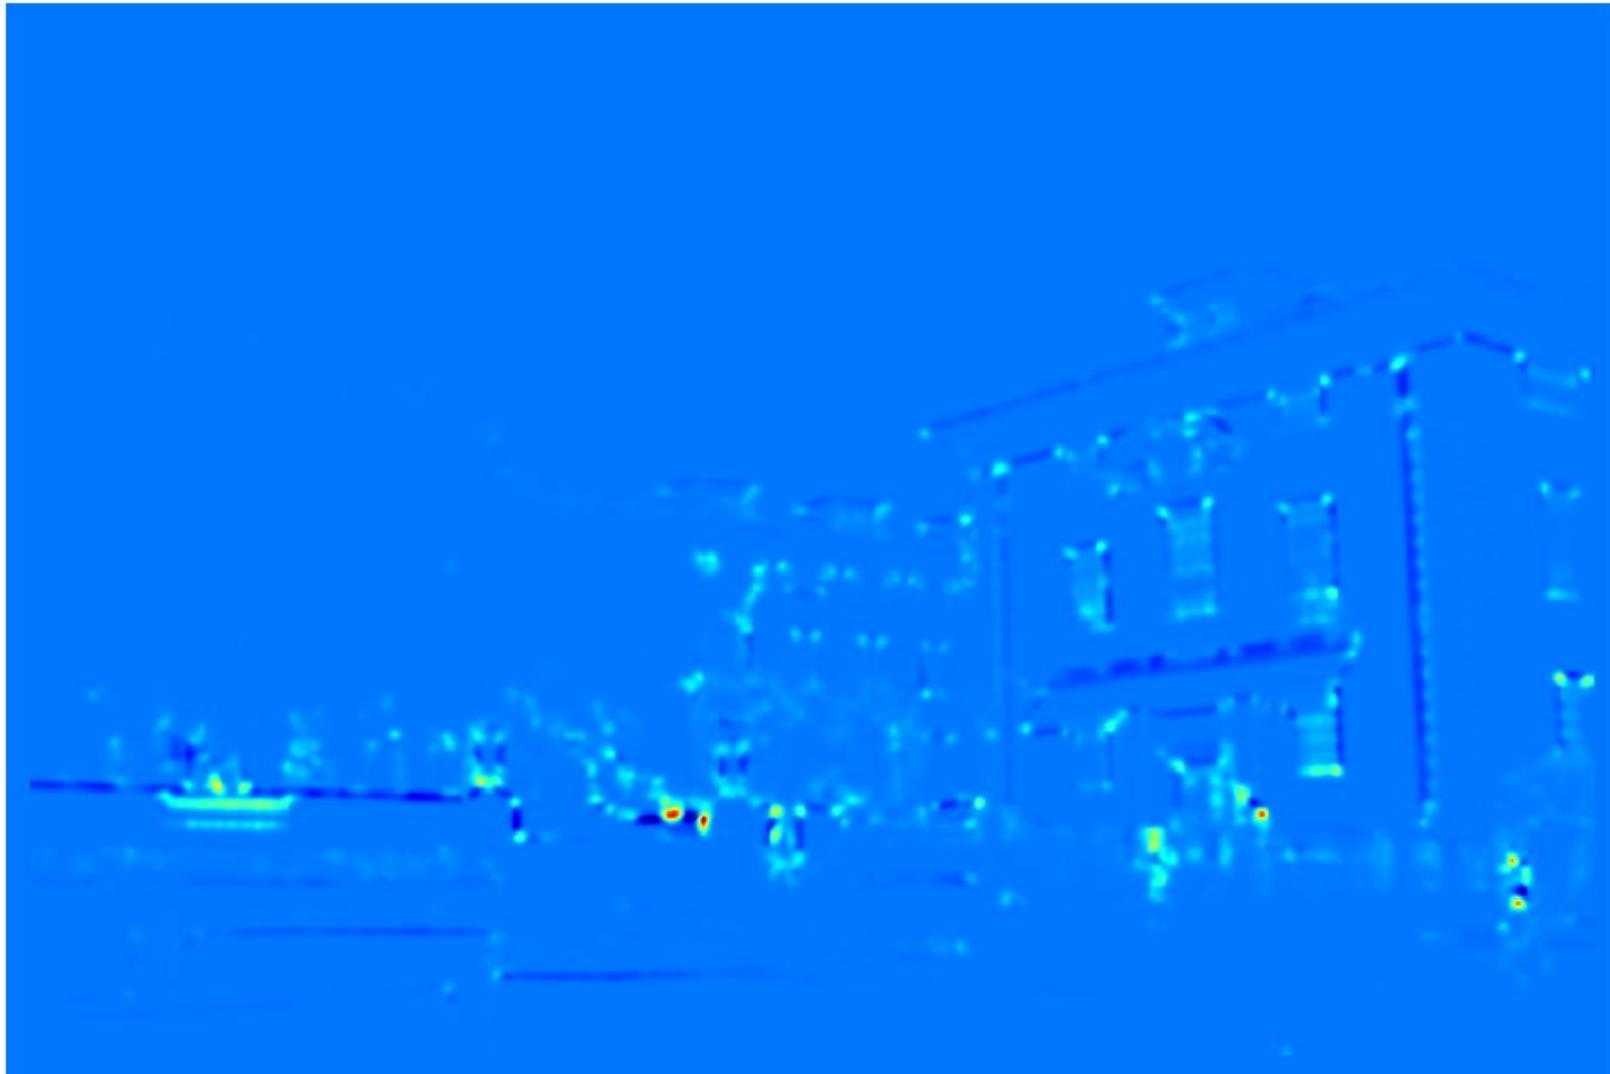
1. Cornerness scores
  - Compute M matrix for each image window
2. Threshold scores
  - Find points whose surrounding window gave large corner response  
 $(f > \text{threshold})$
3. Local filter for most cornerlike region
  - Take the points of local maxima  
*i.e., perform non-maximum suppression*

# Example of Harris application



# Example of Harris application

Compute corner response at every pixel.



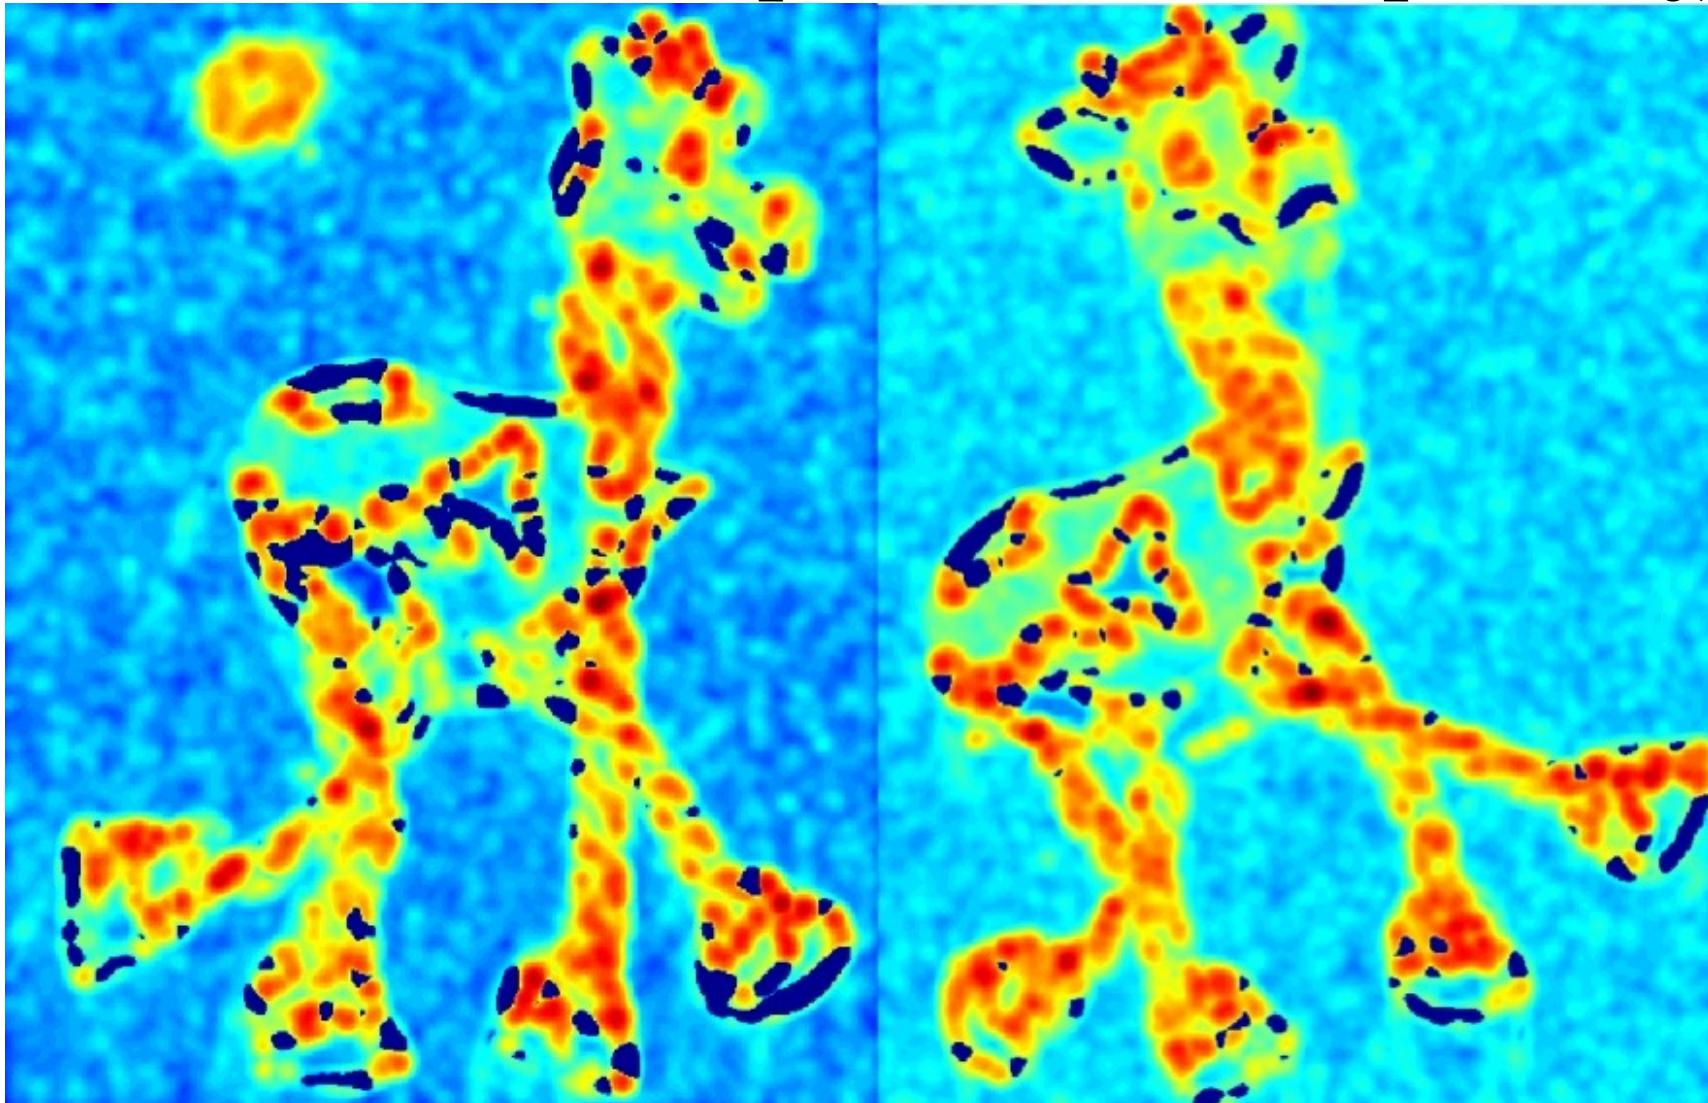
# Example of Harris application



# Harris Detector: Steps



# 1. Harris Detector: compute corner response ( $f$ )

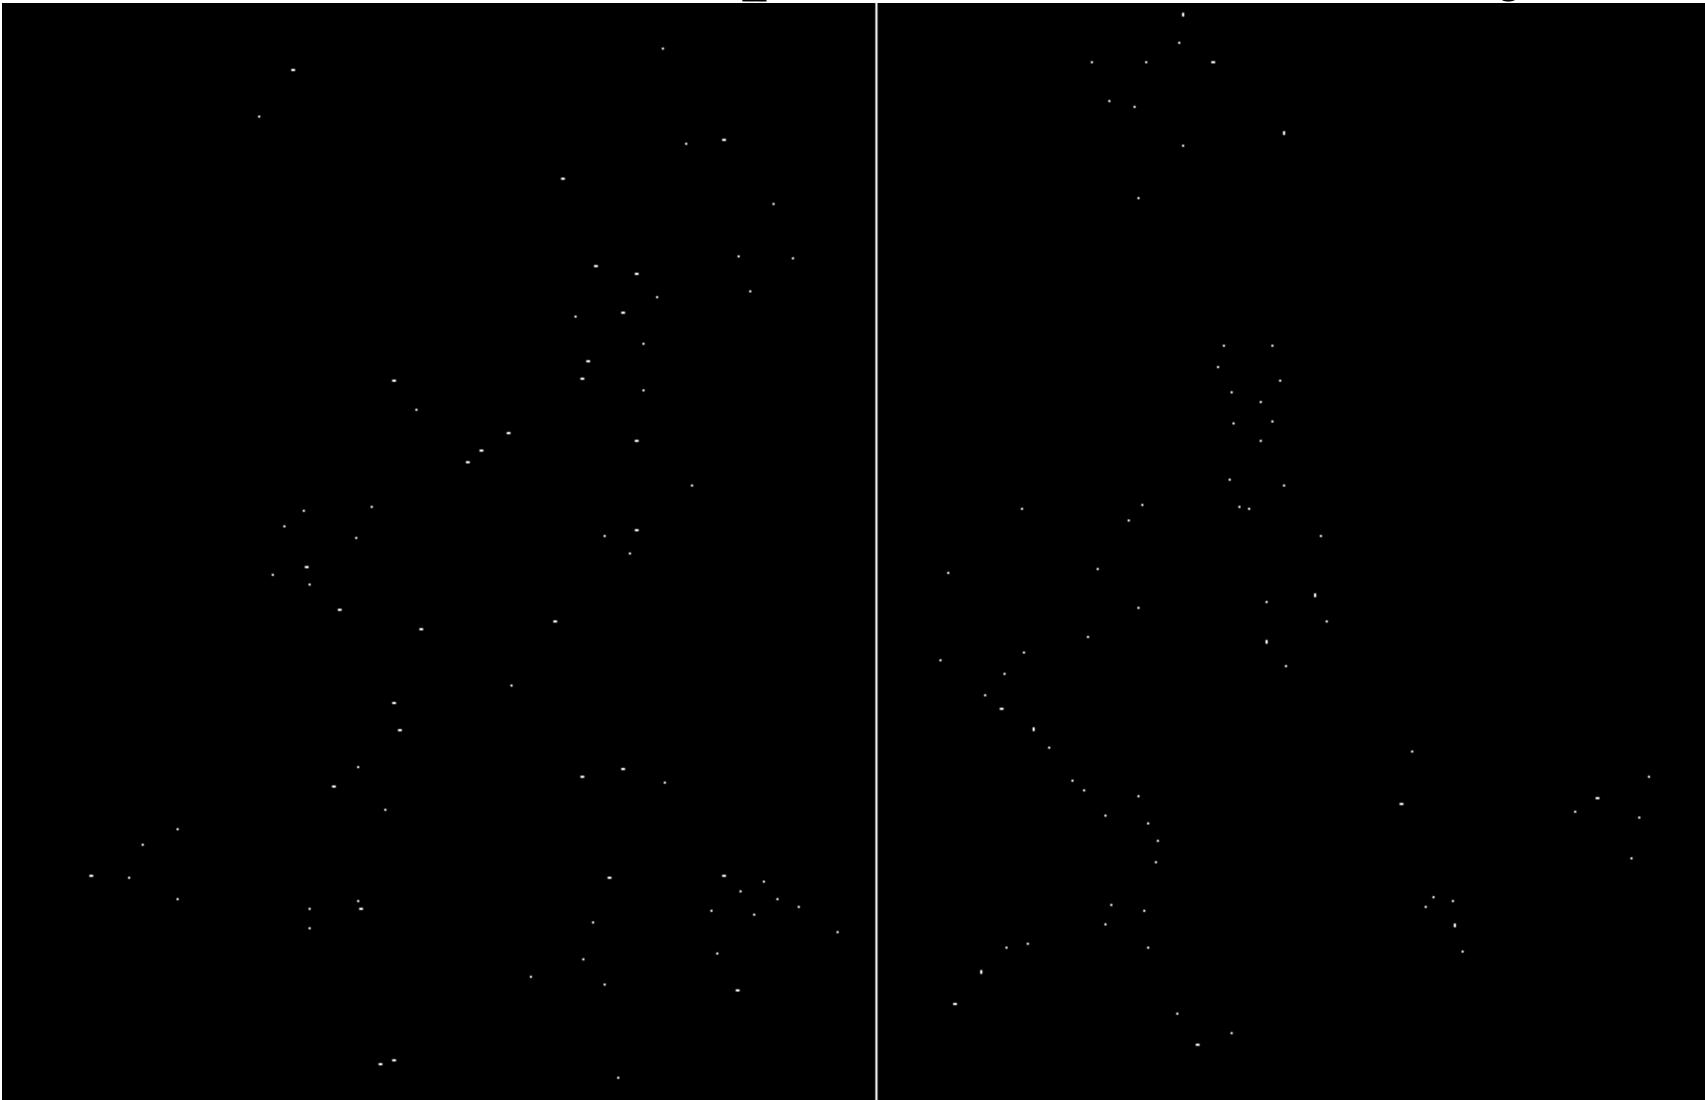


## 2. Harris: Filter points with large corner response

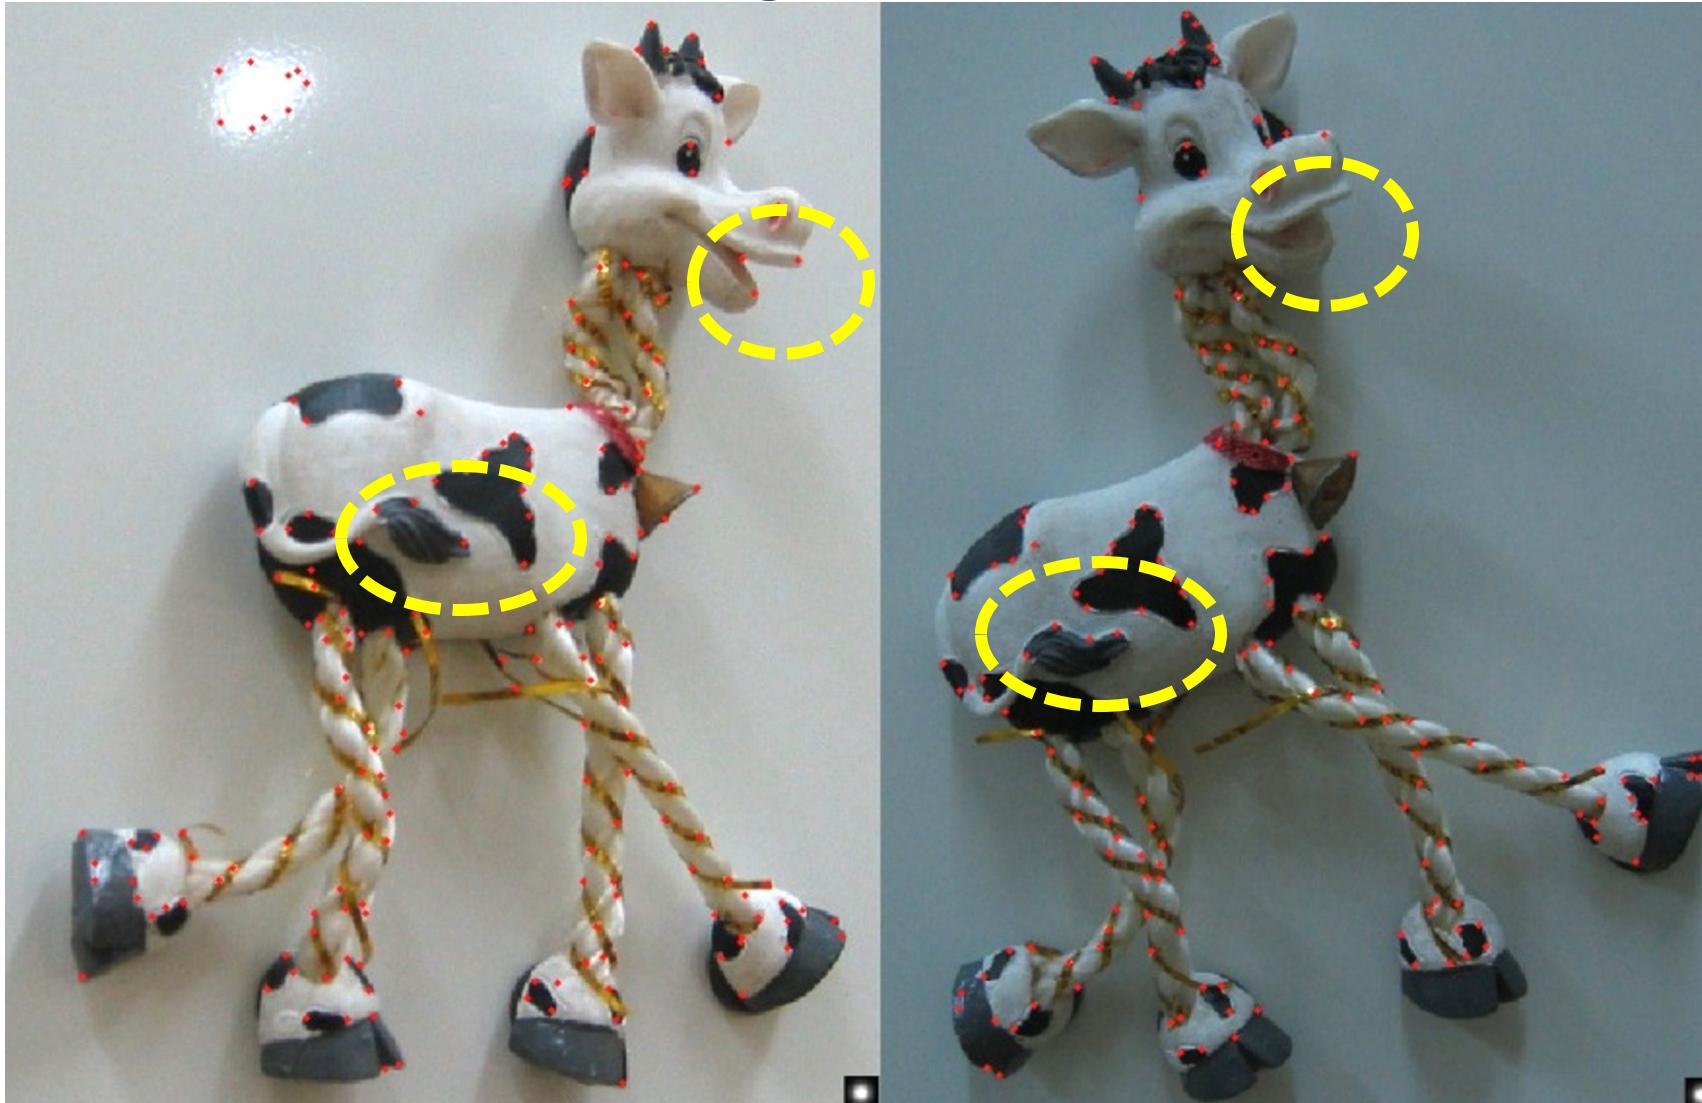
$f > \text{threshold}$



### 3. Harris Detector: Keep local maxima of $f$



# Harris Detector: Finding Matches



# Keypoints by Harris Corner Detector



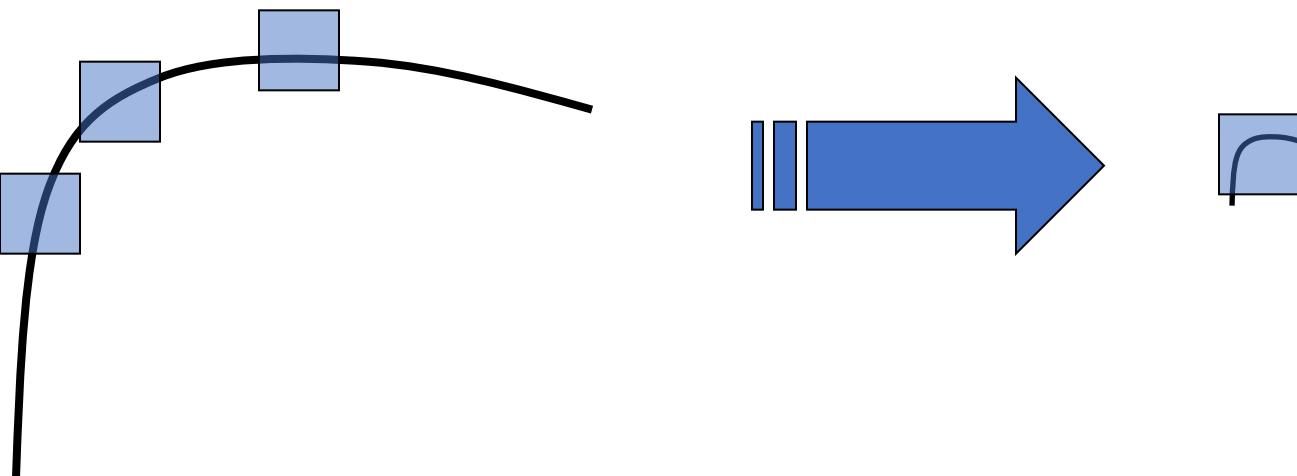
# Keypoints by Harris Corner Detector



# Properties of the Harris corner detector

- Rotation invariant? Yes

- Scale invariant? No



All points will be  
classified as edges

Corner !

# Local Feature Detection

Blob Detector with Automatic Scale Selection

# Scale invariant interest points

How can we independently select interest points in each image, such that the detections are repeatable across different scales?



# Scale invariant interest points

How can we independently select interest points in each image, such that the detections are repeatable across different scales?



# Scale invariant interest points

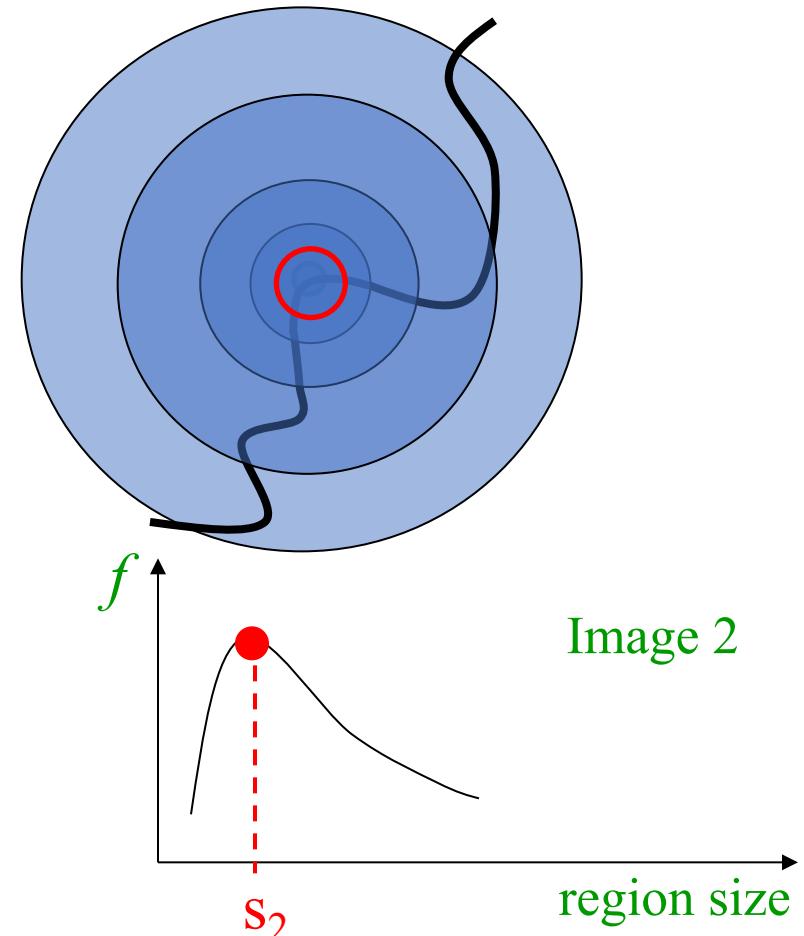
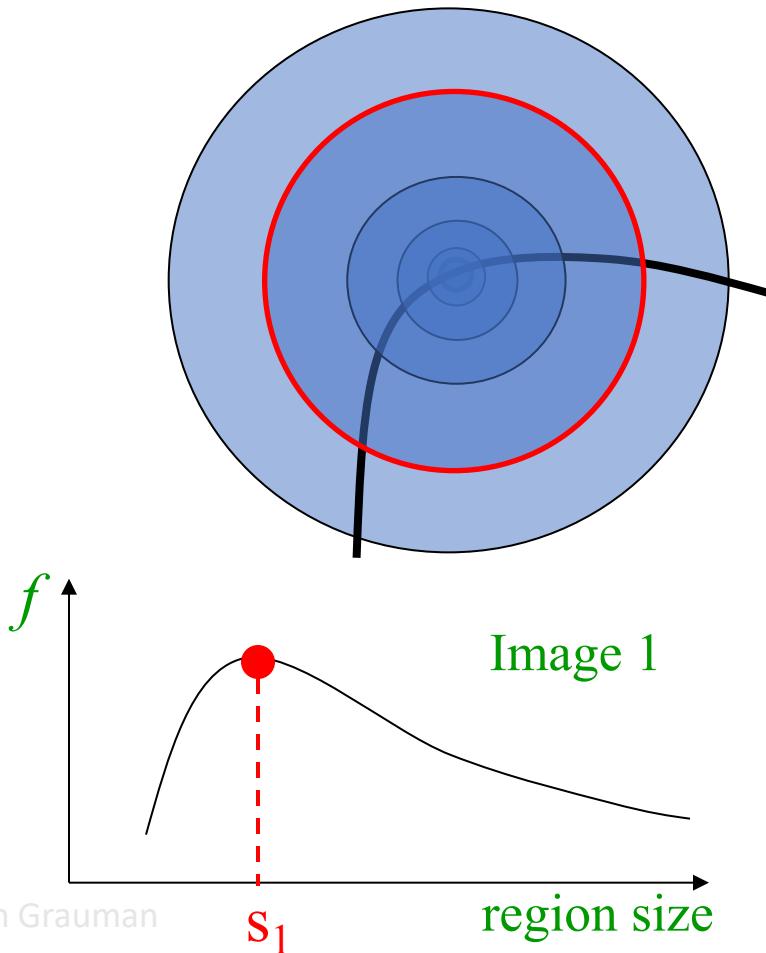
How can we independently select interest points in each image, such that the detections are repeatable across different scales?



# Automatic scale selection

## Intuition:

- Find scale that gives local maxima of some function  $f$  in both position and scale.



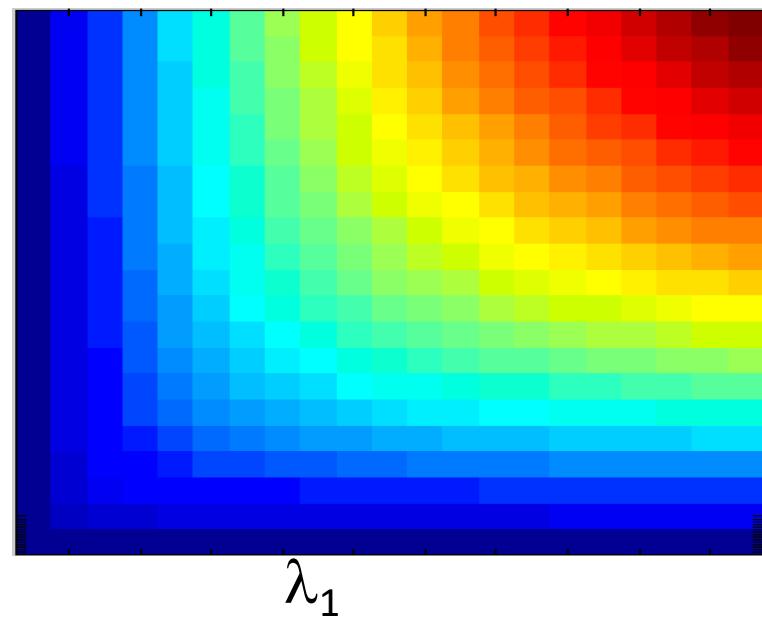
# Choosing $f$

- What can be the “signature” function?

$$M = X \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} X^T$$

$$f = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2}$$

$\lambda_2$



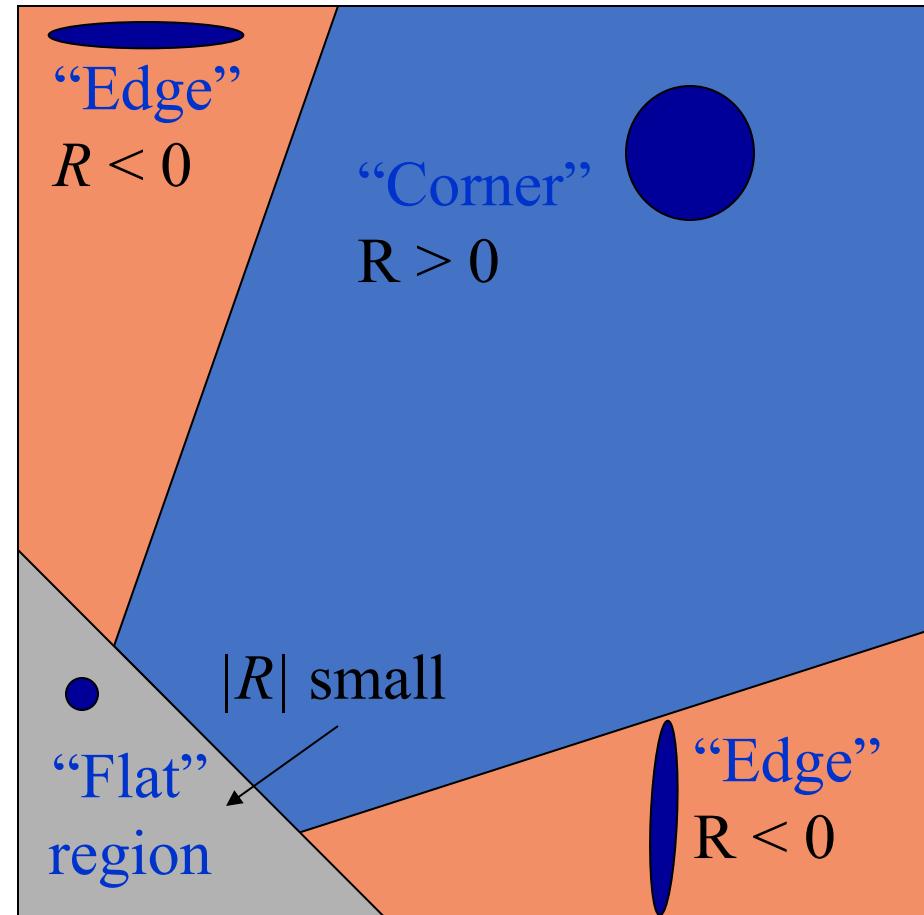
# Choosing $f$

$$\begin{aligned} R &= \det(\mathbf{M}) - \alpha \operatorname{trace}(\mathbf{M})^2 \\ &= \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2 \end{aligned}$$

$\alpha$ : constant (0.04 to 0.06)

Calculating eigenvalues is usually slow!

Determinant, trace have closed form solution

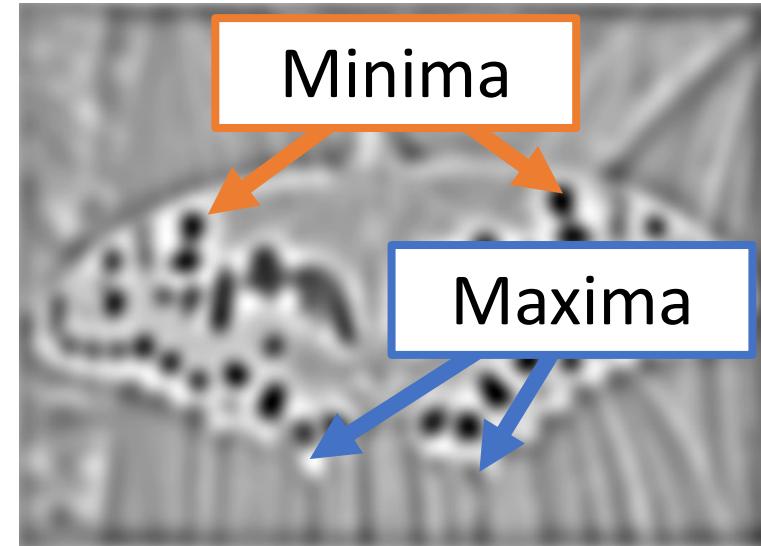


# Blob Detection

Another detector (has some nice properties)

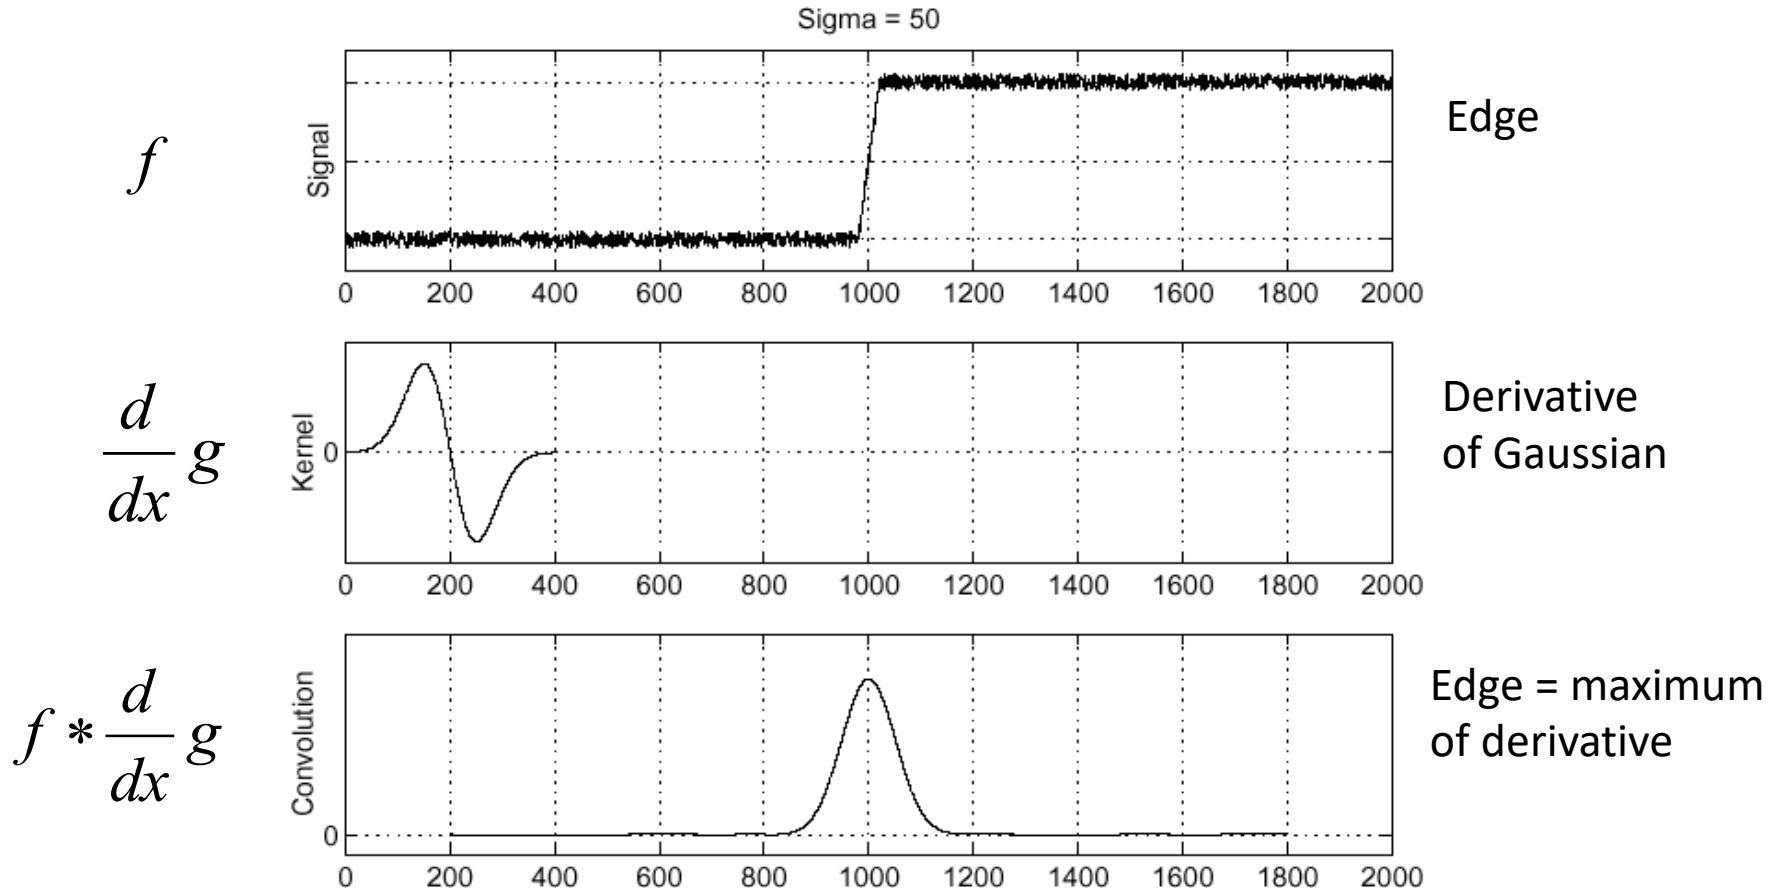


$$* \quad \bullet =$$



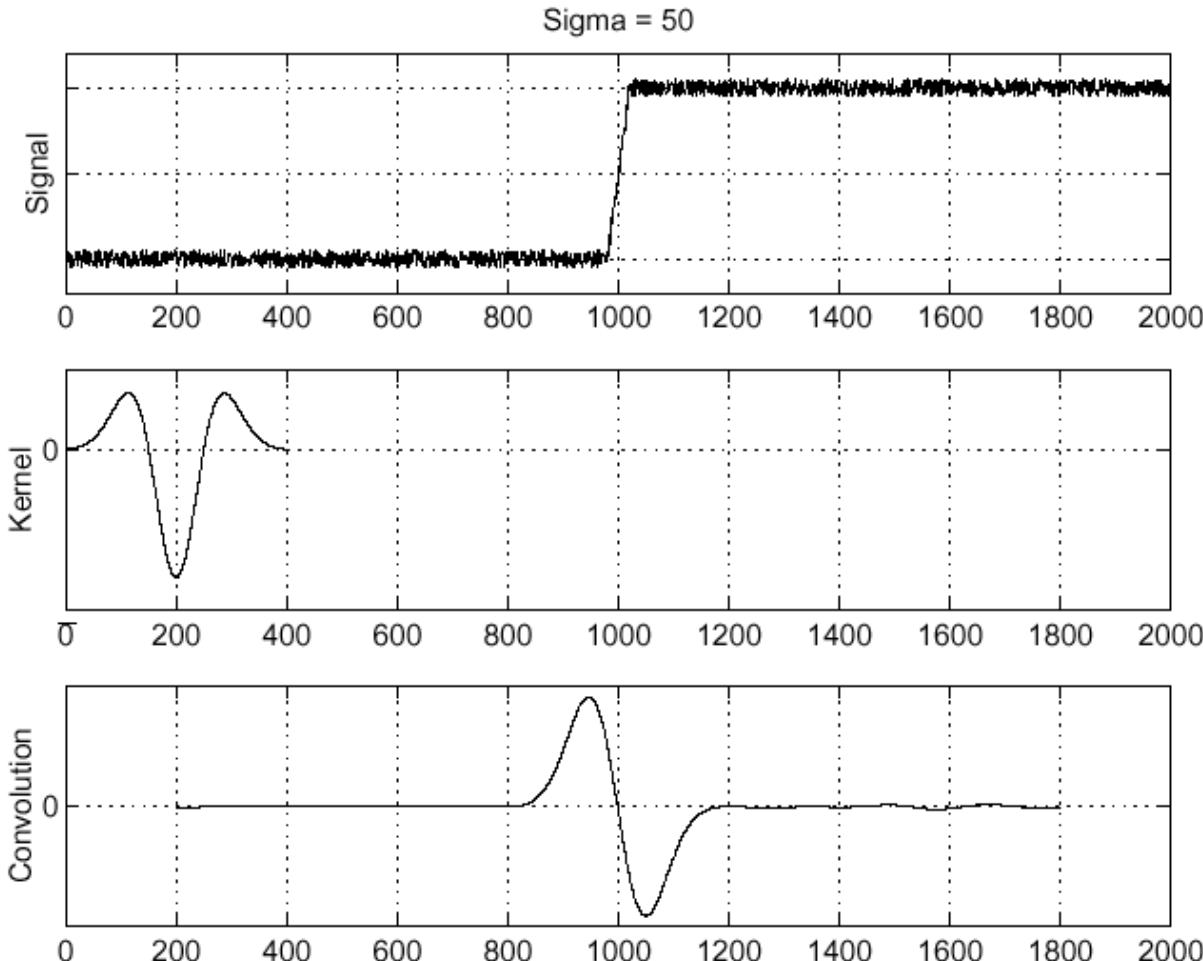
Find maxima *and minima* of blob filter response in scale  
*and space*

# Recall: Edge Detection



# Recall: Edge Detection

$$f * \frac{d^2}{dx^2} g$$



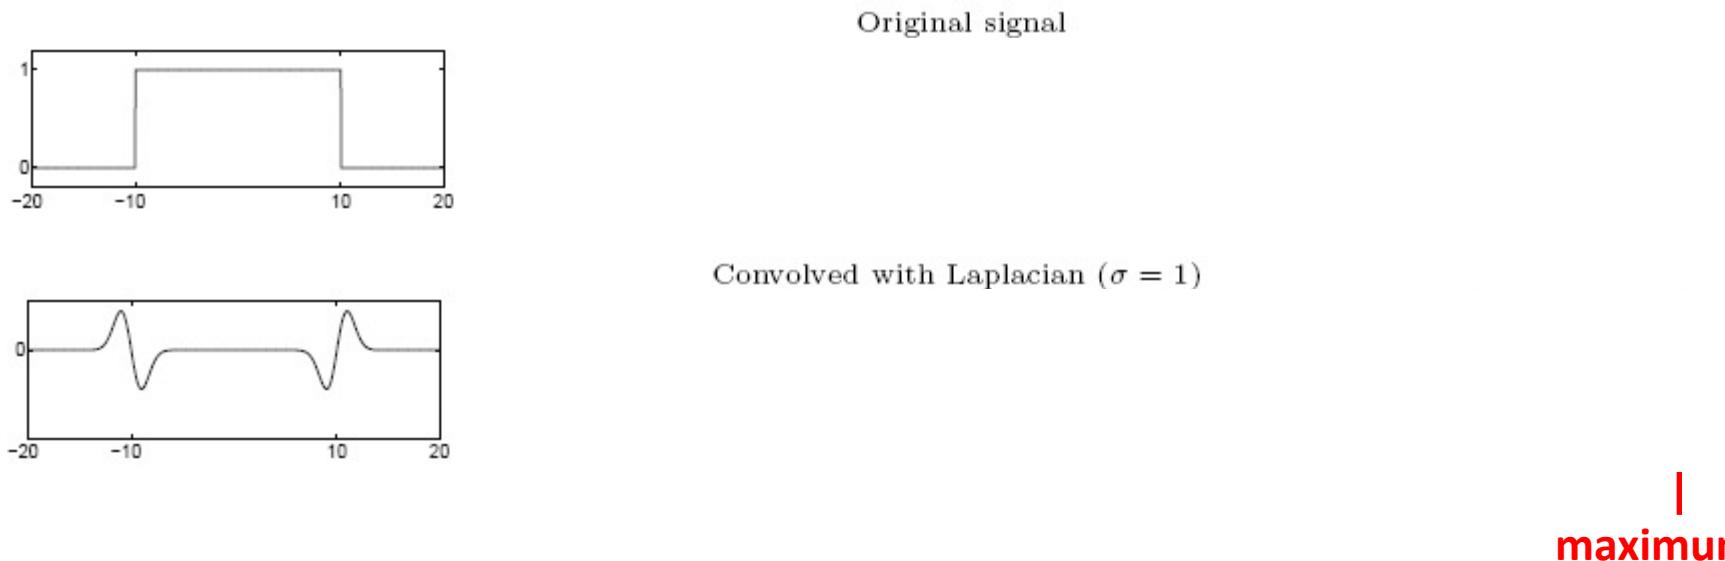
Edge

Second derivative  
of Gaussian  
(Laplacian)

Edge = zero crossing  
of second derivative

# From edges to blobs

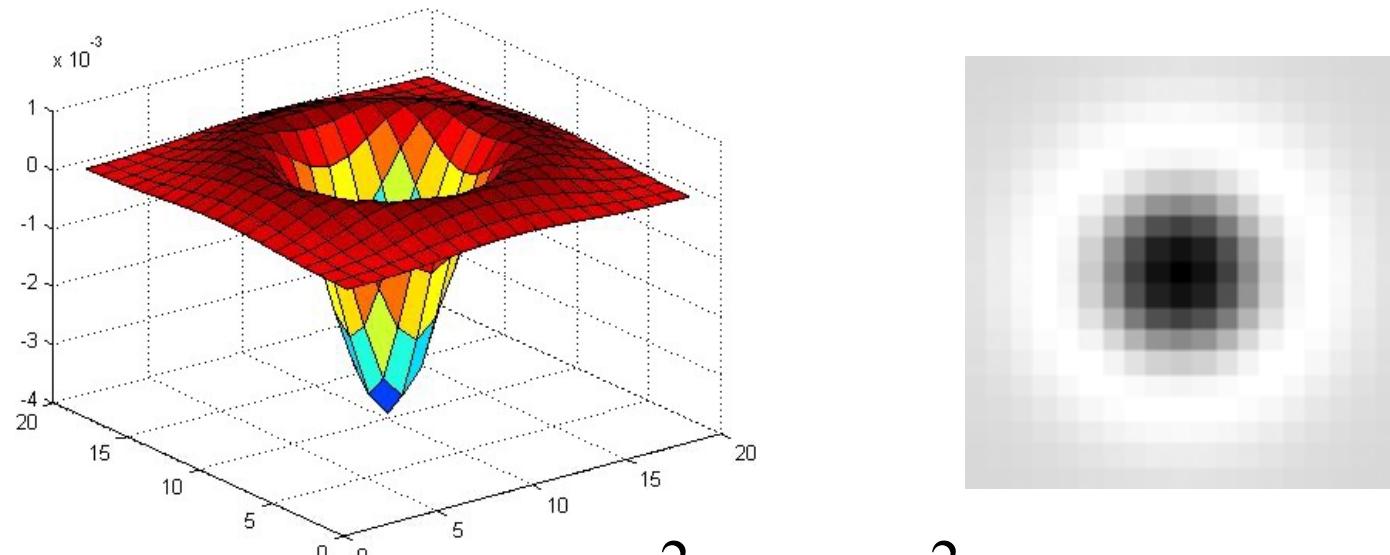
- Edge = ripple
- Blob = superposition of two ripples



**Spatial selection:** the **magnitude** of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

# Blob detection in 2D

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

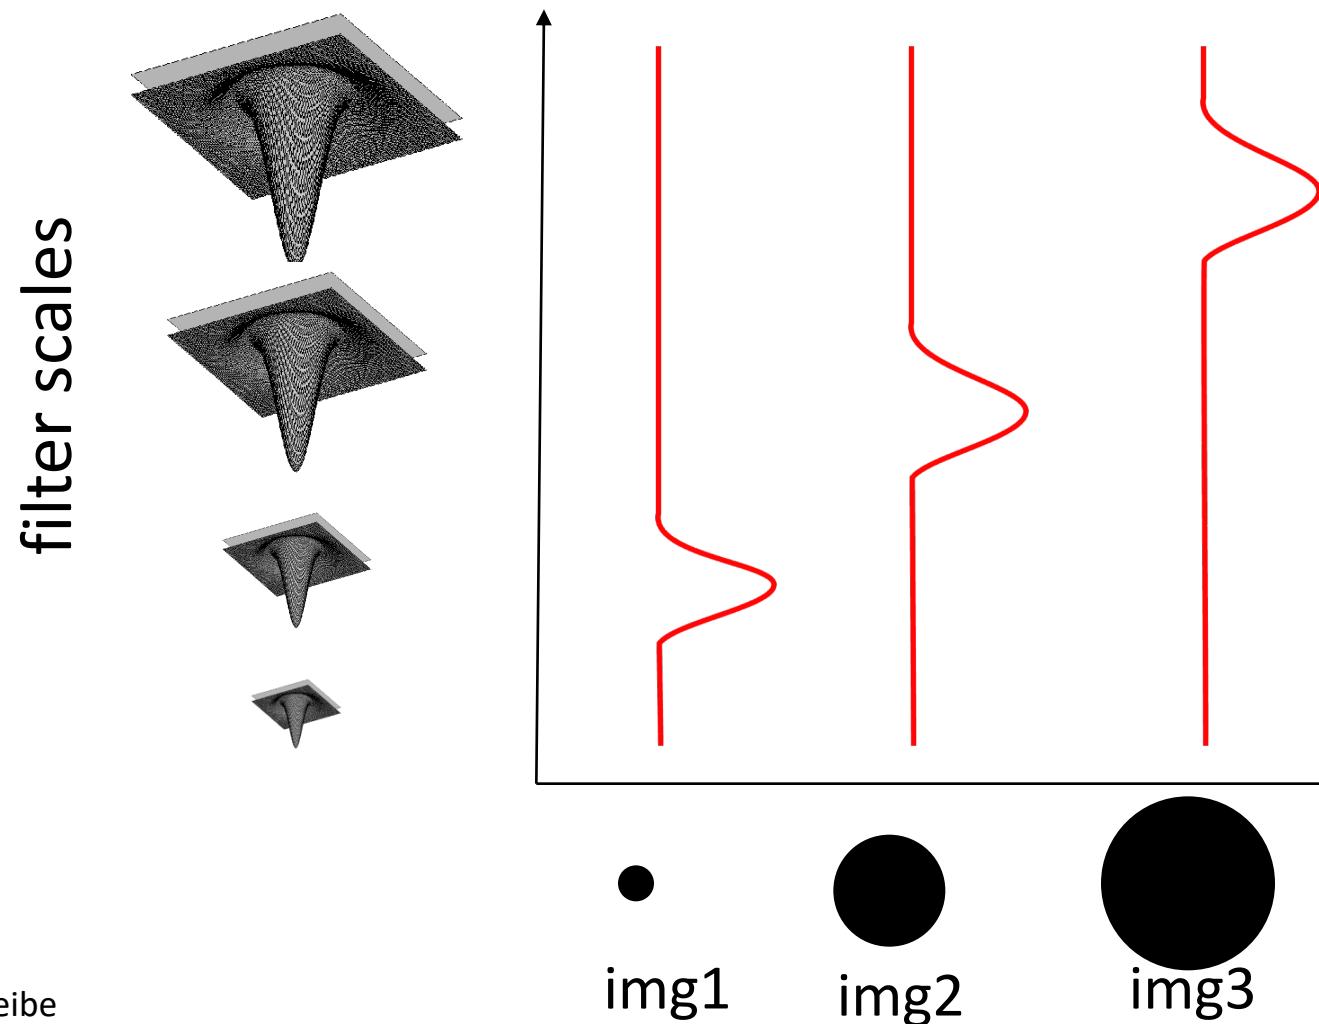


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Blob detection in 2D: scale selection

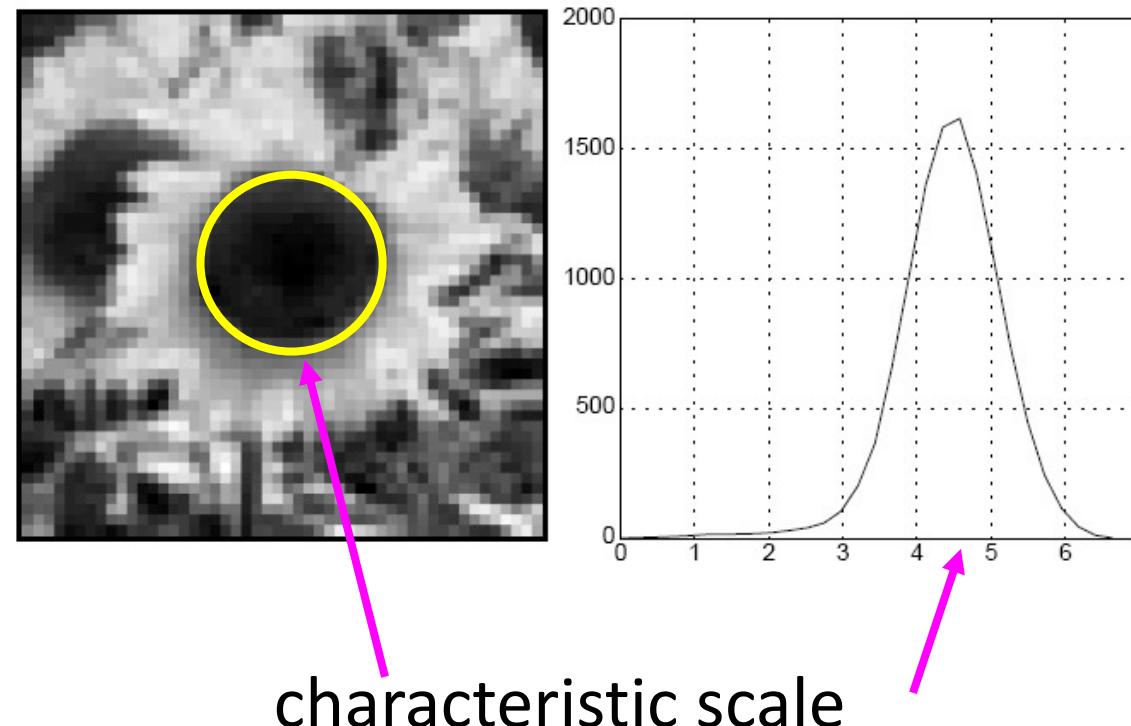
Laplacian-of-Gaussian = “blob” detector

$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$



# Blob detection in 2D

- We define the *characteristic scale* as the scale that produces peak of Laplacian response

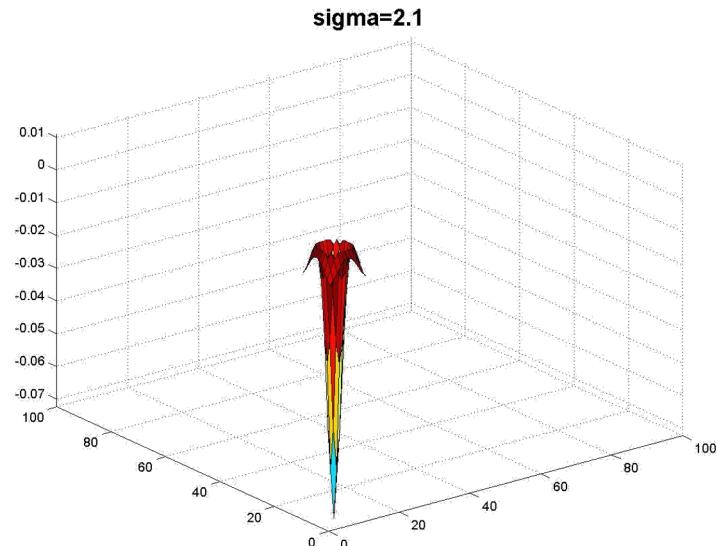


# Example

Original image at  
 $\frac{3}{4}$  the size

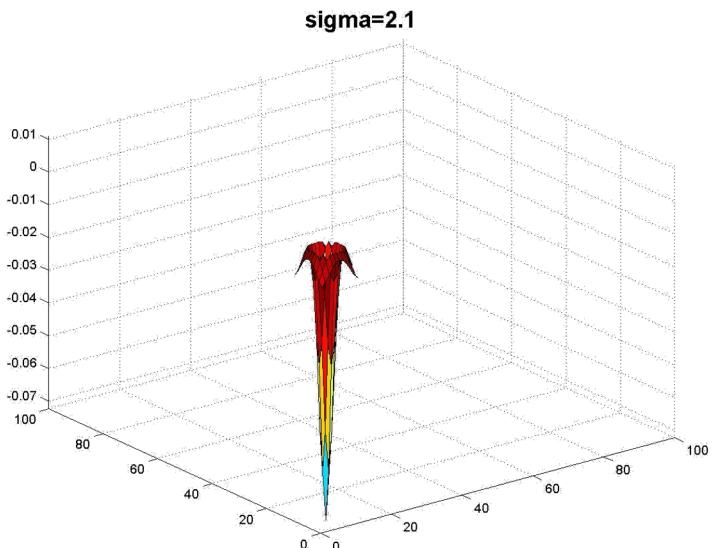
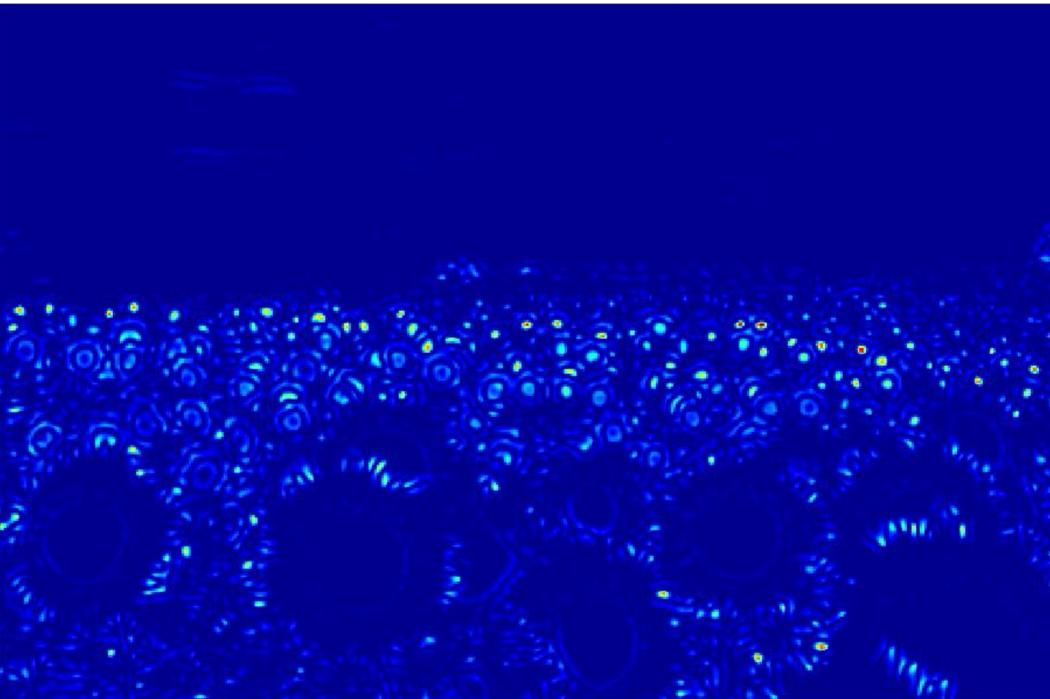
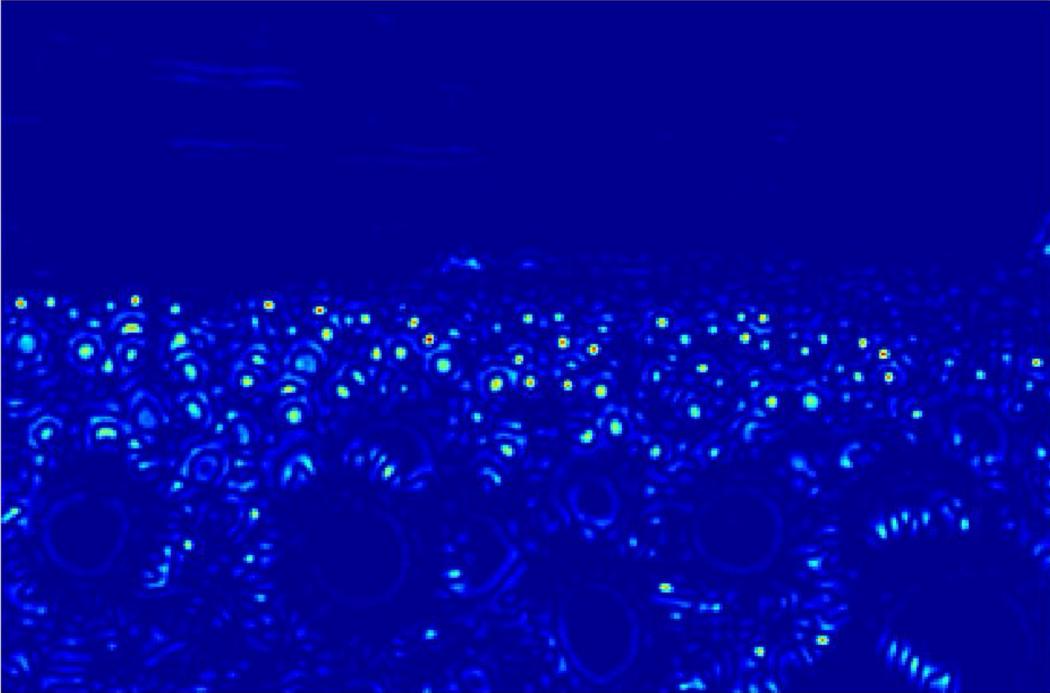


What will we recognize?

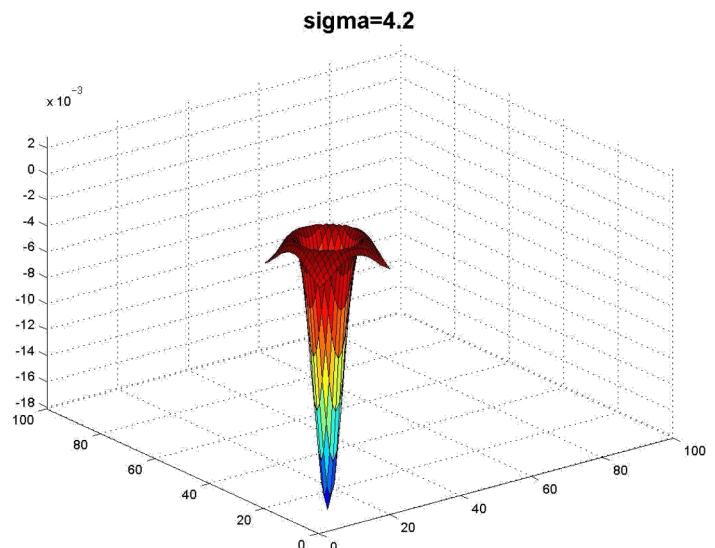
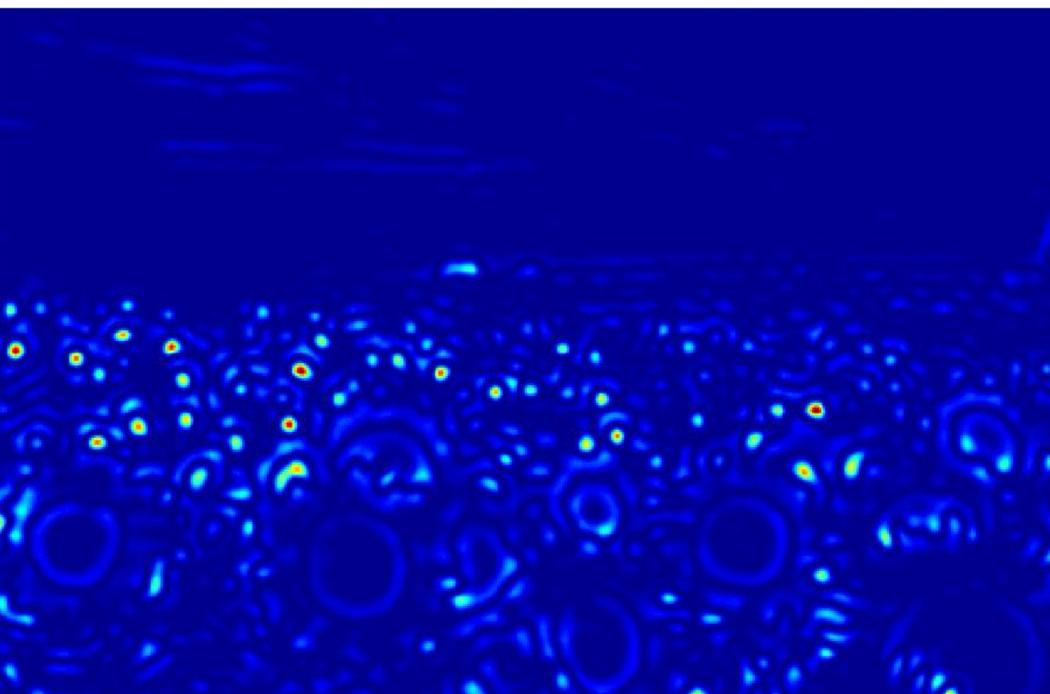
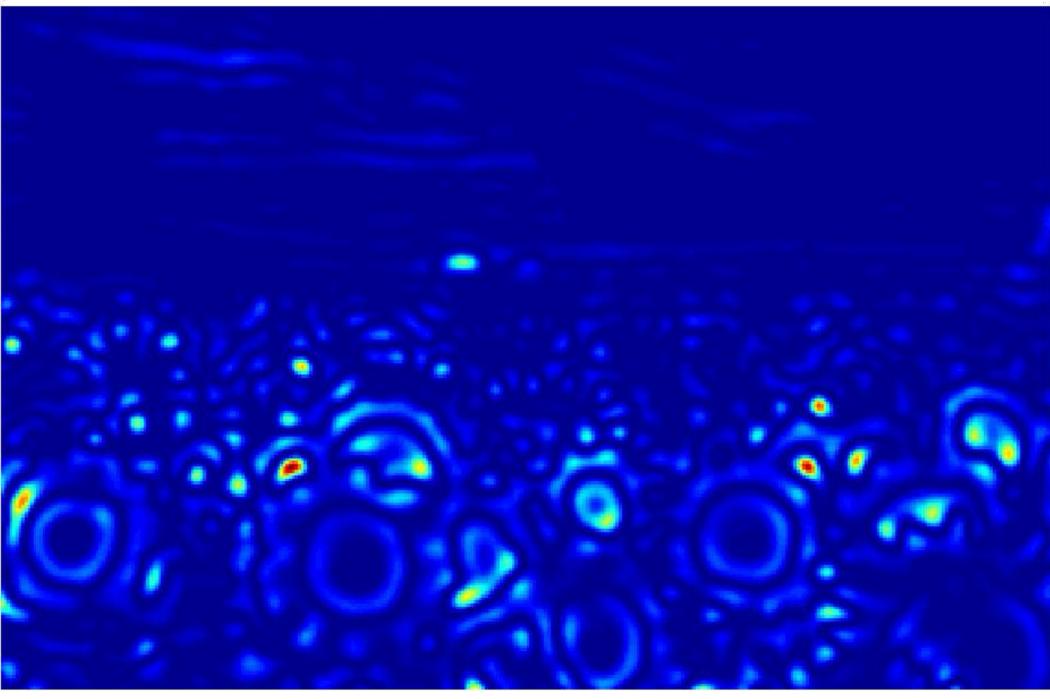


# Example

Original image at  
 $\frac{3}{4}$  the size

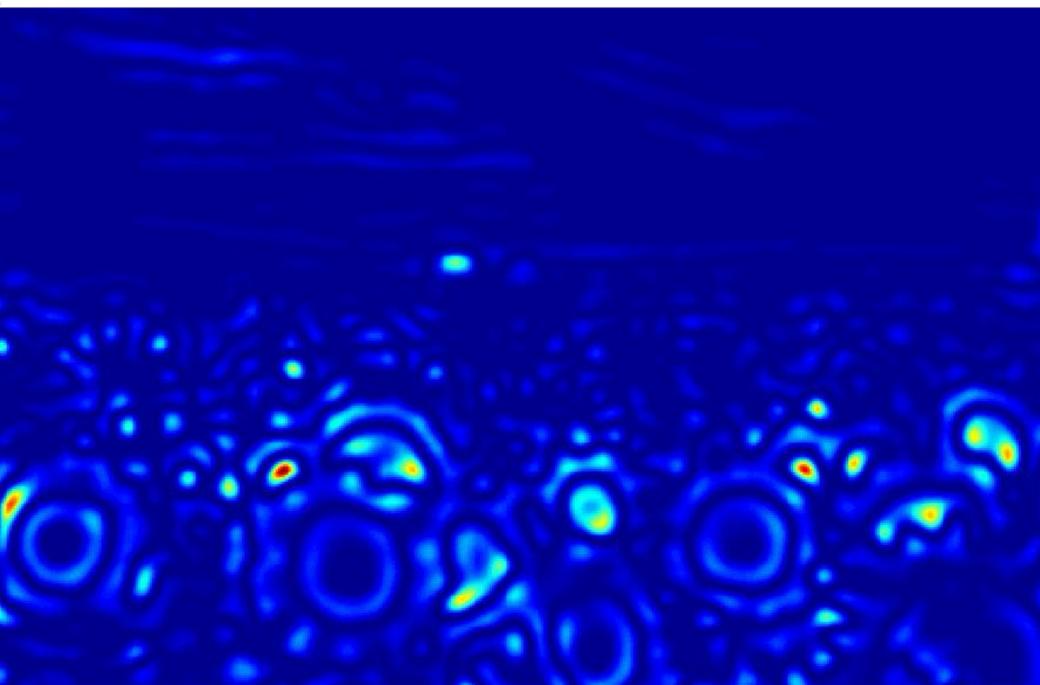
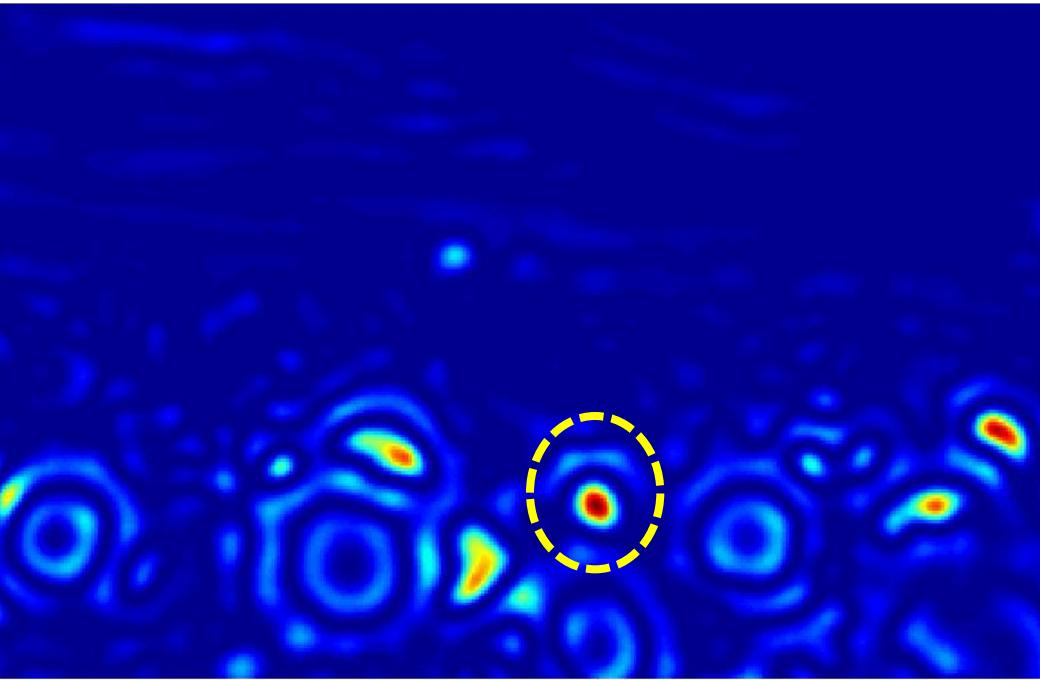
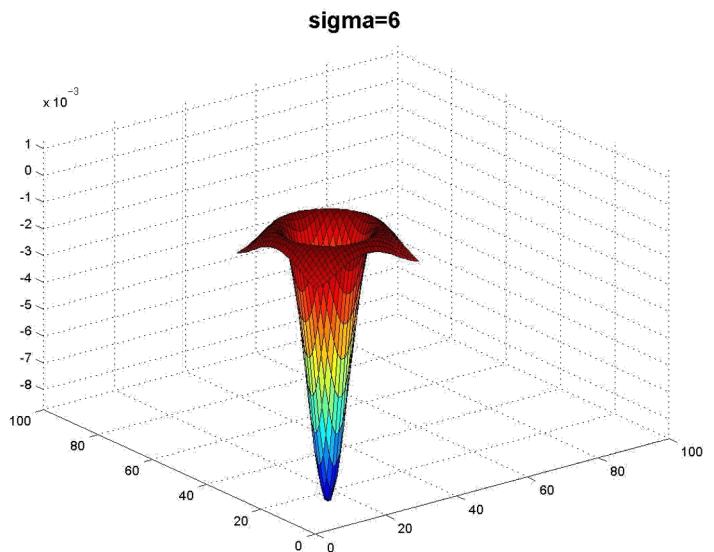


# Example

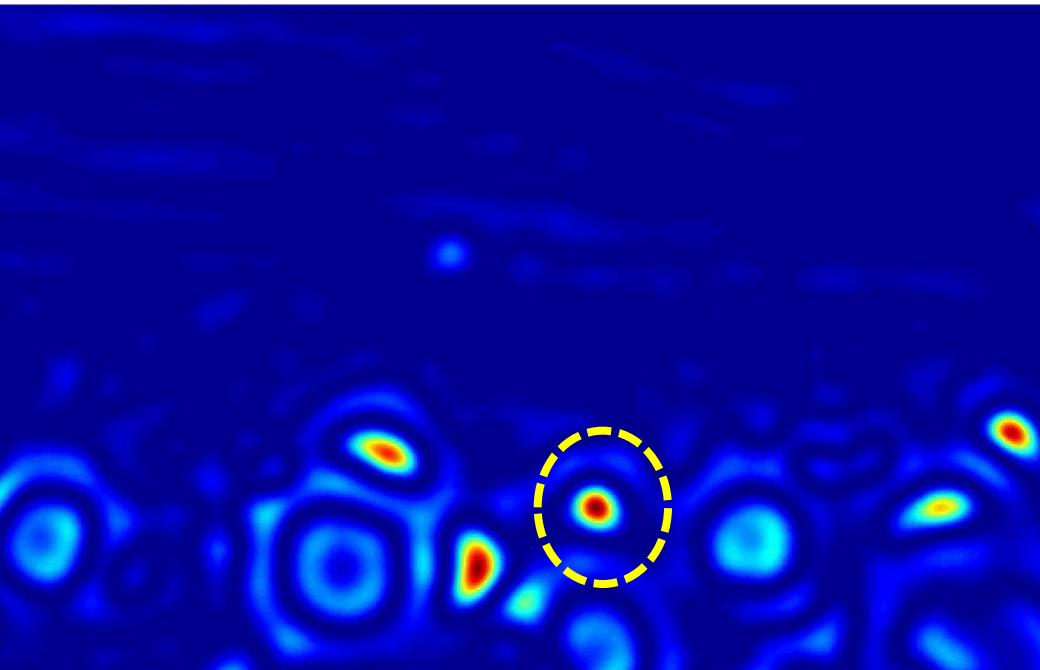
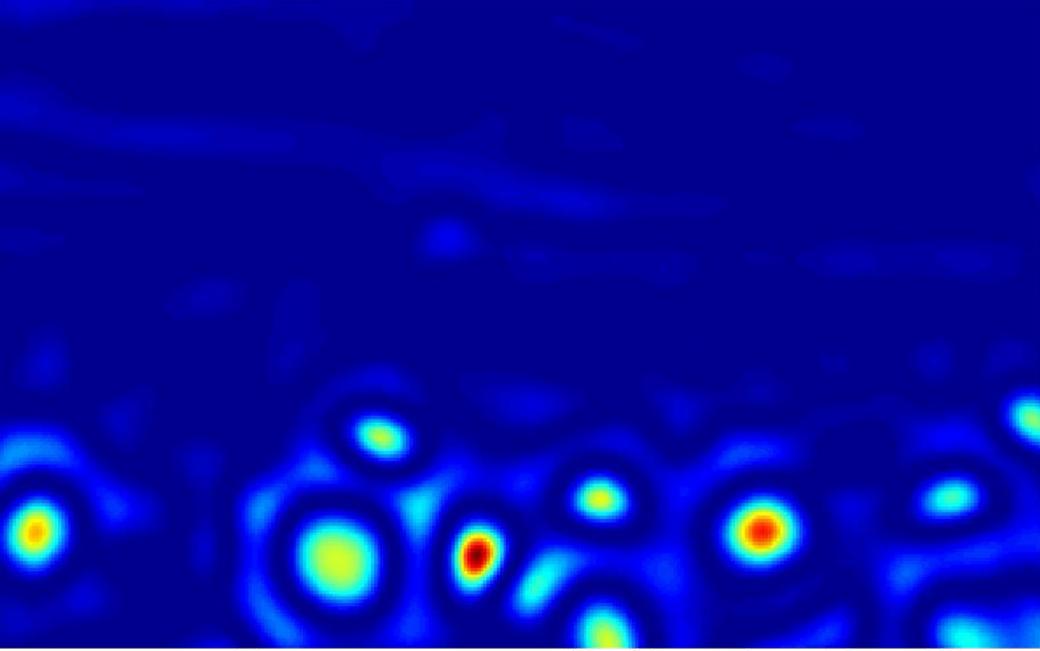
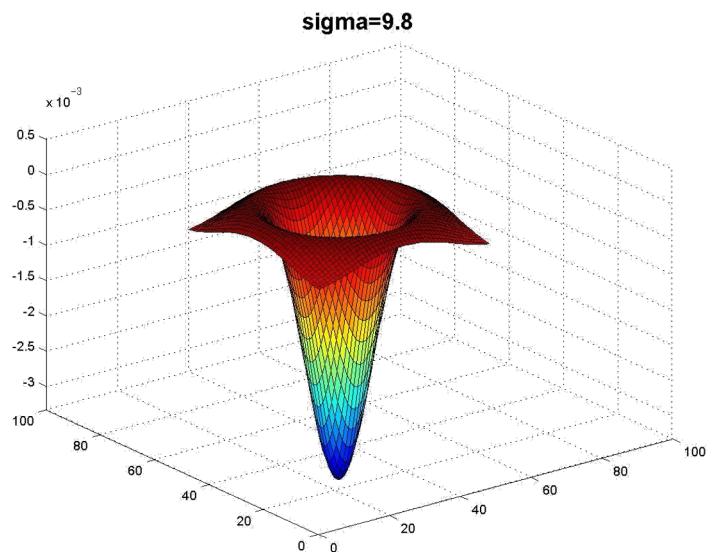


# Example

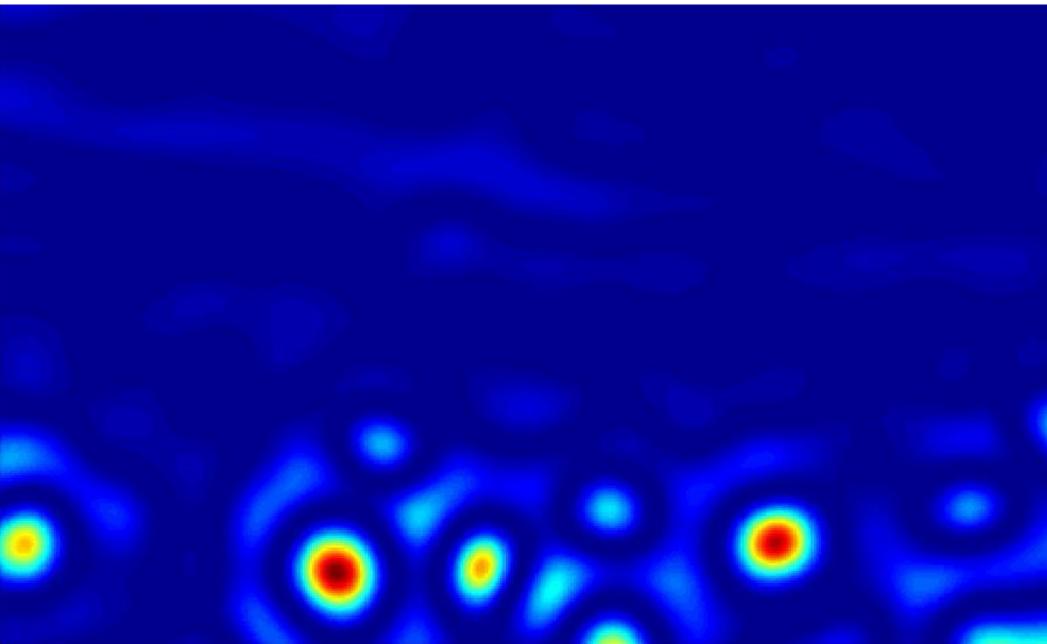
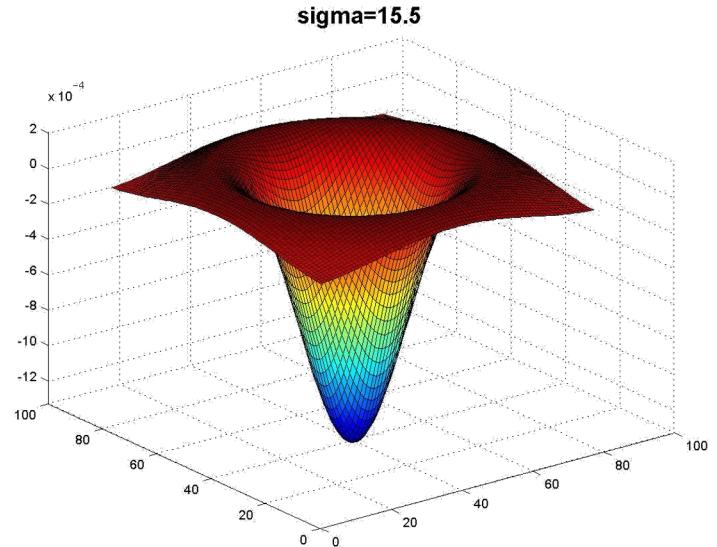
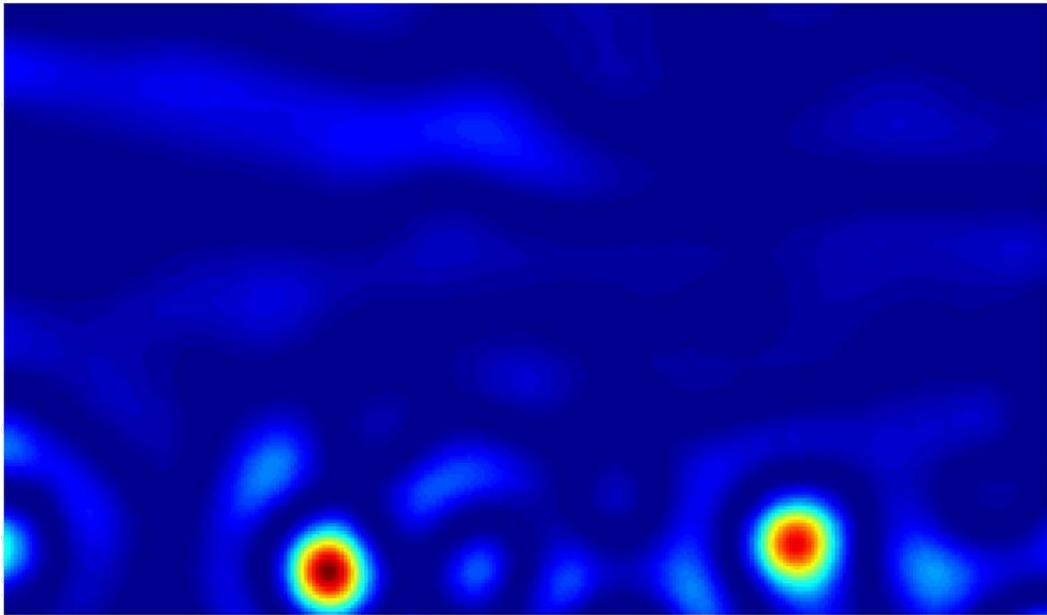
Do we recognize the same points in the two images?



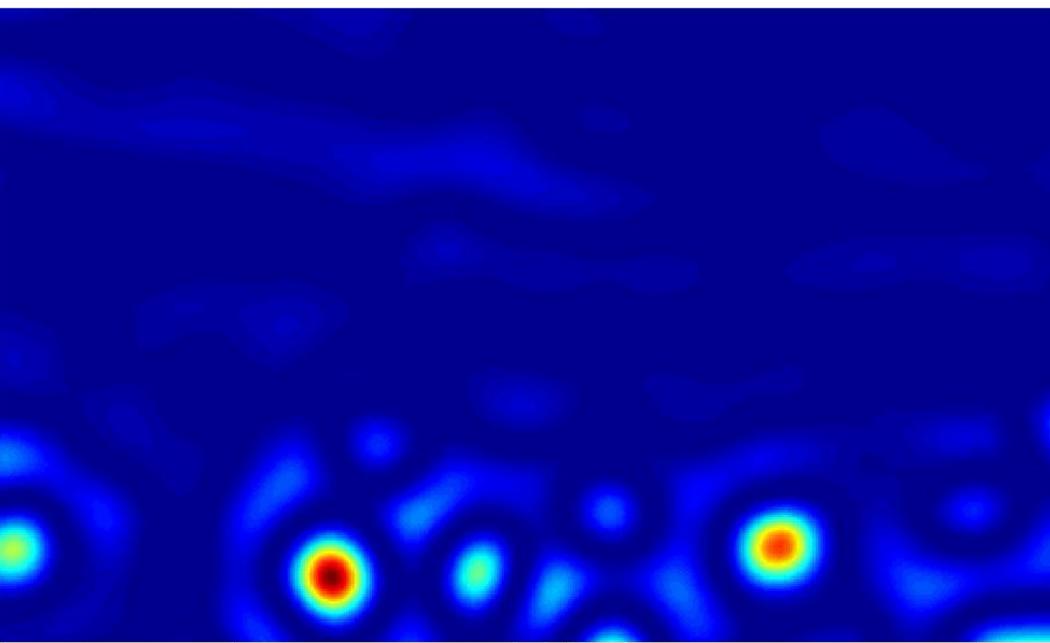
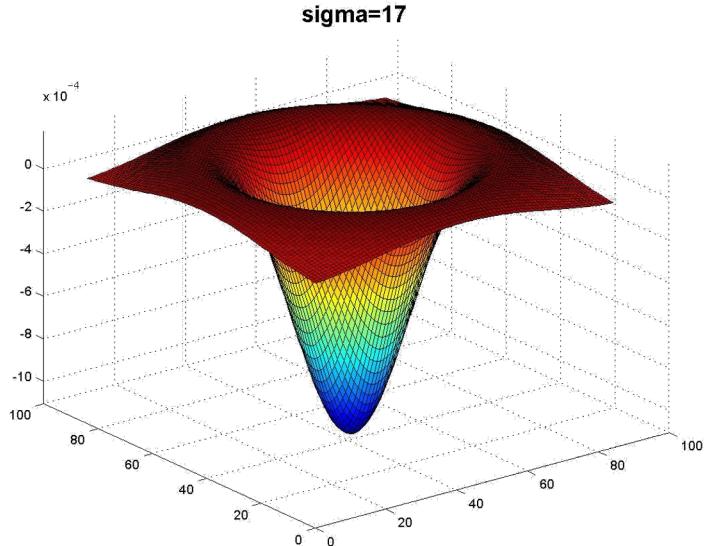
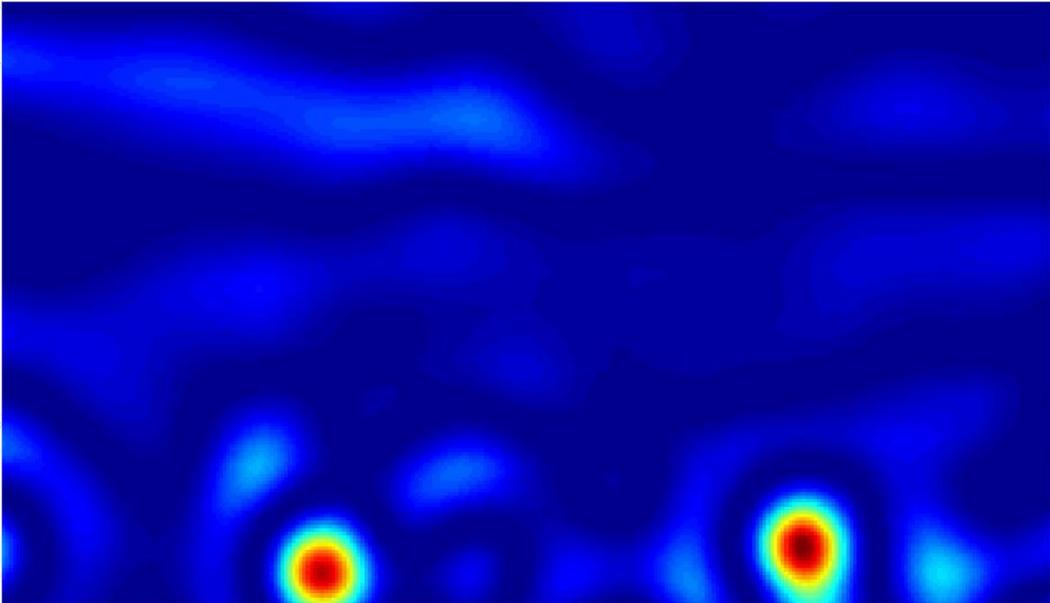
# Example



# Example

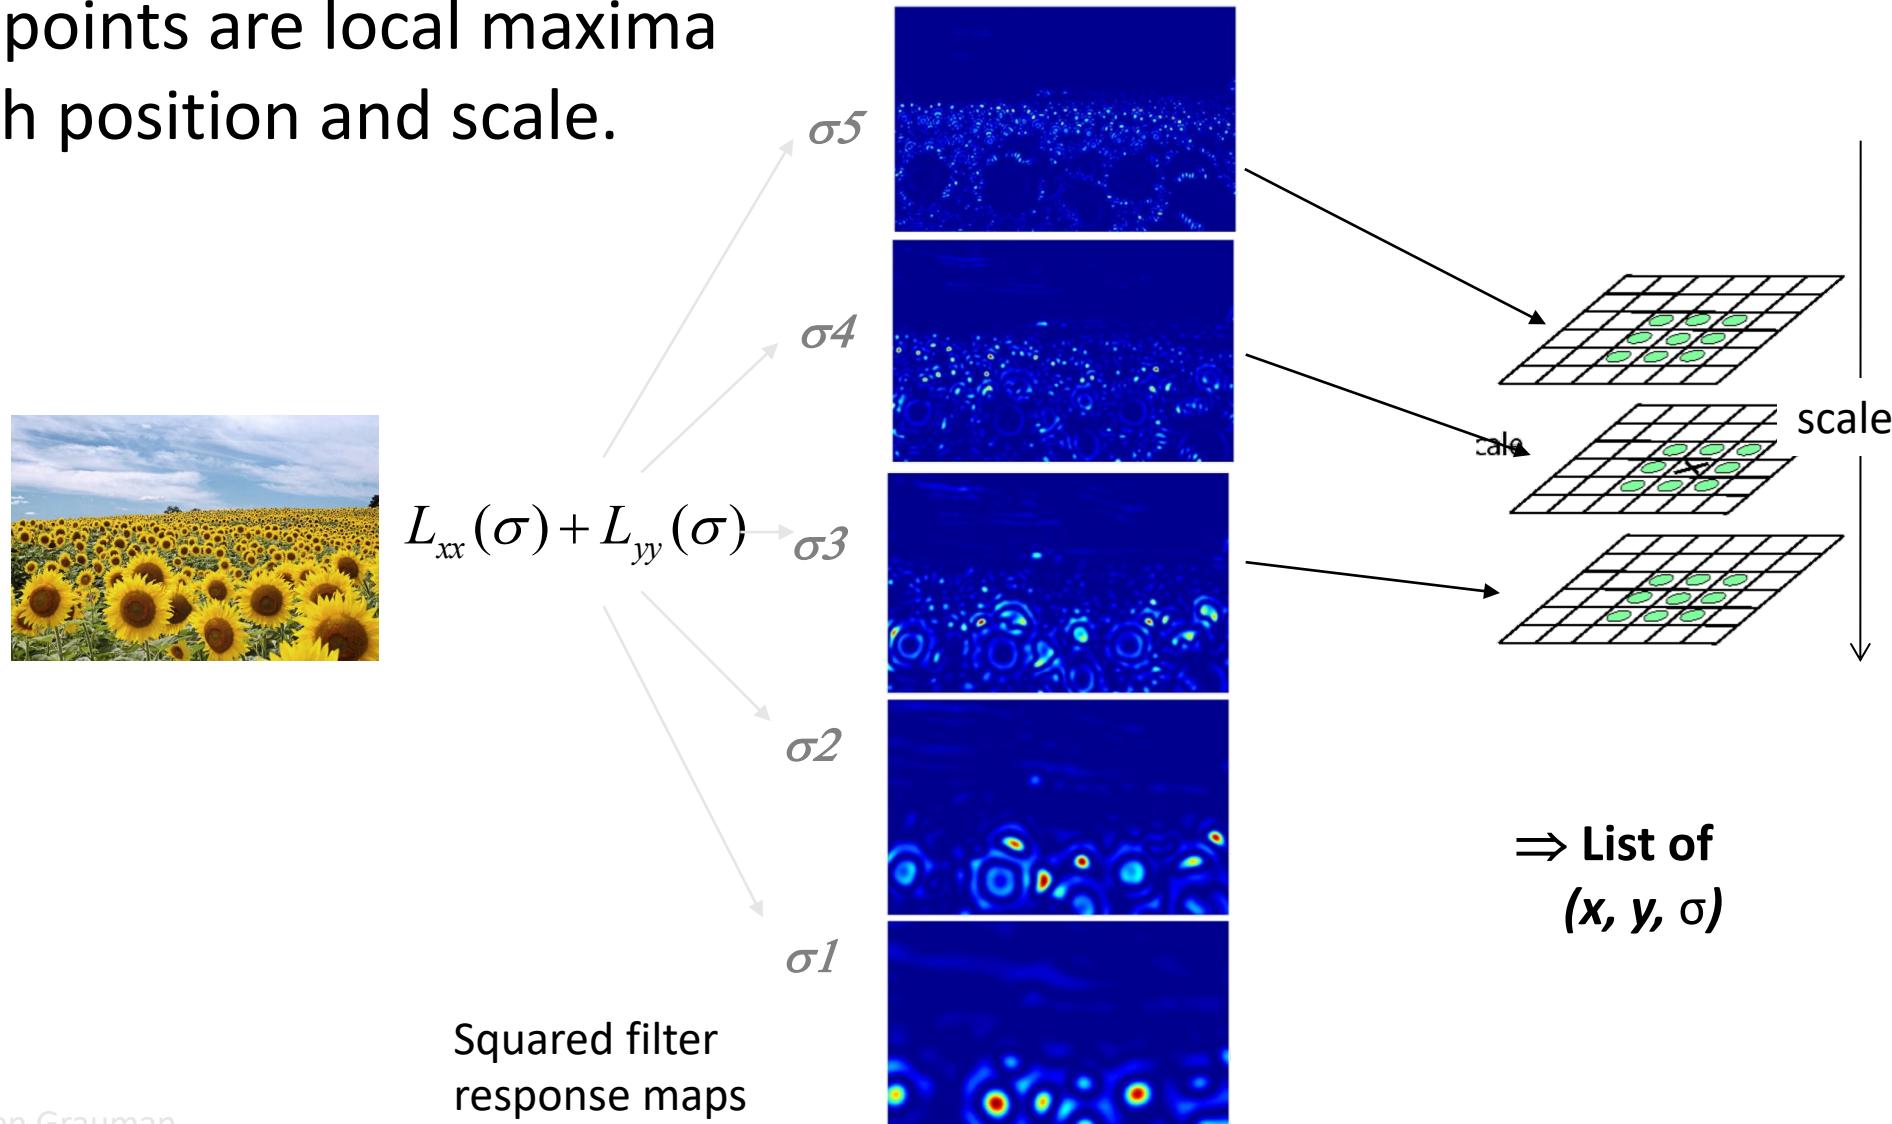


# Example



# Scale invariant interest points

Interest points are local maxima  
in both position and scale.

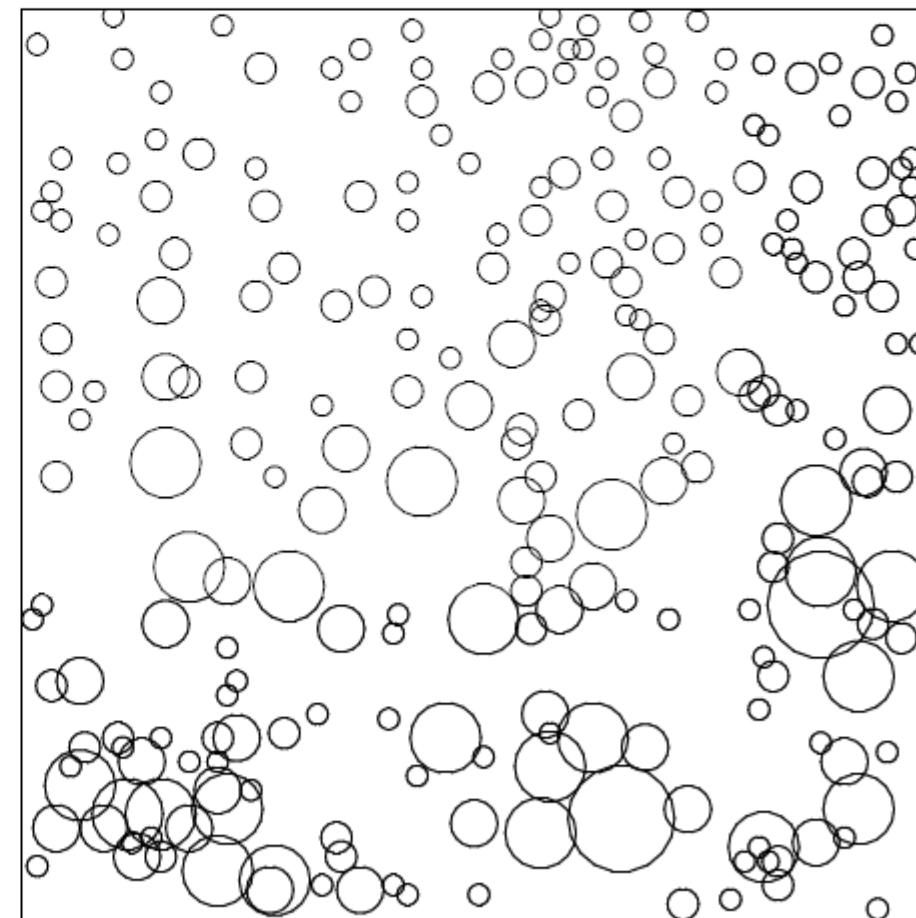


# Scale-space blob detector: Example

*original image*



*scale-space maxima of  $(\nabla_{norm}^2 L)^2$*



T. Lindeberg. Feature detection with automatic scale selection. IJCV 1998.

# Scale-space blob detector: Example



# Questions?

See you next time!