

# CS 6476 Computer Vision:

## 7. Fitting & Hough Transform

Humphrey Shi

9/12/2024



# Outline

- Logistics & Course Recap
- Fitting & Hough Transform

# Logistic & Course Recap

# Assignments

## Assignment 2: SIFT Local Feature Matching

SEP 8, 2024 10:30 AM

SEP 22, 2024 11:59 PM

10

Late Due Date: SEP 27, 2024 11:59 PM

## Assignment 1: Image Filtering

AUG 26, 2024 8:00 AM

SEP 8, 2024 11:59 PM

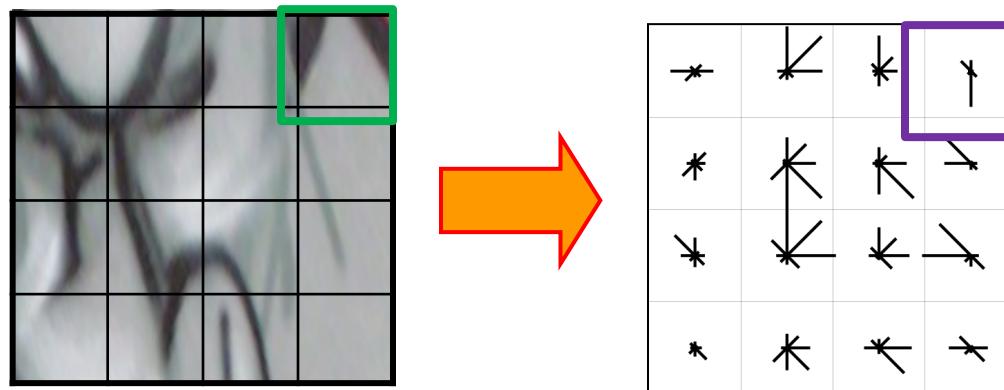
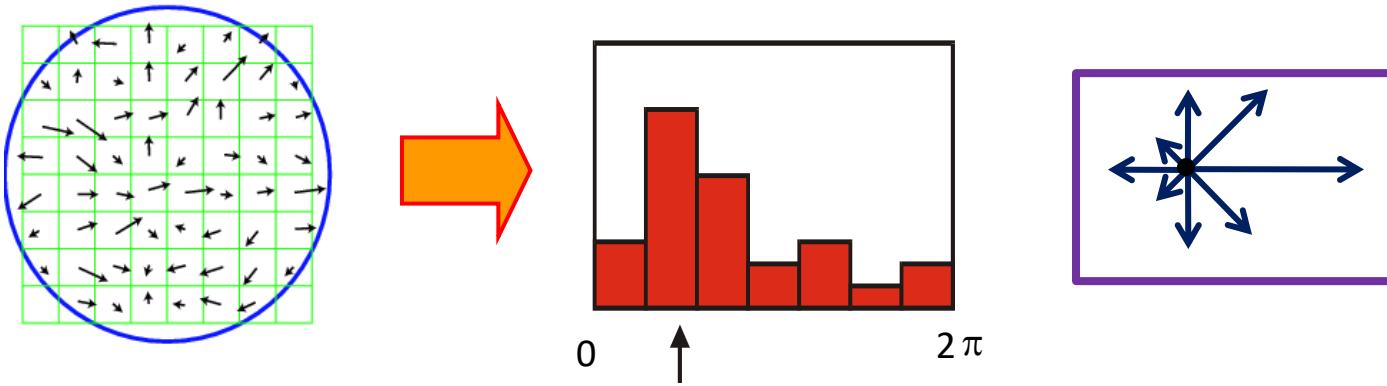
230

Late Due Date: SEP 13, 2024 11:59 PM

- You have five "late days" for the whole course without penalty.
- You will lose 10% each day for late projects, and late submissions will not be accepted after 1 week.
- There are extra credit points for several projects and the total score for those can be up to 110% of 14pts.

# Recap: SIFT descriptor

- Use histograms to bin pixels within sub-patches according to their gradient orientation.

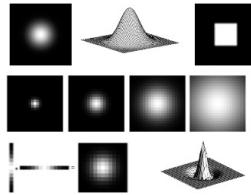


*Why subpatches?*  
*Why does SIFT have some illumination invariance?*  
*Why scale/rotation invariant?*

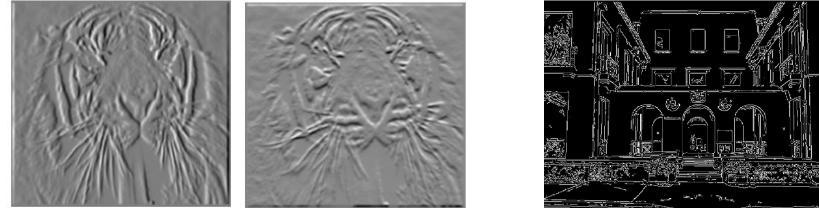
# Fitting & Hough Transform

# Our Vision Journey So Far

- Image Filtering



- Edge Detection



- Corner/Blob Detection



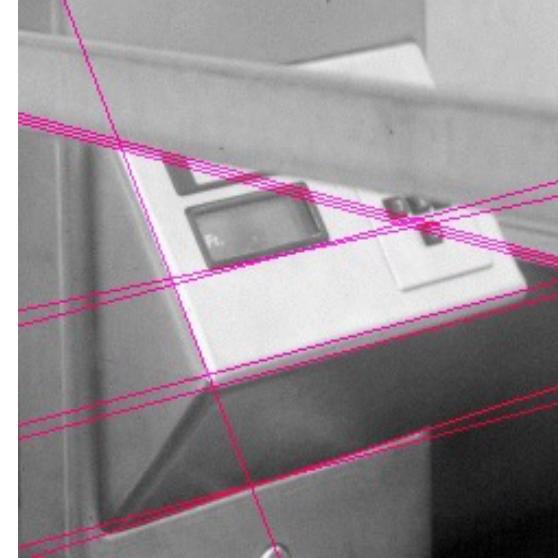
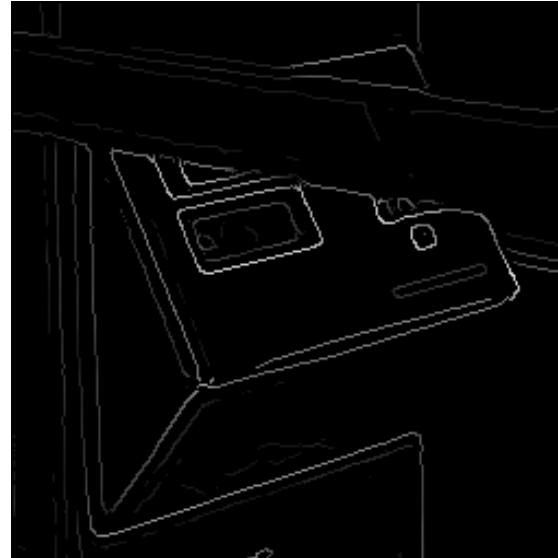
- SIFT



- What should we do next?

# Fitting

- Want to associate a model with observed features



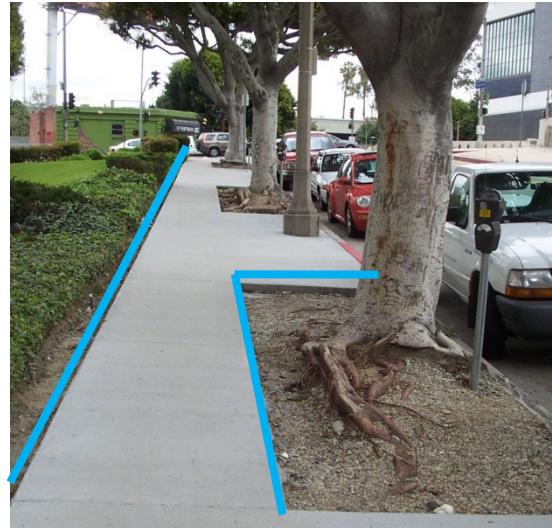
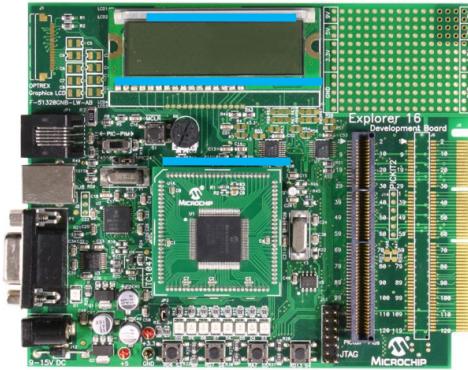
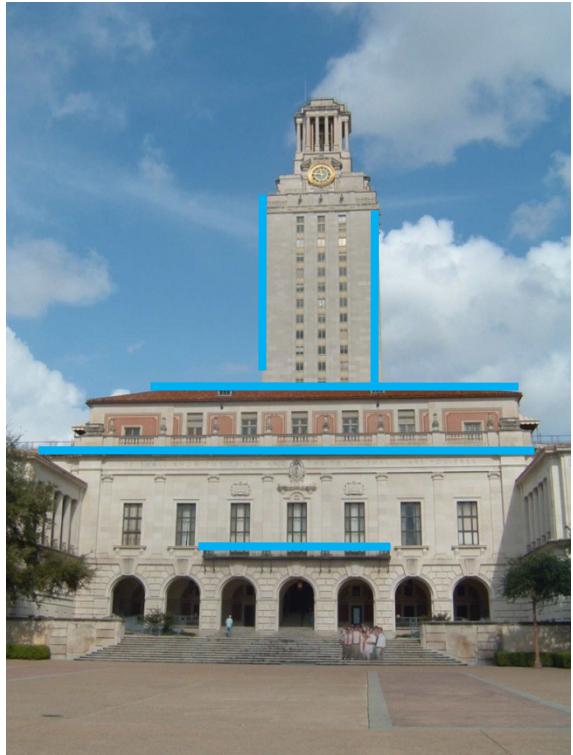
For example, the model could be a line, a circle, or an arbitrary shape.

# Fitting: Main idea

- Choose a **parametric model** to represent a **set of features**
- Membership criterion is not local
  - Can't tell whether a point belongs to a given model just by looking at that point
- Three main questions:
  - What model represents this set of features best?
  - Which of several model instances gets which feature?
  - How many model instances are there?
- Computational complexity is important
  - It is infeasible to examine every possible set of parameters and every possible combination of features

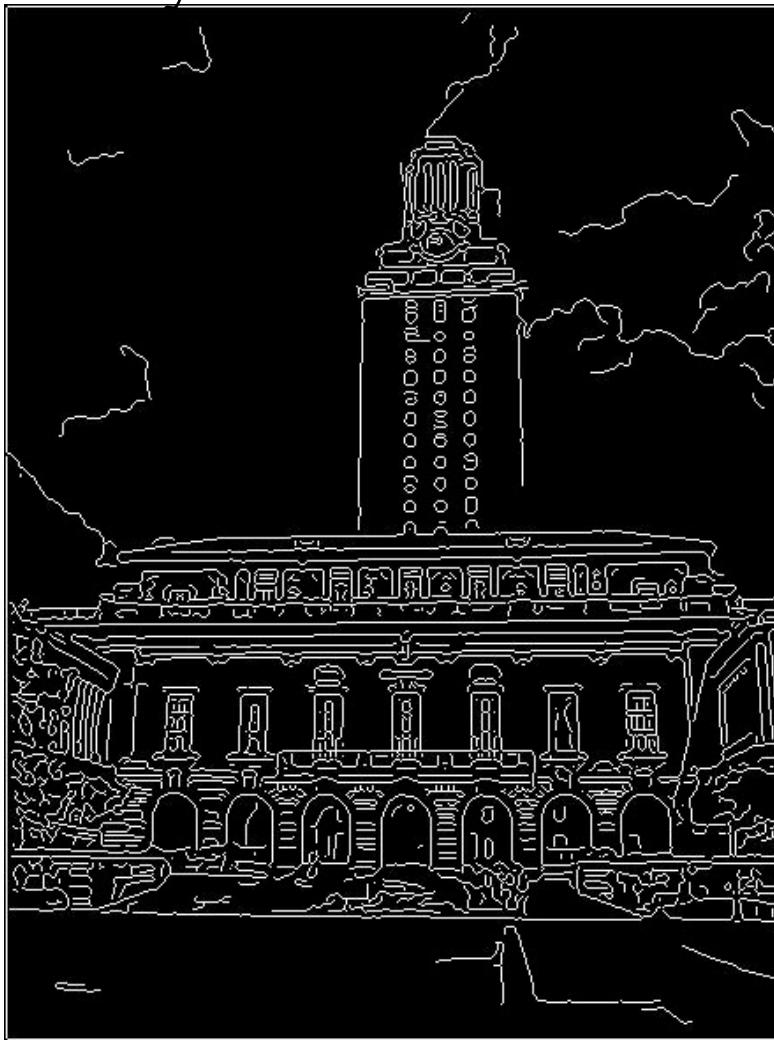
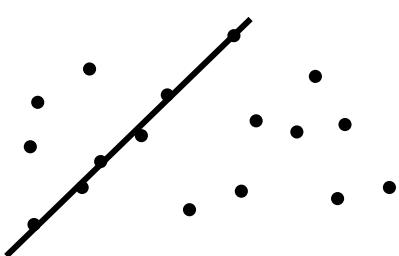
# Example: Line Fitting

- Why fit lines? Many objects characterized by presence of straight lines



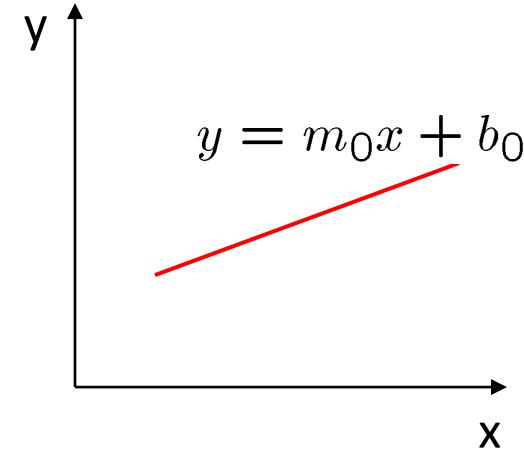
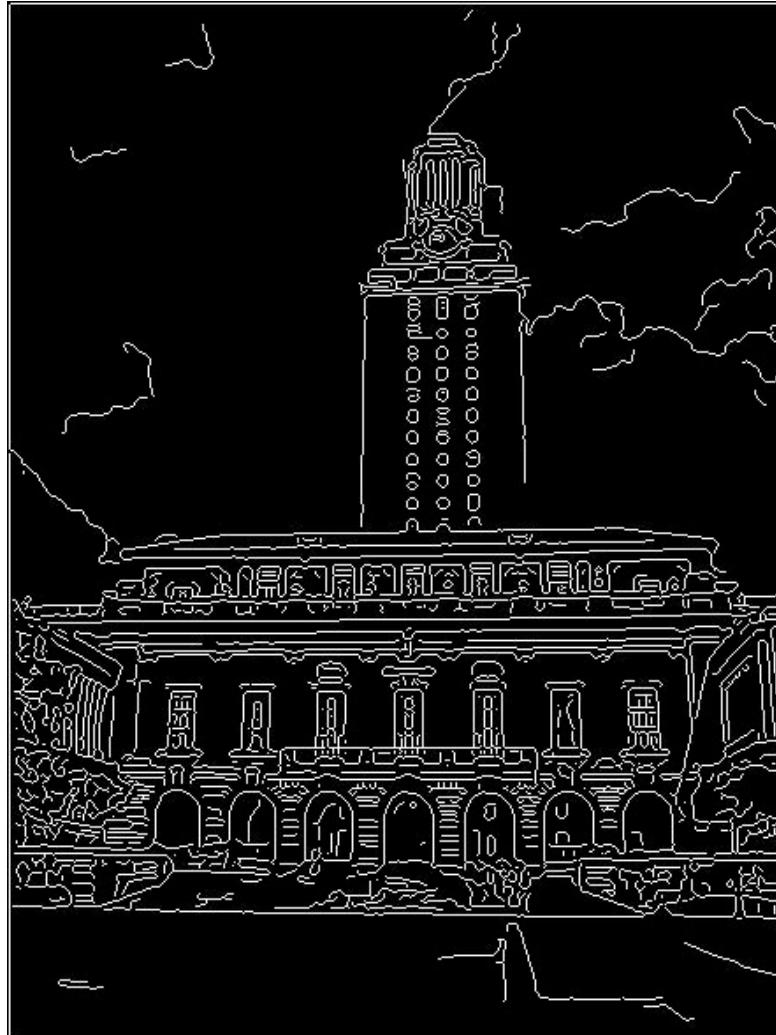
Wait, why aren't we done just by running edge detection?

# Difficulty of Line Fitting



- **Extra edge points (clutter), multiple models:**
  - which points go with which line, if any?
- Only some parts of each line detected, and some parts are **missing**:
  - how to find a line that bridges missing evidence?
- **Noise** in measured edge points, orientations:
  - how to detect true underlying parameters?

# How to Fit Lines to Images?



What are the challenges if we do this naively?

# Voting

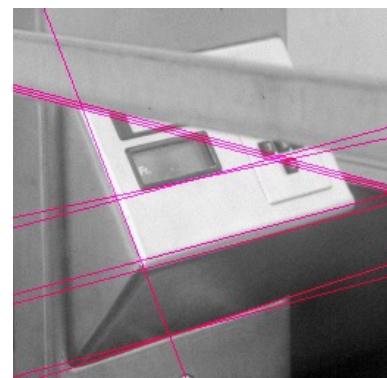
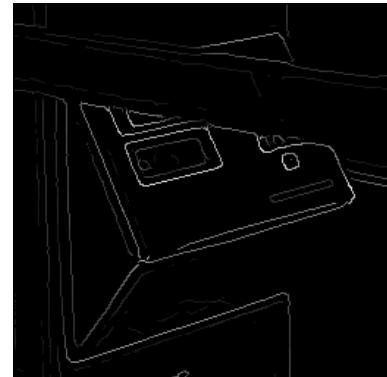
- It's not feasible to check all combinations of features by fitting a model to each possible subset.
- **Voting** is a general technique where we let the features *vote for all models that are compatible with it*.
  - Cycle through features, cast votes for model parameters.
  - Look for model parameters that receive a lot of votes.
- Noise & clutter features will cast votes too, *but* typically their votes should be inconsistent with the majority of “good” features.

# Fitting Lines: Hough transform

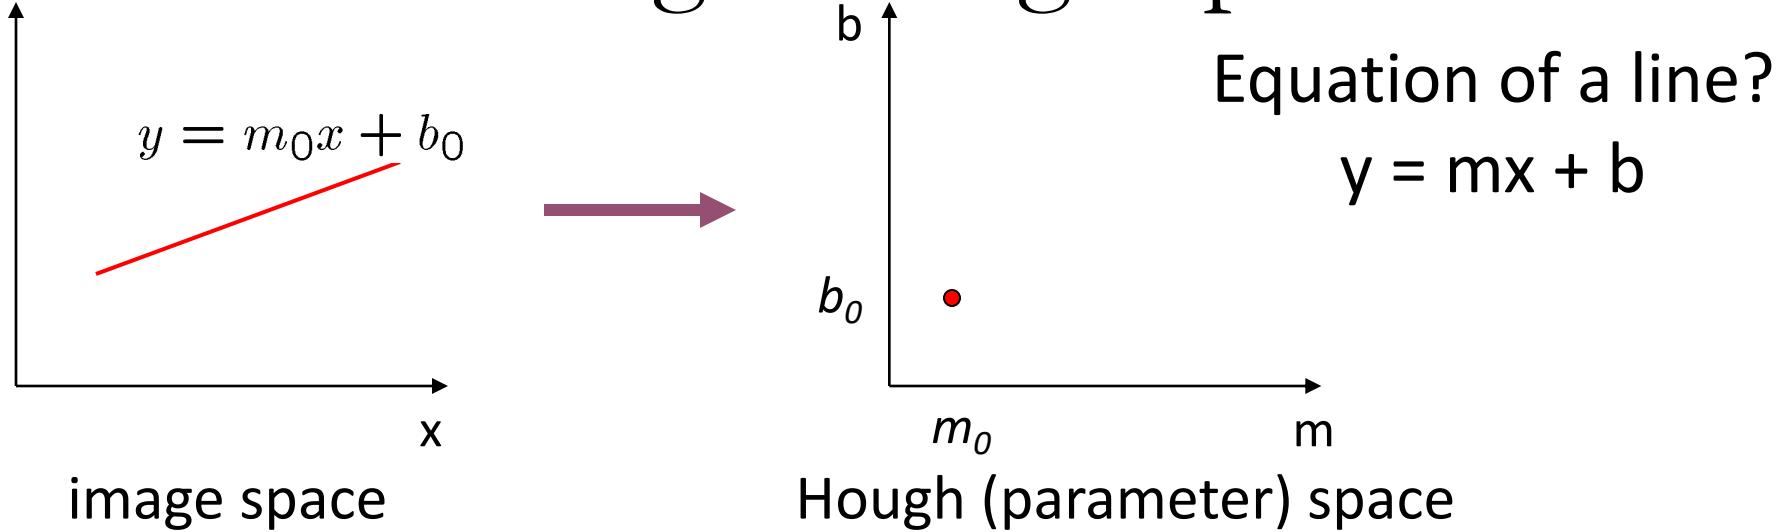
- Given points that belong to a line, what is the line?
- How many lines are there?
- Which points belong to which lines?
- **Hough Transform** is a voting technique that can be used to answer all of these questions.

Main idea:

1. Record vote for each possible line on which each edge point lies.
2. Look for lines that get many votes.



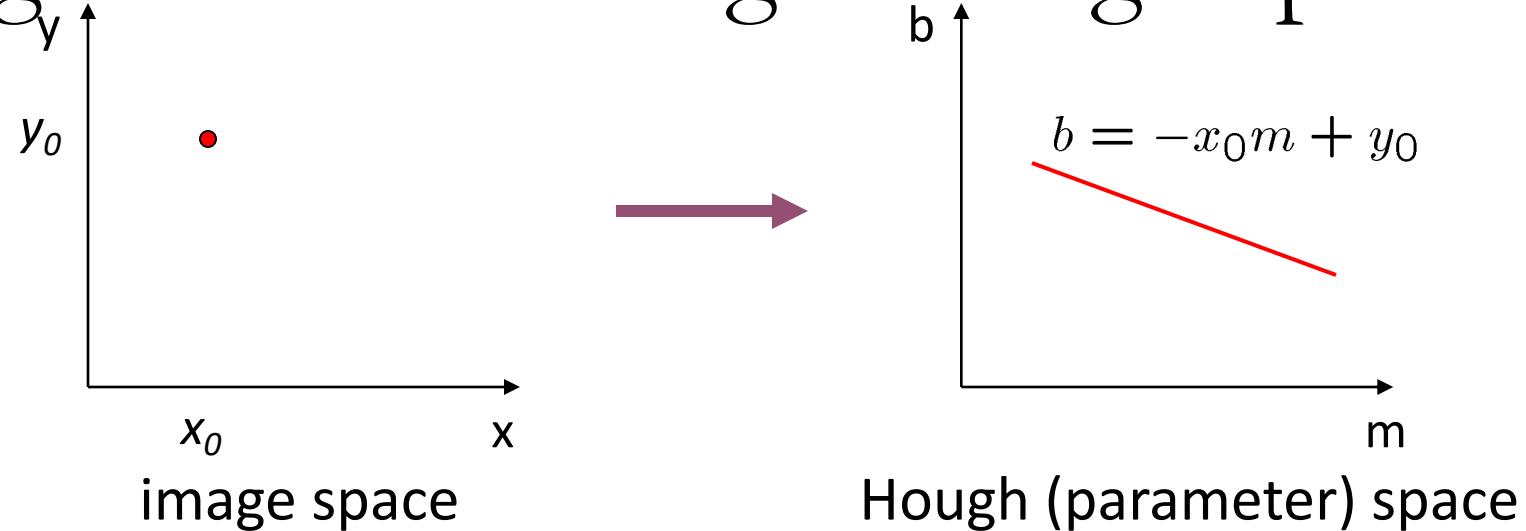
# Finding lines in an image: Hough space



Connection between image  $(x,y)$  and Hough  $(m,b)$  spaces

- A line in the image corresponds to a point in Hough space
- To go from image space to Hough space:
  - given a set of points  $(x,y)$ , find all  $(m,b)$  such that  $y = mx + b$

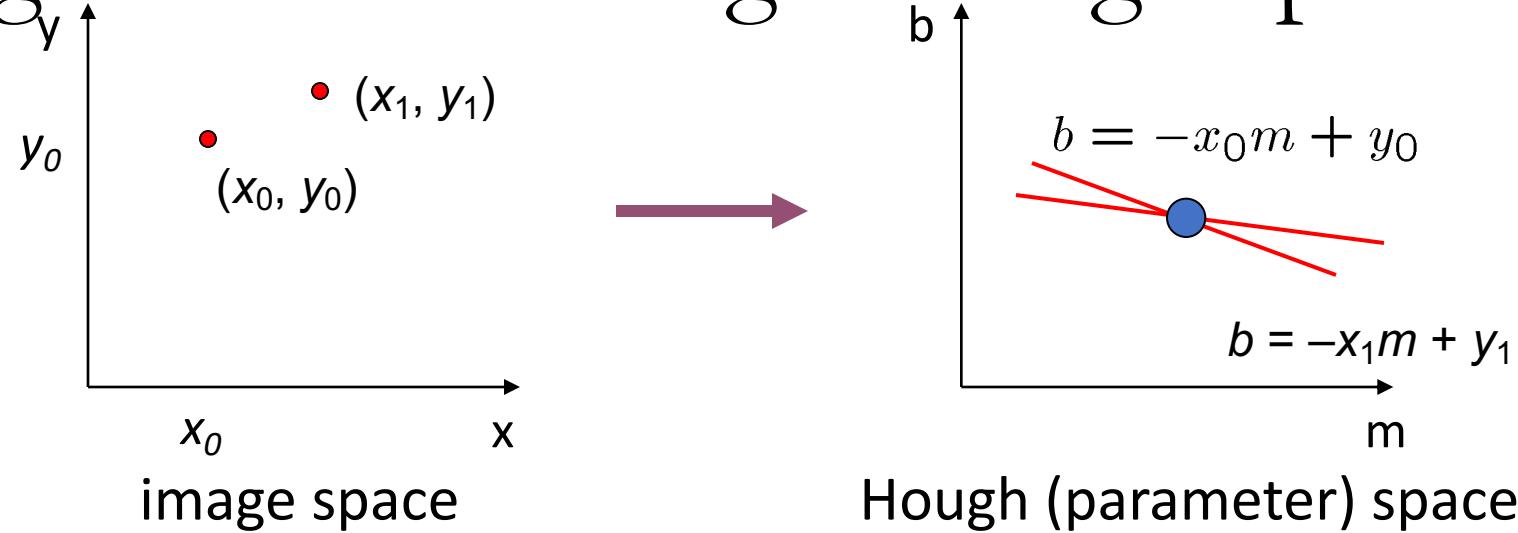
# Finding lines in an image: Hough space



Connection between image  $(x,y)$  and Hough  $(m,b)$  spaces

- A line in the image corresponds to a point in Hough space
- To go from image space to Hough space:
  - given a set of points  $(x,y)$ , find all  $(m,b)$  such that  $y = mx + b$
  - What does a point  $(x_0, y_0)$  in the image space map to?
    - Answer: the solutions of  $b = -x_0 m + y_0$
    - this is a line in Hough space

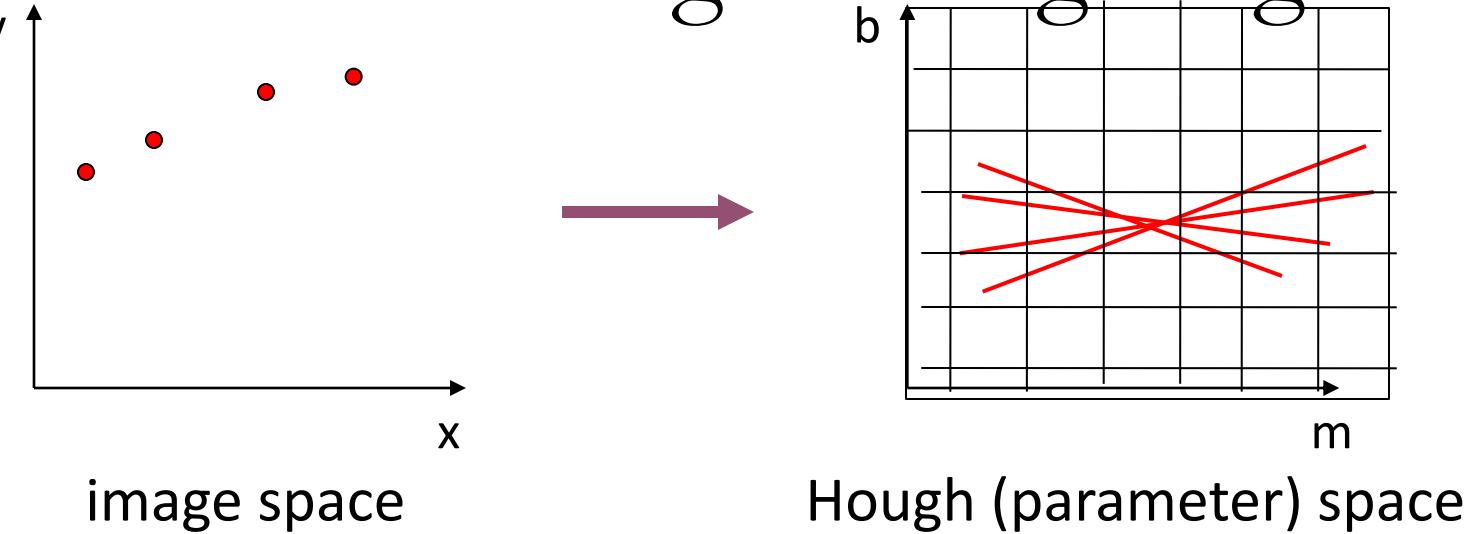
# Finding lines in an image: Hough space



What are the line parameters for the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?

- It is the intersection of the lines  $b = -x_0m + y_0$  and  $b = -x_1m + y_1$

# Finding lines in an image: Hough algorithm

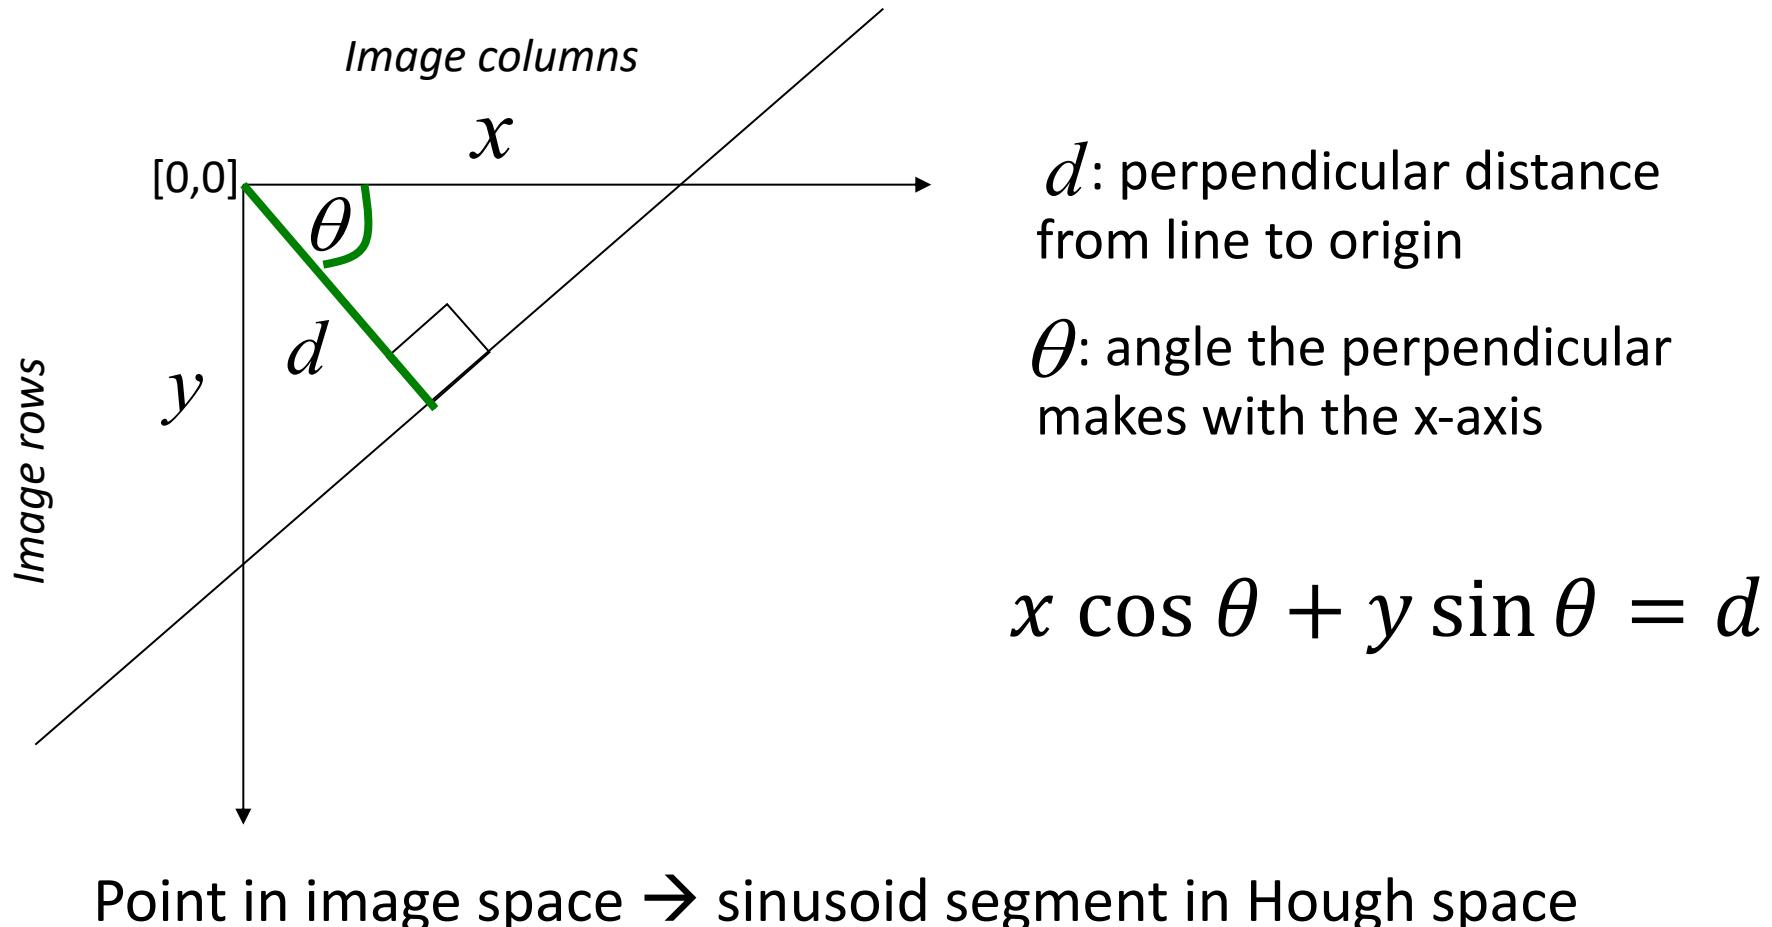


How can we use this to find the most likely parameters ( $m, b$ ) for the most prominent line in the image space?

- Let each edge point in image space *vote* for a set of possible parameters in Hough space
- Accumulate votes in discrete set of bins; parameters with the most votes indicate line in image space.

# Polar representation for lines

Issues with usual  $(m, b)$  parameter space: can take on infinite values, undefined for vertical lines.



# Hough transform algorithm

Using the polar parameterization:

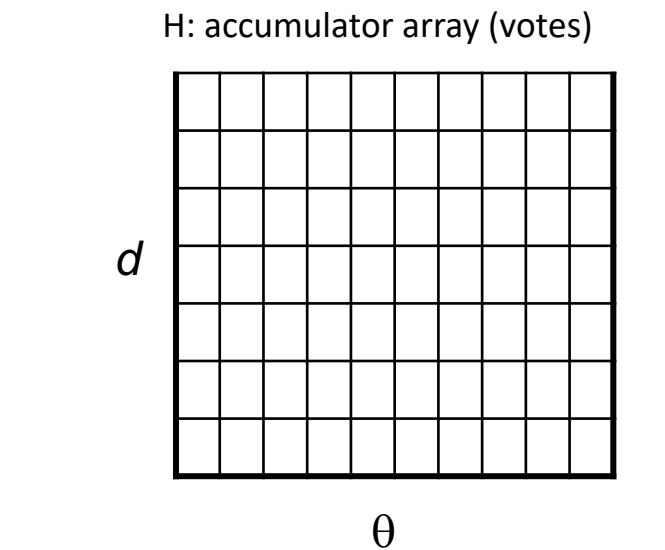
$$x \cos \theta + y \sin \theta = d$$

Basic Hough transform algorithm

1. Initialize  $H[d, \theta] = 0$
2. for each edge point  $I[x, y]$  in the image  
    for  $\theta = [\theta_{\min} \text{ to } \theta_{\max}]$  // some quantization

$$x \cos \theta + y \sin \theta = d$$
$$H[d, \theta] += 1$$

3. Find the value(s) of  $(d, \theta)$  where  $H[d, \theta]$  is maximum
4. The detected line in the image is given by



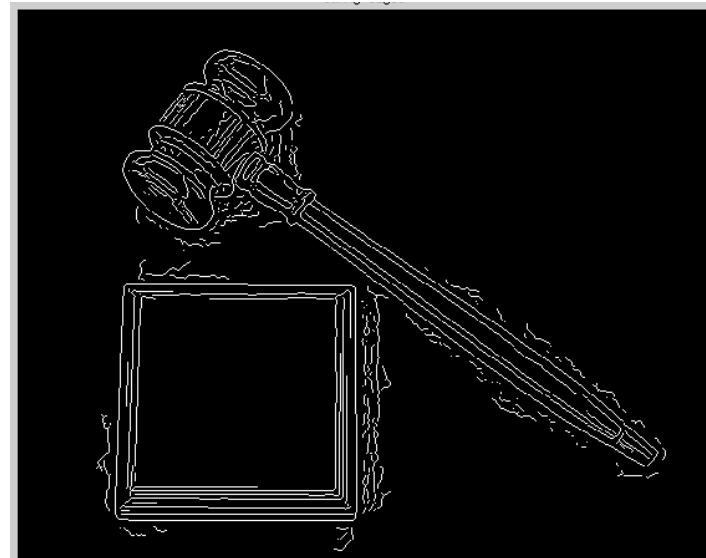
$$x \cos \theta + y \sin \theta = d$$

# Hough Transform Example

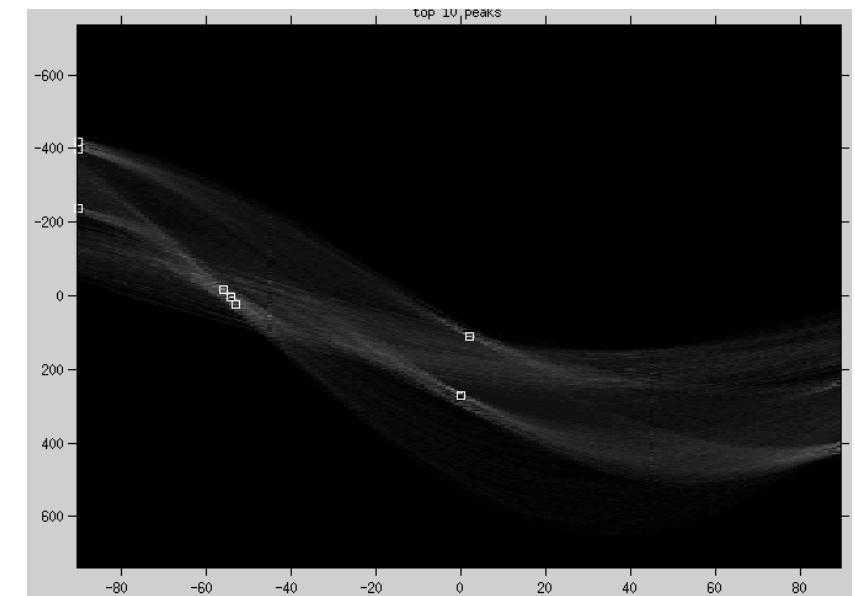
Original image

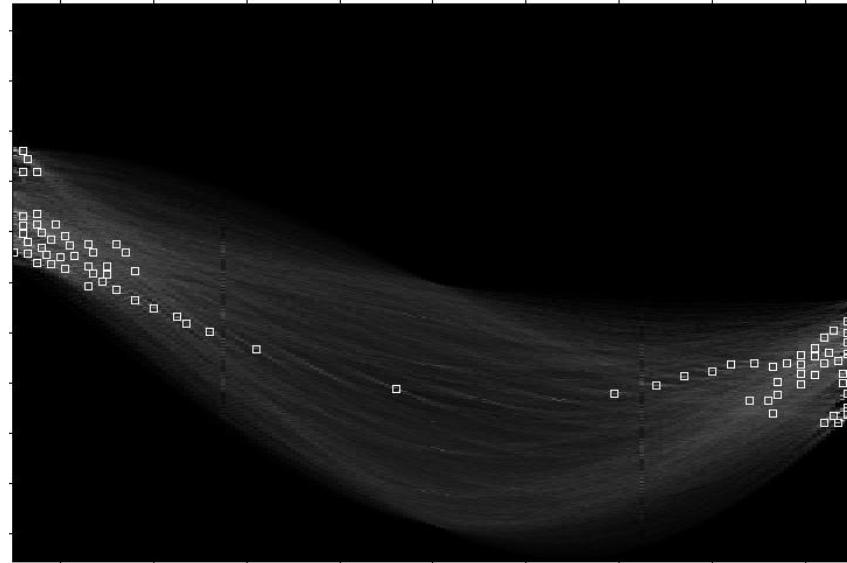
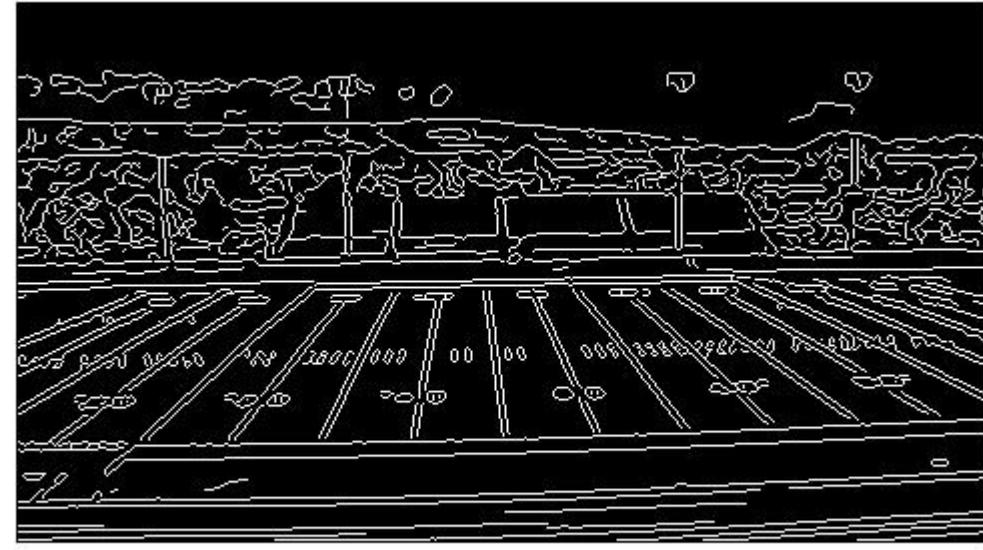


Canny edges



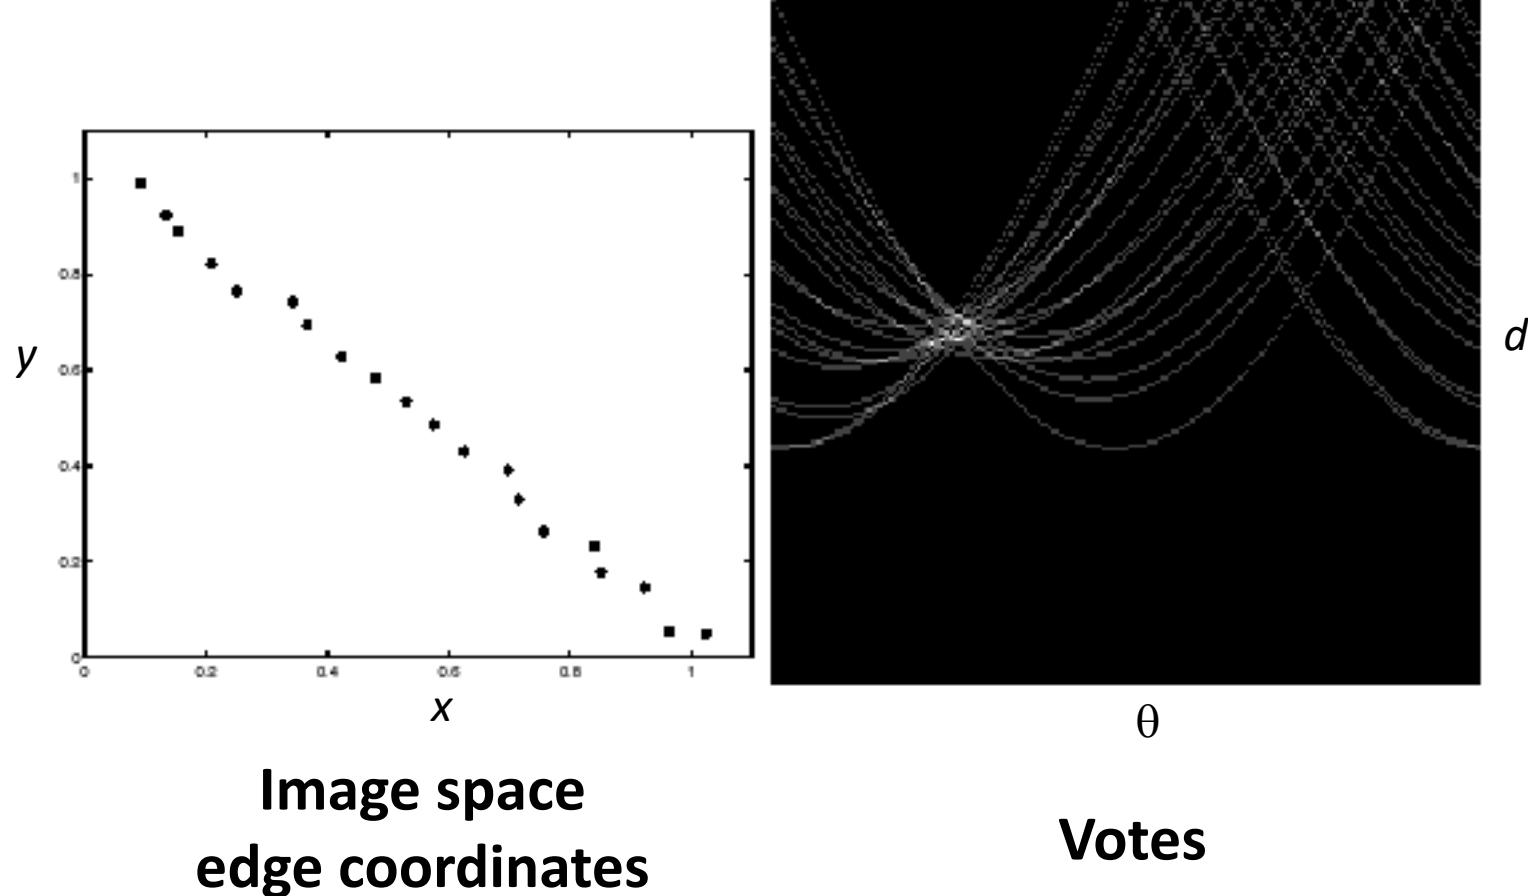
Vote space and top peaks





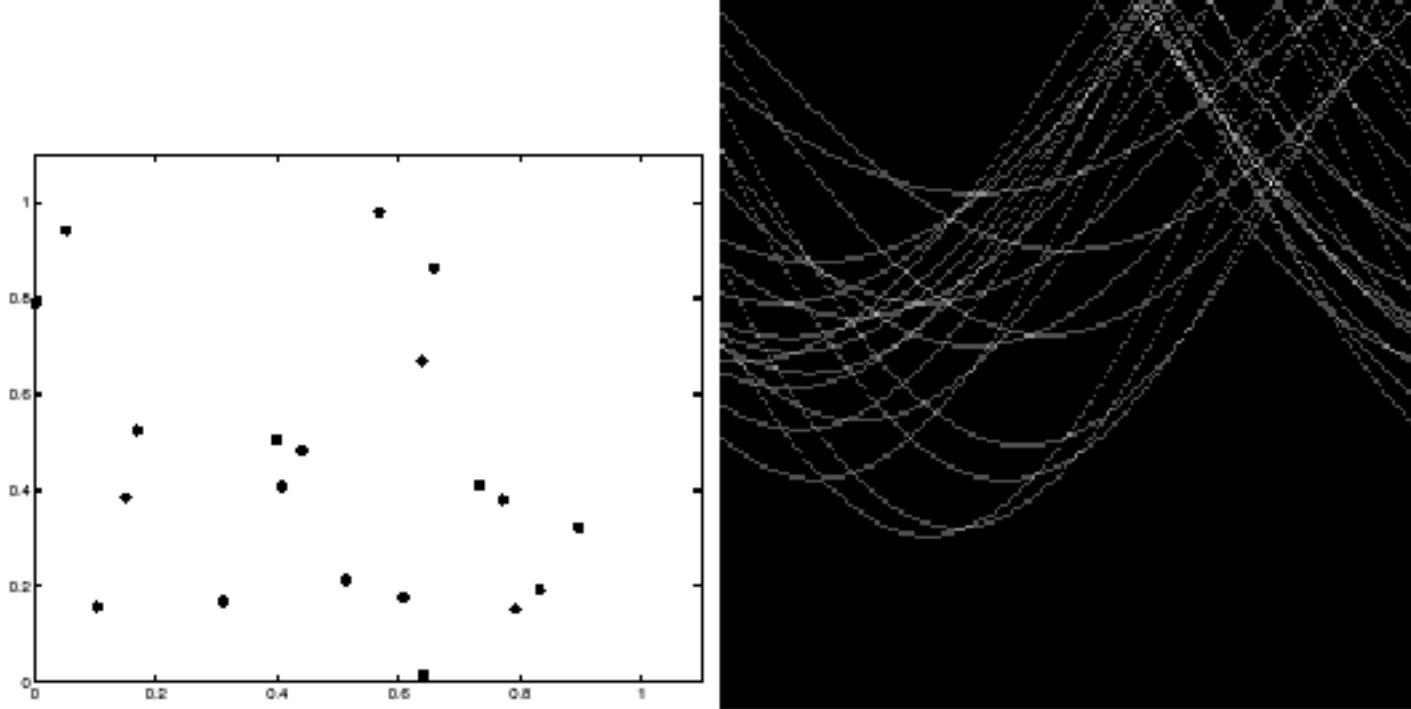
Showing longest segments found

# Impact of noise on Hough



What difficulty does this present for an implementation?

# Impact of noise on Hough



**Image space  
edge coordinates**

**Votes**

Here, everything appears to be “noise”, or random edge points, but we still see peaks in the vote space.

# Extensions

Extension 1: Use the image gradient

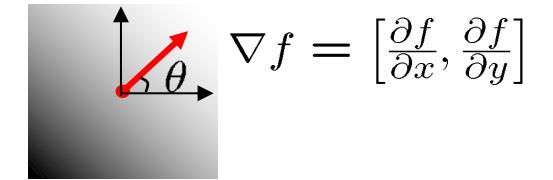
1. Initialize  $H[d, \theta] = 0$
2. for each edge point  $I[x, y]$  in the image  
 $\theta$  = gradient at  $(x, y)$

$$x \cos \theta + y \sin \theta = d$$

$$H[d, \theta] += 1$$

3. Find the value(s) of  $(d, \theta)$  where  $H[d, \theta]$  is maximum
4. The detected line in the image is given by

$$x \cos \theta + y \sin \theta = d$$



$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

Reduces degree of freedom

Extension 2

- give more votes for stronger edges:  $H[d, \theta] += \text{magnitude}(\text{gradient})$

Extension 3

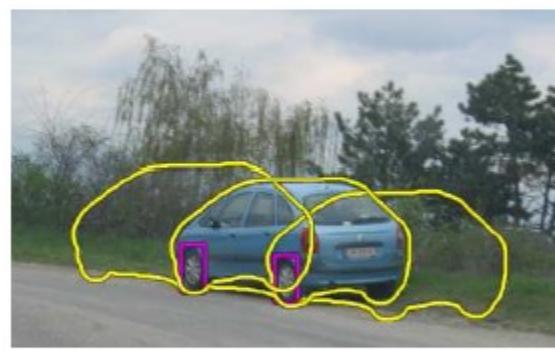
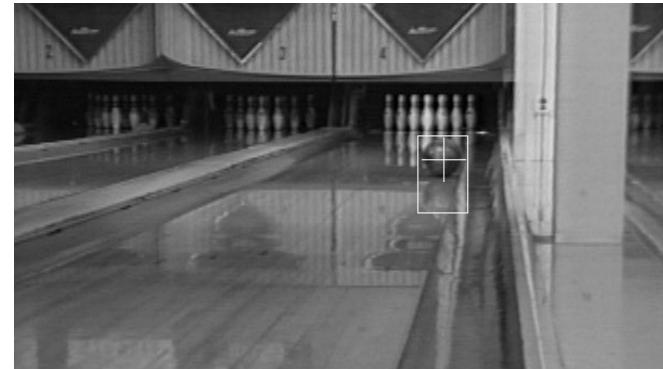
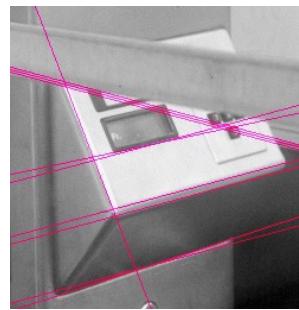
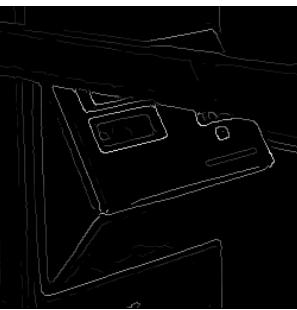
- change the sampling of  $(d, \theta)$  to give more/less resolution

Extension 4

- The same procedure can be used with circles, squares, or any other shape...

# How to fit Circles or any Shapes?

- Want to associate a model with observed features



[Fig from Marszalek & Schmid, 2007]

For example, the model could be a line, a circle, or an arbitrary shape.

# Hough transform for circles

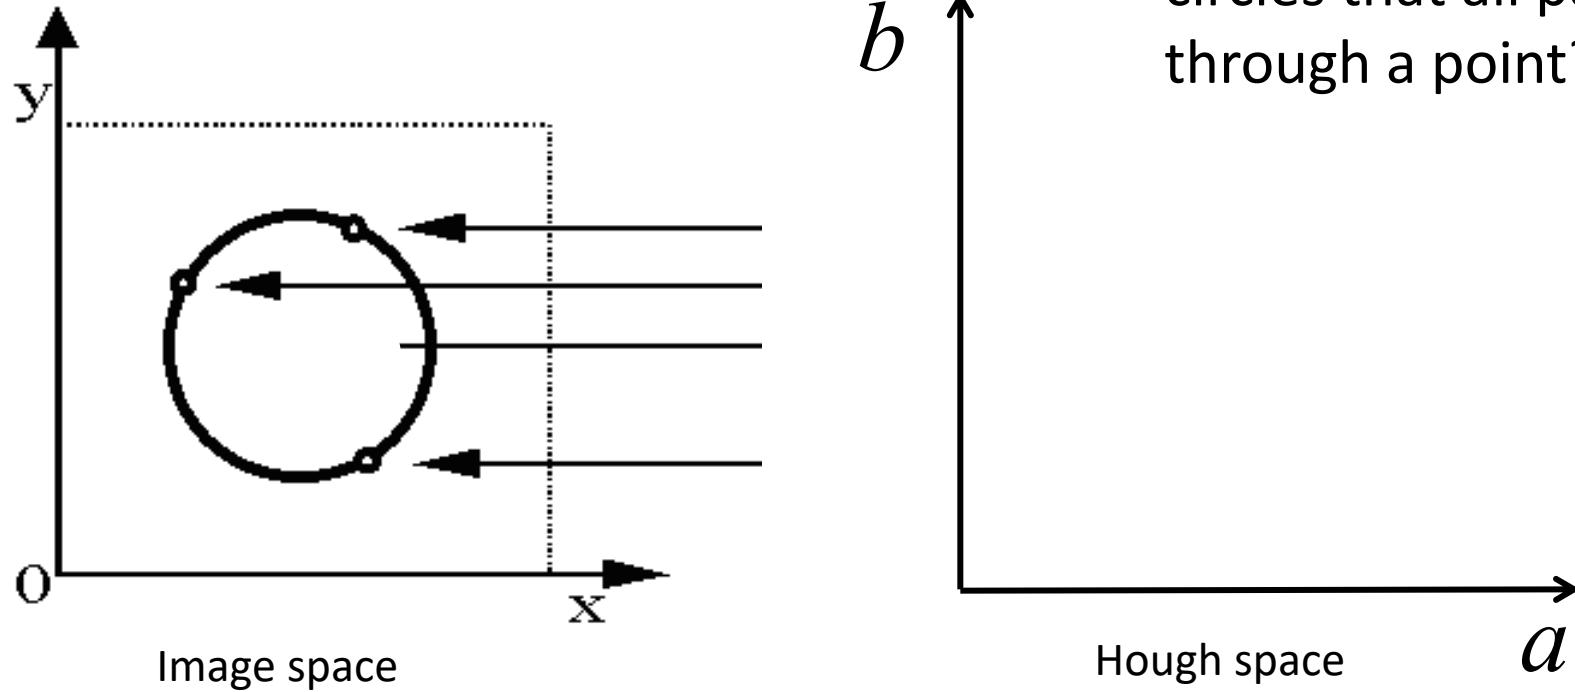
- Circle: center  $(a, b)$  and radius  $r$

Equation of circle?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius  $r$

Equation of set of circles that all pass through a point?

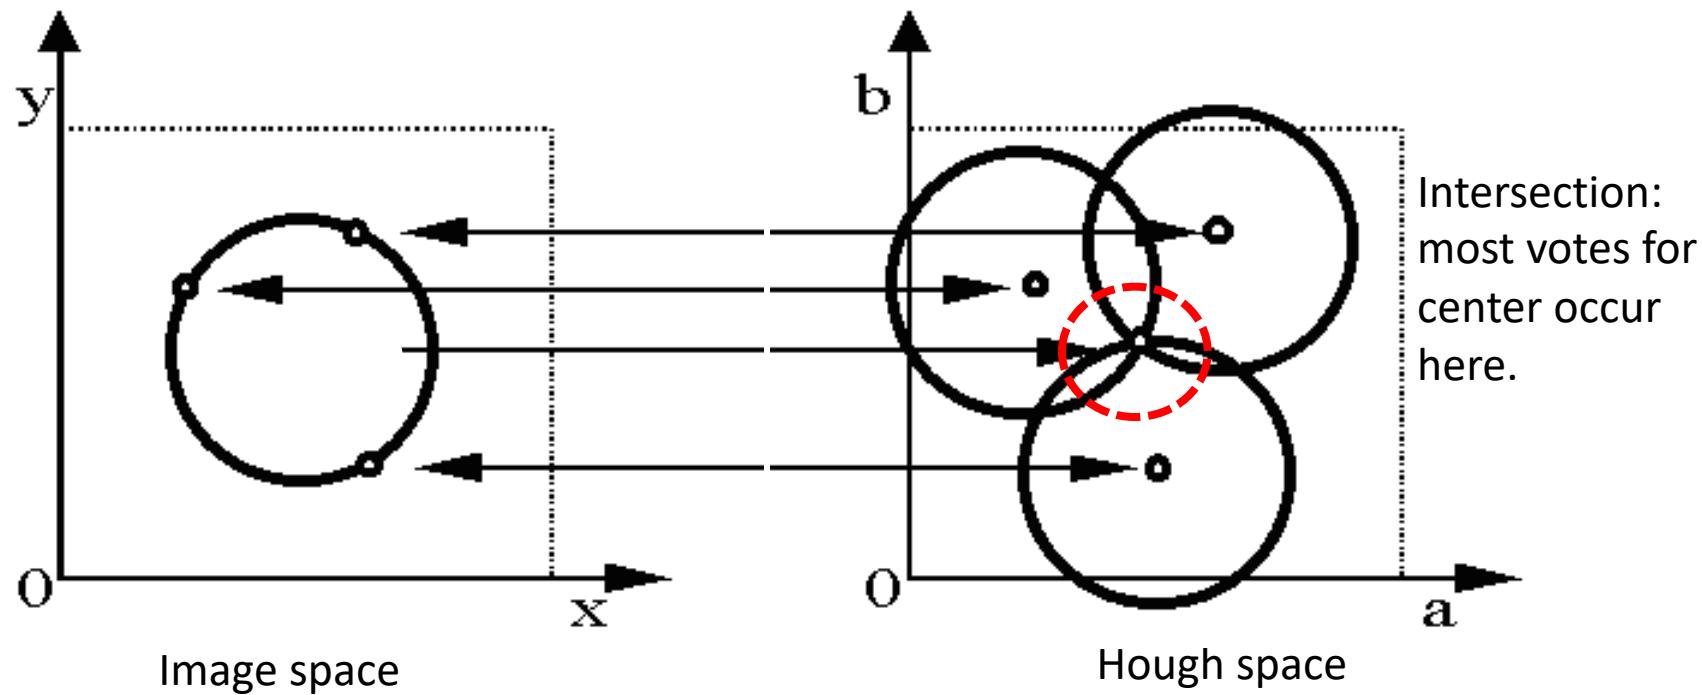


# Hough transform for circles

- Circle: center  $(a, b)$  and radius  $r$

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius  $r$

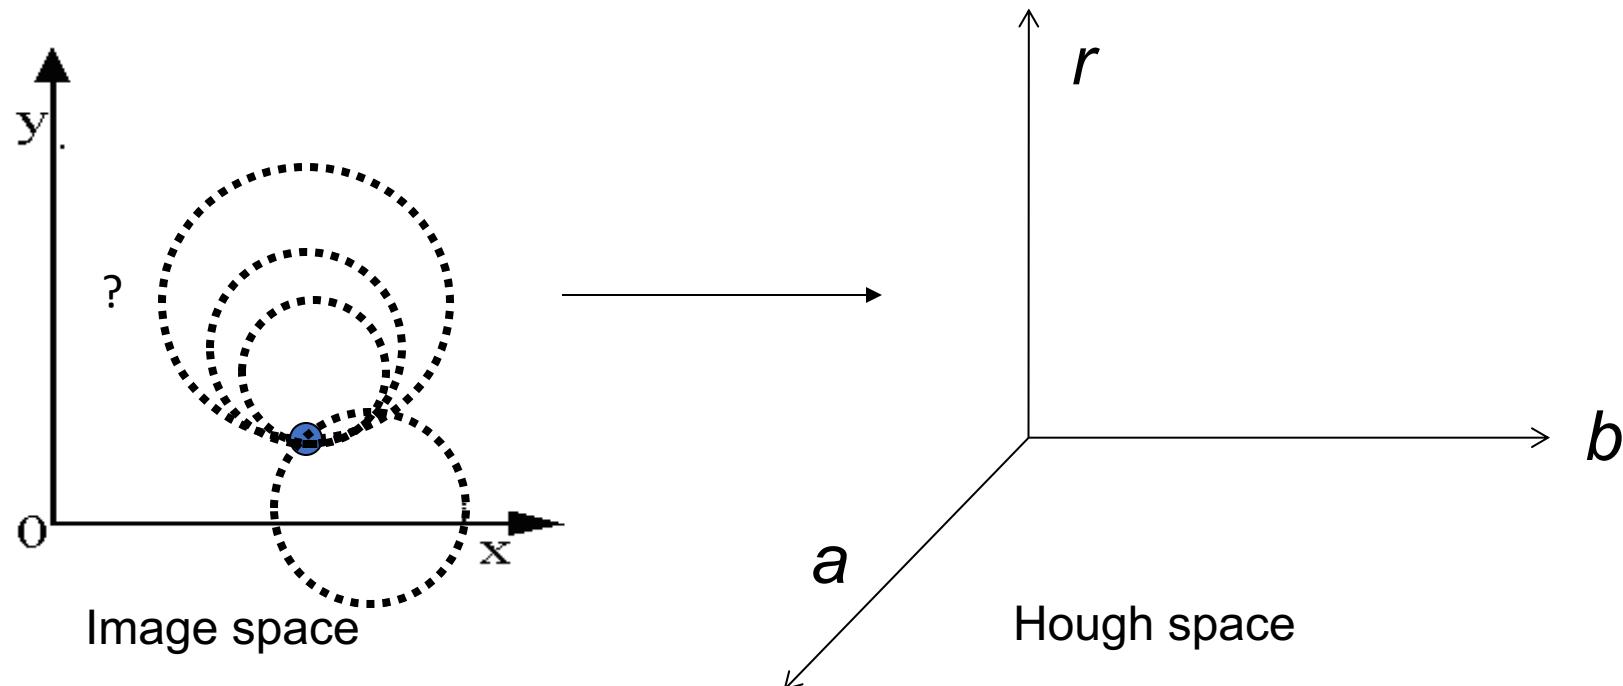


# Hough transform for circles

- Circle: center  $(a, b)$  and radius  $r$

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius  $r$

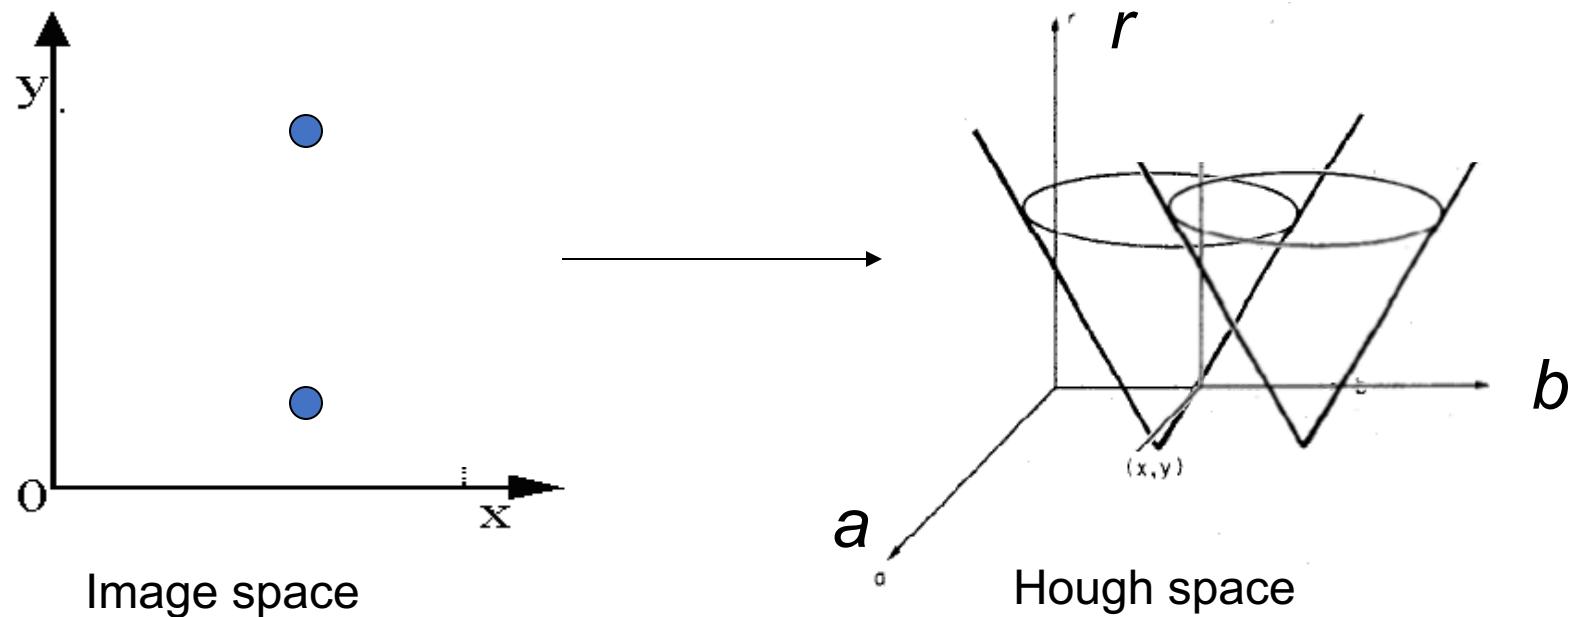


# Hough transform for circles

- Circle: center  $(a, b)$  and radius  $r$

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius  $r$

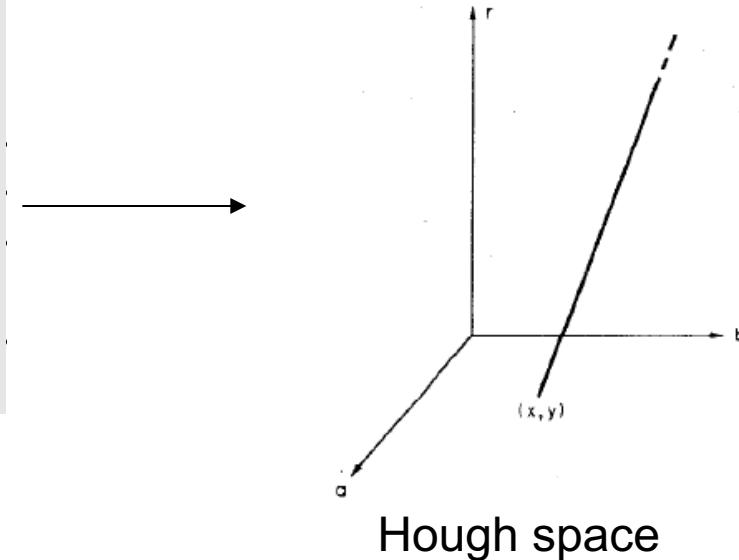
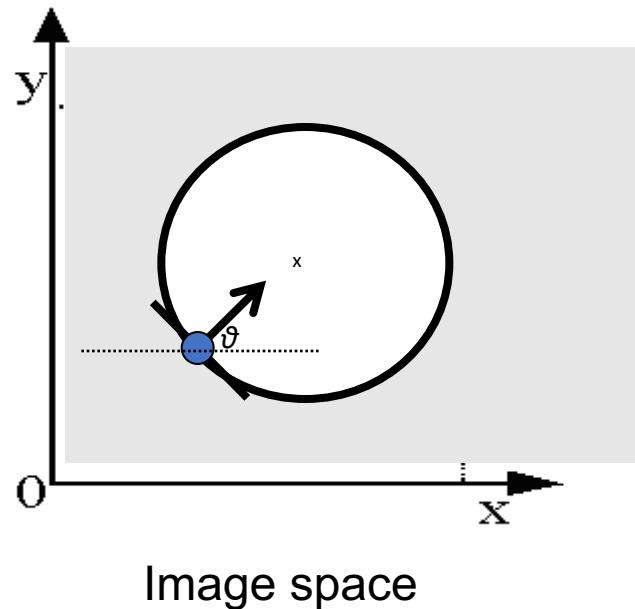


# Hough transform for circles

- Circle: center  $(a, b)$  and radius  $r$

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius  $r$ , **known** gradient direction



# Hough transform for circles

For every edge pixel  $(x,y)$  :

    For each possible radius value  $r$ :

        For each possible gradient direction  $\theta$ :

*// or use estimated gradient at  $(x,y)$*

$a = x - r \cos(\theta)$  // column

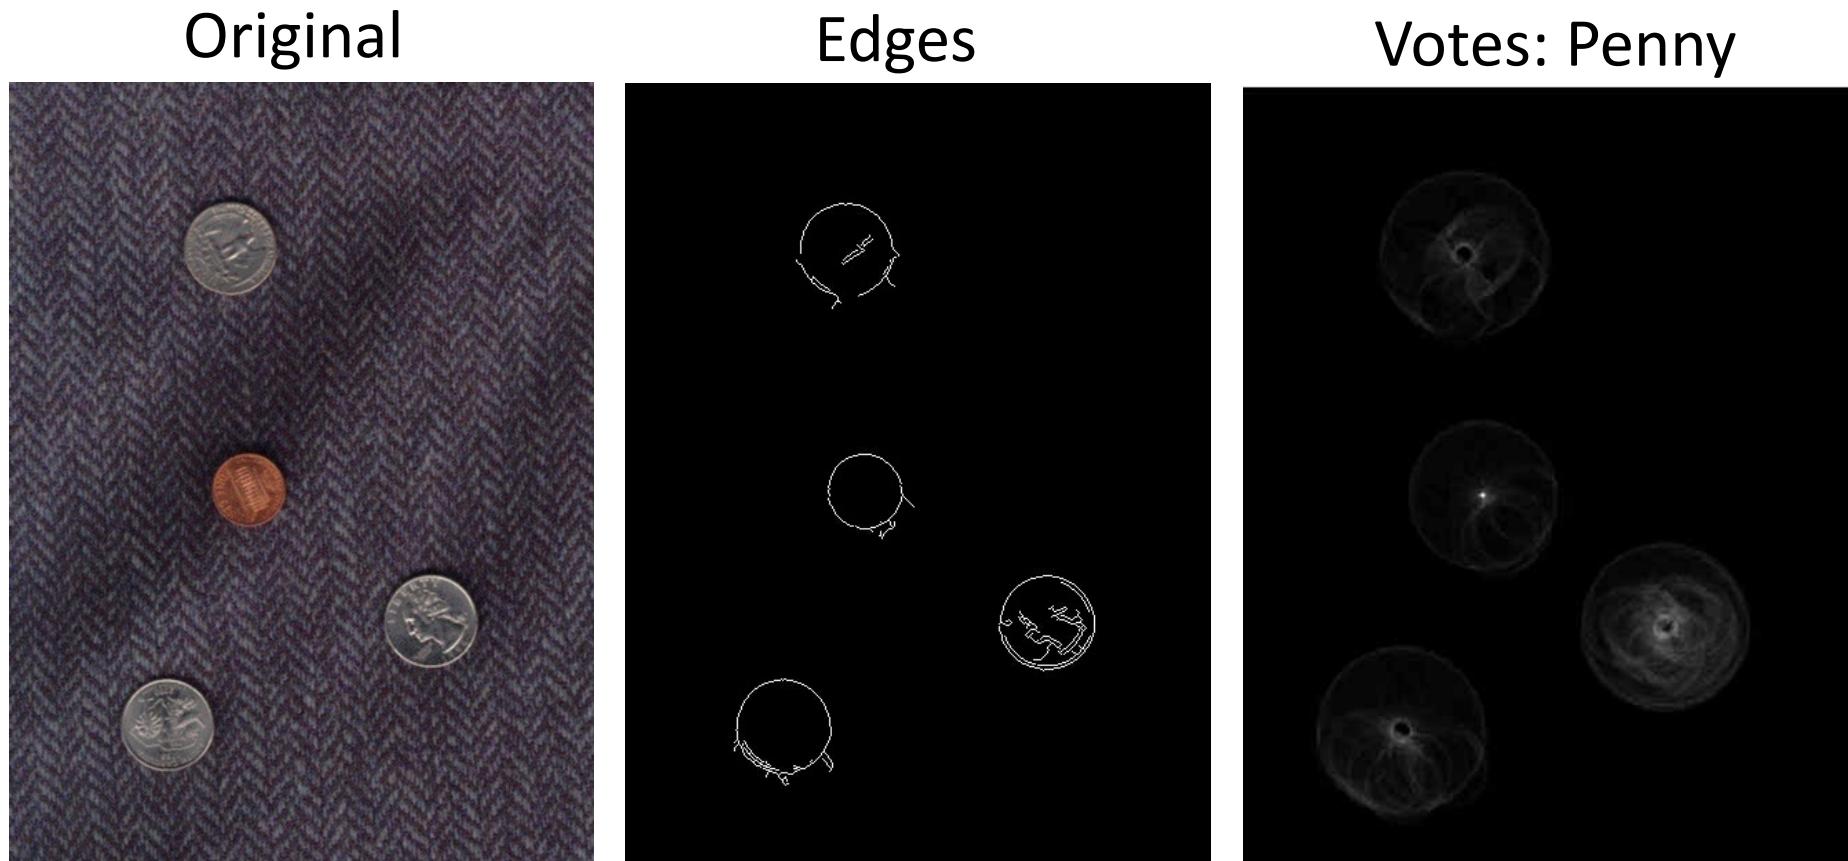
$b = y - r \sin(\theta)$  // row

$H[a,b,r] += 1$

    end

end

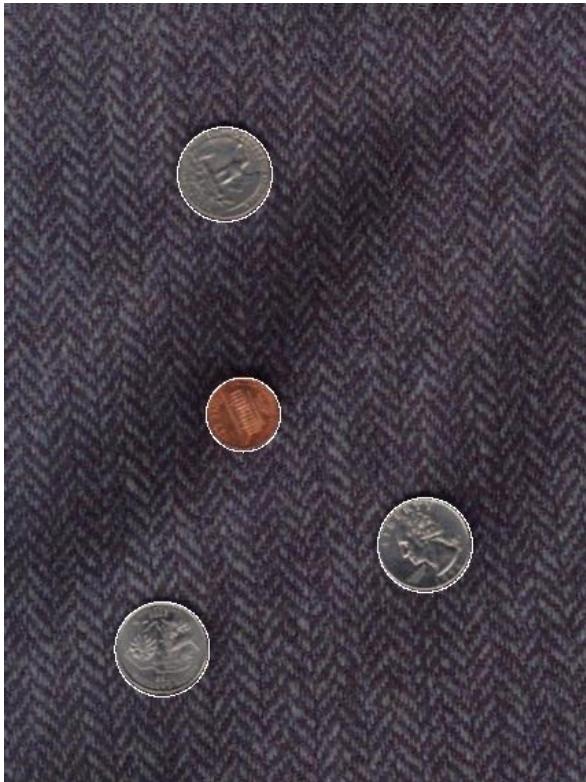
# Example: detecting circles with Hough



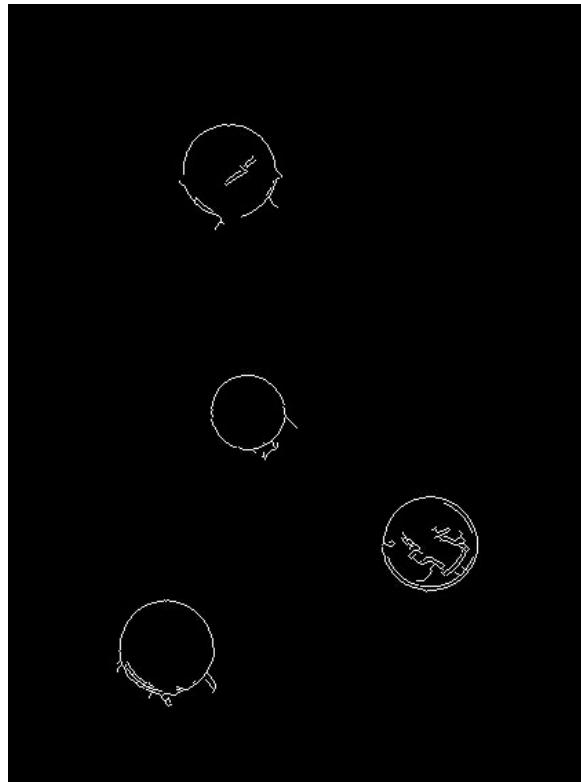
Note: a different Hough transform (with separate accumulators) was used for each circle radius (quarters vs. penny).

# Example: detecting circles with Hough

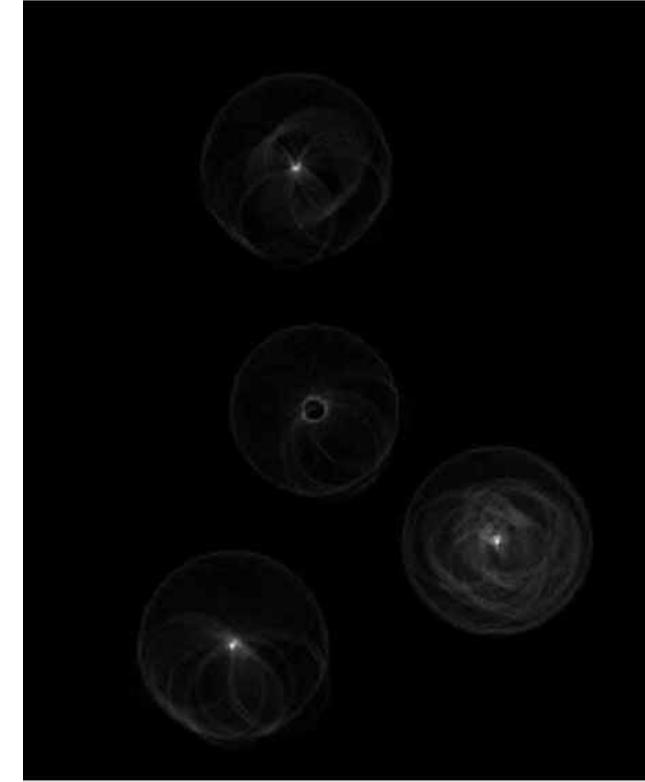
Combined detections



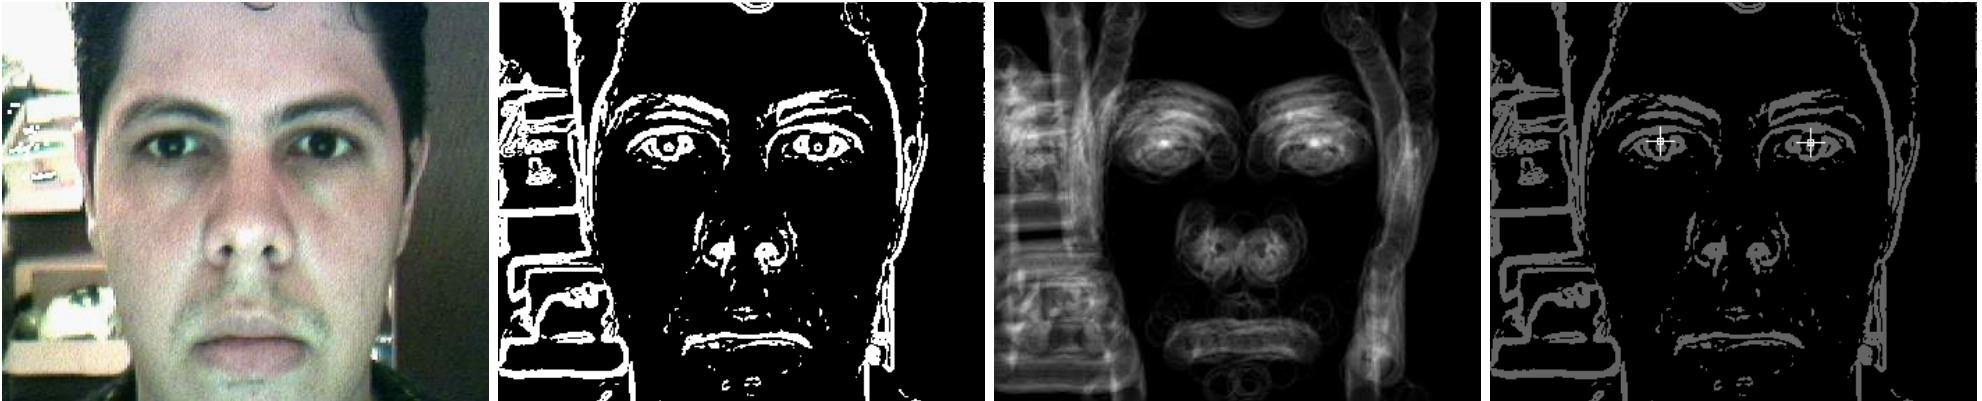
Edges



Votes: Quarter



# Example: iris detection



Gradient+threshold

Hough space  
(fixed radius)

Max detections

- Hemerson Pistori and Eduardo Rocha Costa <http://rsbweb.nih.gov/ij/plugins/hough-circles.html>

# Example: iris detection



Figure 2. Original image



Figure 3. Distance image



Figure 4. Detected face region

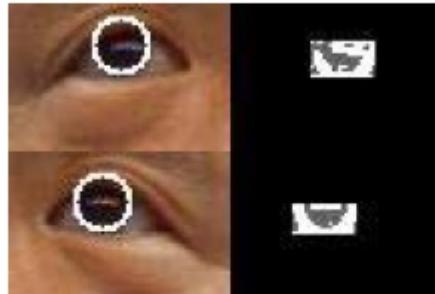


Figure 14. Looking upward

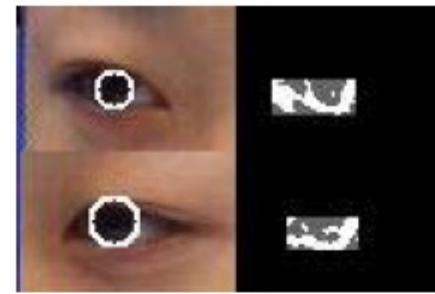


Figure 15. Looking sideways

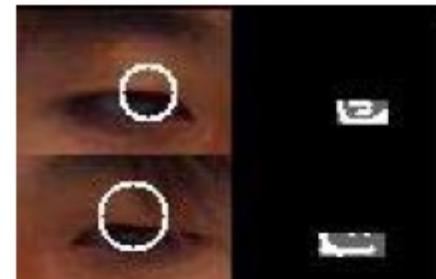


Figure 16. Looking downward

- An Iris Detection Method Using the Hough Transform and Its Evaluation for Facial and Eye Movement, by Hideki Kashima, Hitoshi Hongo, Kunihito Kato, Kazuhiko Yamamoto, ACCV 2002.

# Voting: practical tips

- Minimize irrelevant tokens first
- Choose a good grid / discretization



- Vote for neighbors, also (smoothing in accumulator array)
- Use direction of edge to reduce parameters by 1
- Keep track of which points contributed to a vote
  - Allows you to quickly see which points correspond to a “winning” vote.

# Hough transform: pros and cons

## Pros

- All points are processed independently, so can cope with occlusion, gaps
- Some robustness to noise: noise points unlikely to contribute *consistently* to any single bin
- Can detect multiple instances of a model in a single pass

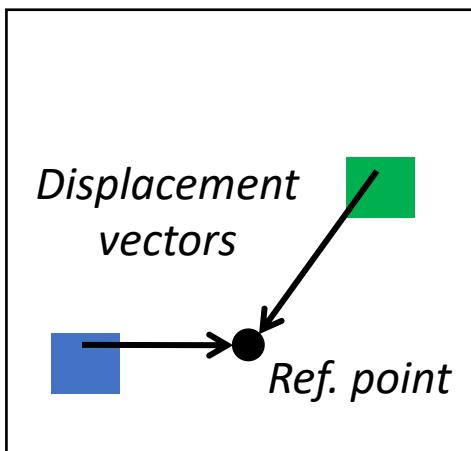
## Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization: can be tricky to pick a good grid size

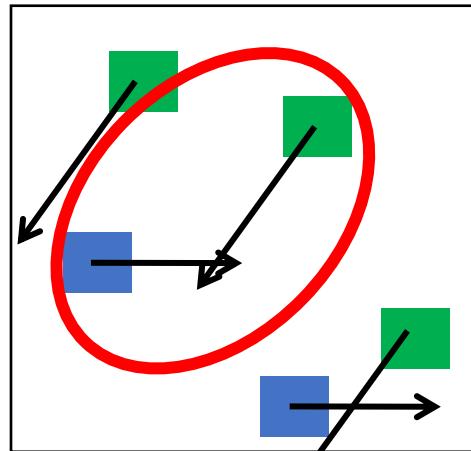
# Generalized Hough Transform

What if we want to detect arbitrary shapes?

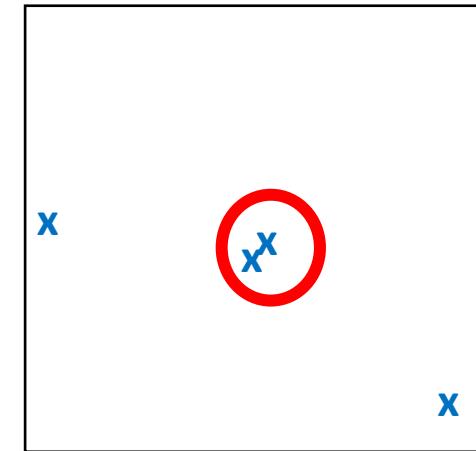
## Intuition:



Model image



Novel image

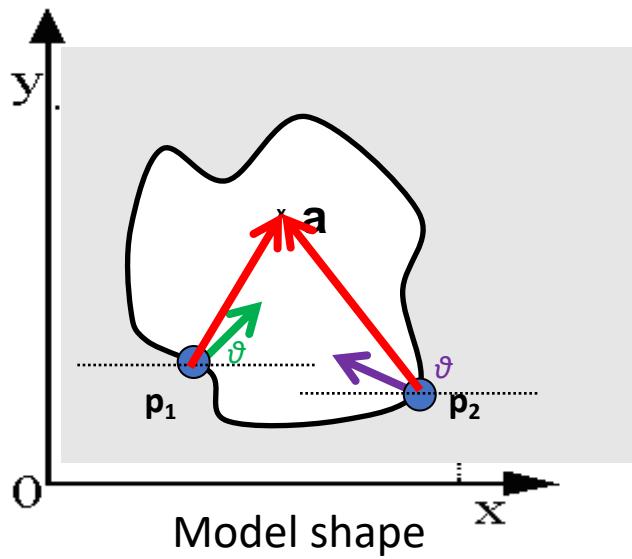


Vote space

Now suppose those colors encode gradient directions...

# Generalized Hough Transform

- Define a model shape by its boundary points and a reference point.



		...
		...
:		

## Offline procedure:

At each boundary point, compute displacement vector:  $r = a - p_i$ .

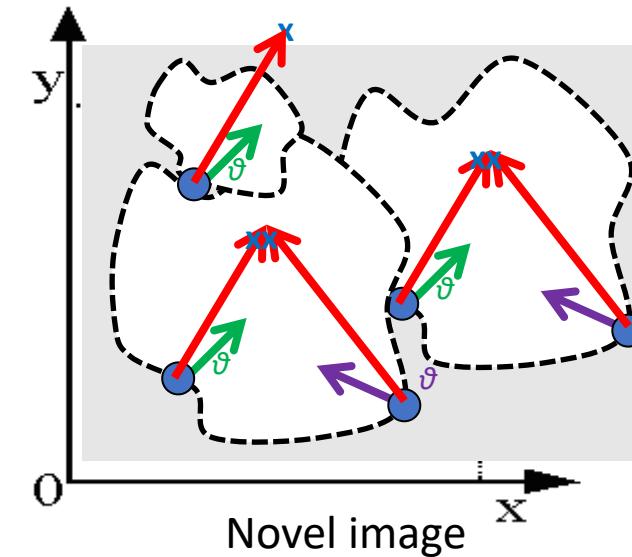
Store these vectors in a table indexed by gradient orientation  $\theta$ .

# Generalized Hough Transform

## Detection procedure:

For each edge point:

- Use its gradient orientation  $\theta$  to index into stored table
- Use retrieved  $r$  vectors to vote for reference point

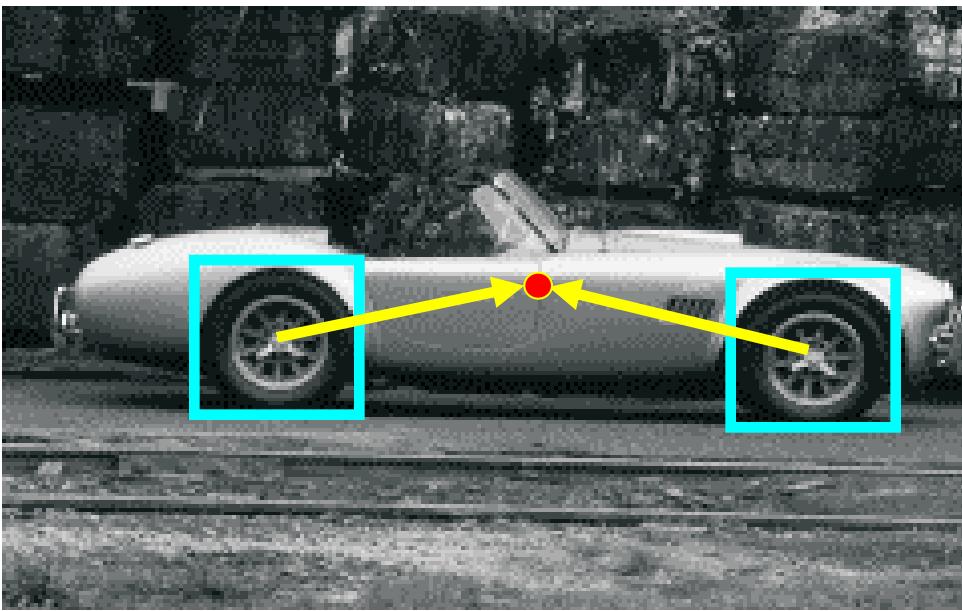


		...
		...
:		

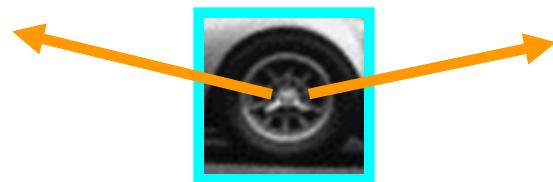
Assuming translation is the only transformation here, i.e., orientation and scale are fixed.

# Generalized Hough for object detection

- Instead of indexing displacements by gradient orientation, index by matched local patterns.



training image



“visual codeword” with  
displacement vectors

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

# Generalized Hough for object detection

- Instead of indexing displacements by gradient orientation, index by “visual codeword”



test image

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

# Summary

- **Fitting** problems require finding any supporting evidence for a model, even within clutter and missing features.
  - associate features with an explicit model
- **Voting** approaches, such as the **Hough transform**, make it possible to find likely model parameters without searching all combinations of features.
  - Hough transform approach for lines, circles, ..., arbitrary shapes defined by a set of boundary points, recognition from patches.

# Questions?

See you next time!